

Machine Learning

Kon Yi

September 28, 2026

Preface

These are my notes for *Machine Learning* (機器學習), taught by Hsuan-Tien Lin (林軒田) at National Taiwan University in the fall of 2026. This preface reproduces the course-introduction handout of 09/09/2026, so that the ground rules sit next to the material they govern; everything after it is the course itself.

Five Reasons for NOT Taking the Course

The course opens, by tradition, with an argument against taking it.

Complicated contents. From a Taiwanese student taking the MIT ML class (translated):

The professor started writing math equations as if he was using some writing accelerator. After class I always felt feeble. The worst part is: I needed to understand the contents as soon as I can. Otherwise I cannot finish the homework and cannot follow up in the next class.

The NTU ML class is designed to be as good as the best classes in the world, so similar things will happen to you. *If you are not willing to be so miserable, ...*

Strict instructor. Will you give me a second chance if I copy homework from other people? No. Could you let me pass because I will be kicked out by the 1/2 rule? No. Will you change my score from F to C? No. How many will pass? Any, if necessary. *If you do not like a strict instructor, ...*

Huge loads. From a student taking ML 2010, posted on BBS (translated):

Only 1 problem requires running 100 test experiments? ($100 \times 10 \text{ min} = 16 \text{ hr}$); I double-checked multiple times — the one who designed this homework is heartless.

This class is two to six times harder than a normal one at NTU: six homework sets are planned, plus the final project, each due within two weeks in general. Homework 0 arrives in the first week — already hard, worth very few points, and there to remind you of the prerequisites. *If you do not want to spend so much time on homework, ...*

Online-learnable. Two of the course's incarnations were invited by NTU as Massive Online Open Courses on NTU-Coursera, *Machine Learning Foundations* and *Machine Learning Techniques*, which cover 70% of what we teach; it is much easier to simply take the recorded version at home. Indeed we will do pseudo-flipped teaching. *If you want to learn at your own pace, ...*

Shallow and old-fashioned. This is not a deep learning class — it surrounds shallow (old) machine learning. Yes, many of the materials are ten or more years old; and yes, AI has knowledge of those materials, just as AI can write code. *If you want to chase newer buzzwords, ...*

From a student in ML 2013, in the final feedback (translated): *“this is a class that makes you experience all kinds of devastation and desperation.”*

May the Brave Ones Stay.

Basic Information

- **Mode:** physical, with a live YouTube screencast and Slido questions (#HTML2026FALL), if everything goes smoothly.
- **Instructor:** Hsuan-Tien Lin, htlin@csie.ntu.edu.tw. Office hour after class, or by appointment.
- **Public course infopage:** <https://www.csie.ntu.edu.tw/~htlin/course/ml26fall/>.
- **Semi-private course platform:** NTU COOL, for announcements, recorded videos and the like — be sure to check emails from NTU COOL frequently.
- **Class chats:** Discord (invitation URL on NTU COOL).
- **Homework grading** (registered students only): Gradescope, linked from NTU COOL.
- **Languages:** teaching in Mandarin; Slido questions in English, traditional or simplified Chinese; homework written in English or Chinese; TA and instructor hours in Mandarin, with English welcome. Don't be afraid of English.

The Course and Its History

- 2008–2012: 3 credits, taught in English.
- 2013–2014: 3 credits, Mandarin, with MOOC.
- 2015: 4 credits, Mandarin, with MOOC.
- 2016–2020: mostly 2 + 2 credits (MLF + MLT), Mandarin, with MOOC.
- 2021 Fall and 2023 Spring: 3 credits, English.
- 2023 Fall: 3 credits, Mandarin, with MOOC offline and physical discussions stimulated by ChatGPT.
- 2024: 3 credits, English, TAICA.
- 2026: 3 credits, Mandarin, TAICA, with MOOC offline and discussions stimulated by ChatGPT.

The reason for moving the lectures offline is that many students prefer watching a recording on their own time, which frees the precious class time to encourage speculation rather than lecturing. Three things are new this semester.

Cross-university teaching (2nd time). Taiwan needs more AI talent, so the course accepted an invitation from TAICA (Taiwan AI College Alliance) to be offered to 50+ universities — with thanks to the MOE for more TAs, a YouTube screencast to break the time and space boundary, and some online TA hours to be planned. An attempt to further increase the impact in this new AI age.

Final exam. AI can solve all of your homework with a high success rate. That is not a totally bad thing for teaching, but it makes fair grading exponentially difficult, so we will experiment with a closed-book final exam bearing some similarity to the homework — an old-fashioned yet fair grading protocol, until something better is designed.

AI-assisted TAs. AI can answer many questions with 90%+ correctness in a friendly manner; the question is how to best help learning. We are building an AI agent to answer students' questions, guarded by the TAs' regular human checking — human-computer collaboration towards better teaching and learning.

On enrollment: per school policy, students who have taken another version of the course (ML Foundations, ML Techniques, or the 4-credit ML) are generally *not* allowed to take it again. We will not check, but the school may, and multiple takings may not look nice on a transcript anyway. NTU students sign up through the link on the course infopage, processed progressively; a limited quota is reserved for TAICA students from other universities. Auditing is welcome in general, as long as it does not interfere with official members. *Leave as soon as possible — give your classmates a chance to be miserable.*

The course is supported by sixteen teaching assistants, reachable at html_ta@csie.ntu.edu.tw.

- | | |
|-----------------------------|-------------------------|
| • Cheng-Han Hsieh (*), CSIE | • Yen-Hung Hsueh, EE |
| • Guan-Chen Chen, CSIE | • Li-Cheng Hsu, EE |
| • Chun-Yang Cheng, CSIE | • Jung-Sheng Hung, GINM |
| • Sheng-Yuan Chang, CSIE | • Yun-Shi Lee, IM |
| • Tsung-Lun Chiang, EE | • Yi-Ting Lo, CSIE |
| • Ting-Chieh Chiu, CSIE | • Han-Hwa Shih, GICE |
| • Ting-Wei Chuang, CSIE | • Yen-Chu Tsai, CSIE |
| • Kang-Yu Fan, CSIE | • Yi-Ming Yao, Math |

THE Book

Learning from Data: A Short Course

Y. Abu-Mostafa (Caltech), M. Magdon-Ismael (RPI), H.-T. Lin (NTU)

The idea was initiated in 2008. Its five chapters are closely needed for the first part of the class, and further e-Chapters are used in the second part; the teaching follows the book, with suggested reading inside it. Copies can be found at the NTU Library (a reserved copy in the shared course material area) and in R536 (shared copies to be read in the room, if open); the Chuan-Hwa Book Company has imported some copies, and Amazon/Kindle is the main selling channel in the US. If the book is not affordable to you but you really want to read it, email the instructor and he will see how he can help.

THE Principle

Taking any unfair advantage over other class members is not allowed. It is everyone's responsibility to maximize the level of fairness.

NO CHEATING
NO LYING
NO PLAGIARISM
NO PIRATING of THE BOOK

There are very serious consequences. In practice this means: homework discussions are encouraged, but write the final solutions independently and understand them fully; references — books, notes, the Internet, AI tools — can be consulted but not copied from, and you are strongly encouraged to save your AI prompts and interaction logs. There is no need to lend, borrow, sell, buy or trade a solution, a piece of code or a dialogue, and doing so is not allowed.

Grading

There is no midterm. The main reference is

homework sets	30%
final exam	30%
final project	40%

after which the raw score goes through an order-preserving normalization: a raw score of 80 with term rank A is possible, and so is a raw score of 60 with term rank F; raw scores 80 and 60 both mapping to B is possible but unlikely; raw scores 80 and 60 mapping to F and B respectively is impossible. For non-NTU students the rank will tentatively be converted by the NTU metric back to a score in $[0, 100]$, which the student's own university then converts back to a rank.

Homework uses multiple-choice problems for auto-grading, with photographed or scanned details and code uploaded for selective grading by the TAs — there is a penalty if your details are not logical — and for plagiarism detection. Late parts are worth 90% if twelve hours late, 80% if one day late, and so on; each person is granted four penalty-free late half-days, the gold medals. About a quarter to a third of the problems are programming assignments, in any language, on any platform, with any packages; AI-generated code is fine if you co-wrote it with the AI and your logs justify that.

The *final project* is a teamed in-class competition, with team size to be decided, co-run with Prof. Shou-de Lin's *Machine Learning: Theory and Practice* class — so the two classes cannot both be taken this semester. You will need to sign and abide by a data policy. It can make you rich, if you know what I mean.

Contents

1. The Learning Problem	7
1.1. Course Introduction	7
1.2. What is Machine Learning	8
1.3. Applications of Machine Learning	10
1.4. Components of Machine Learning	13
1.5. Machine Learning and Other Fields	15
2. Learning to Answer Yes/No	18
2.1. Perceptron Hypothesis Set	18
2.2. Perceptron Learning Algorithm (PLA)	21
2.3. Guarantee of PLA	24
2.4. Non-Separable Data	28
3. Types of Learning	31
3.1. Learning with Different Output Space $\mathcal{Y}$	31
3.2. Learning with Different Data Label y_n	34
3.3. Learning with Different Protocol $f \Rightarrow (x_n, y_n)$	37
3.4. Learning with Different Input Space $\mathcal{X}$	39
4. Feasibility of Learning	43
4.1. Learning is Impossible?	43
4.2. Probability to the Rescue	46
4.3. Connection to Learning	48
4.4. Connection to Real Learning	51
5. Training versus Testing	56
5.1. Recap and Preview	56
5.2. Effective Number of Lines	58
5.3. Effective Number of Hypotheses	62
5.4. Break Point	66
6. Theory of Generalization	69
6.1. Restriction of Break Point	69
6.2. Bounding Function: Basic Cases	71
6.3. Bounding Function: Inductive Cases	73
6.4. A Pictorial Proof	77
7. The VC Dimension	82
7.1. Definition of VC Dimension	82
7.2. VC Dimension of Perceptrons	85
7.3. Physical Intuition of VC Dimension	89
7.4. Interpreting VC Dimension	91
8. Noise and Error	95
8.1. Noise and Probabilistic Target	95
8.2. Error Measure	98
8.3. Algorithmic Error Measure	101
8.4. Weighted Classification	104

1. The Learning Problem

Lecture 1

§1.1. Course Introduction

9 Sep.

— slides p.1–5 —

Machine learning is a mixture of theoretical and practical tools, so any course about it has to decide how to mix them. There are two natural extremes, and both of them fail in their own way.

- *Theory oriented*: derive everything deeply, so that every statement rests on a proof and the understanding is solid. The cost is that it is less interesting to a general audience — one may finish such a course able to prove a generalization bound, yet unable to fit a model to a real data set.
- *Techniques oriented*: flash over the sexiest techniques broadly, so that the coverage looks shiny. The cost is that there are simply too many techniques: they become hard to choose between, and hard to use properly, because nothing tells you *why* one of them should work here and not there.

Our approach: foundation oriented.

By *foundation* we mean the part that does not change when the fashionable technique changes. Concretely, the course is a mixture of philosophical illustrations, key theory, core techniques, usage in practice, and hopefully jokes; the selecting principle is always *what every machine learning user should know*. The payoff of insisting on foundations is that they transfer: a student who owns them can pick up a future or untaught technique, or go read deeper theory, without being taught each one explicitly.

The material is organized **story-like**, as a sequence of questions, each one asked only after the previous one has been answered:

1. **When** can machines learn? (illustrative + technical)
2. **Why** can machines learn? (theoretical + illustrative)
3. **How** can machines learn? (technical + practical)
4. **How** can machines learn **better**? (practical + theoretical)

Note 1.1.1. The four questions above are the backbone of the *foundations* half of the story. The NTU version continues it with three more chapters, each one a different answer to “where do good features come from”:

5. How can machines learn by *embedding numerous* features? (support vector machines)
6. How can machines learn by *combining predictive* features? (aggregation, boosting, trees and forests)
7. How can machines learn by *distilling hidden* features? (neural networks and deep learning)

This first chapter sits entirely inside question 1, and it is the most basic piece of it: before asking when machines *can* learn, we must say what the learning problem even *is*. That is the plan of the chapter: we first ask what machine learning is (Section 1.2), then look at where it already lives (Section 1.3), then cut it into its formal components (Section 1.4), and finally place it next to the neighbouring fields it is often confused with (Section 1.5).

Exercise 1.1.2 (Fun Time). Which of the following descriptions of this course is true?

- (1) the course will be taught in Taiwanese;
- (2) the course will tell me the techniques that create the android Lieutenant Commander Data in *Star Trek*;
- (3) the course will be 15 weeks long;
- (4) the course will be story-like.

Reference Answer. (4). As for the others: (1) no, the instructor's Taiwanese is unfortunately not good enough for teaching (yet); (2) no, although what we teach may serve as foundations of those (future) techniques; (3) no, unless you choose to join the next course. So: let us begin the story. ■


§1.2. What is Machine Learning

— slides p.6–11 —

§1.2.i. From Learning to Machine Learning


The safest way into a definition is by analogy with the thing we already do ourselves.

Definition 1.2.1 (Learning). **Learning** is acquiring *skill* with experience accumulated from *observations*.



A machine has no eyes and no childhood, so the only observations it can accumulate are the ones we hand it in recorded form. Replacing *observations* by *data* and *learning* by *computation* gives the machine version.

Definition 1.2.2 (Machine learning, first attempt). **Machine learning** is acquiring skill with experience accumulated/computed from *data*.



The word *data* is concrete enough to be useful. The word *skill* is not: it is exactly the kind of word that cannot be handed to a computer, so the definition is not finished yet.

Question 1.2.3. What is skill?

§1.2.ii. A More Concrete Definition

Intuition . skill $\Leftrightarrow$ improve some *performance measure* (e.g. prediction accuracy).

This is the move that makes the whole subject possible: a vague goal (“be good at recognizing trees”) is replaced by a number that can be computed and watched (“fraction of tree images classified correctly”). Every later chapter is, in one way or another, an argument about which number to watch and what watching it really guarantees.

Definition 1.2.4 (Machine learning). **Machine learning** is improving some *performance measure* with experience computed from *data*.



Example 1.2.5 (An application in computational finance). Take the data to be historical stock data, and the performance measure to be investment gain. Then machine learning here means: use the stock data to compute a trading rule whose gain is higher than what we had before.



§1.2.iii. Why Use Machine Learning?

We now know what machine learning *is*. A separate matter is why one would ever want it, given that a computer can already be programmed directly.

Question 1.2.6. Why use machine learning?

Example 1.2.7 (Tree recognition). Suppose we want a system that looks at a photograph and answers whether there is a tree in it.

- To ‘*define*’ *trees and hand-program* the answer is difficult: try writing down, in code, a condition that accepts every pine, palm and bonsai and rejects every bush and lamppost.
- To *learn from data (observations) and recognize* is easy enough that a 3-year-old can do it, having never received a definition of “tree” in their life — only examples.

So an ‘ML-based tree recognition system’ can be easier to *build* than a hand-programmed one, even though it is harder to *explain*.

ML is an alternative route to build complicated systems.

The useful way to read that moral is as a list of situations where the direct route is blocked:

- when humans cannot program the system manually — navigating on Mars;
- when humans cannot ‘define the solution’ easily — speech/visual recognition;
- when rapid decisions are needed that humans cannot make — high-frequency trading;
- when the system must be user-oriented at a massive scale — consumer-targeted marketing.

Remark 1.2.8. Give a computer a fish, you feed it for a day; teach it how to fish, you feed it for a lifetime. :-)

§1.2.iv. Key Essence of Machine Learning

Machine learning is not always the right tool, and the definition above already tells us when it is not. Reading [Definition 1.2.4](#) backwards, three things must hold before the phrase even makes sense.

1. There exists some **underlying pattern** to be learned — so that a ‘performance measure’ *can* be improved. Without a pattern there is nothing to be better at.
2. But there is no programmable (easy) *definition* of that pattern — so that ‘ML’ is *needed*. If you can just write the rule down, write it down.
3. Somehow there is **data** about the pattern — so that ML has some ‘inputs’ to learn from. A pattern nobody has recorded cannot be computed from.

Key essence: it helps decide whether to use ML at all.

Exercise 1.2.9 (Fun Time). Which of the following is best suited for machine learning?

- (1) predicting whether the next cry of a baby girl happens at an even-numbered minute or not;
- (2) determining whether a given graph contains a cycle;
- (3) deciding whether to approve a credit card to some customer;
- (4) guessing whether the earth will be destroyed by the misuse of nuclear power in the next ten years.

Reference Answer. (3), and the three essences say exactly why each of the others fails.


- (1) Fails essence 1: there is no pattern.
- (2) Fails essence 2: the definition is programmable (run a cycle detection algorithm).
- (3) Passes all three — pattern: customer behaviour; definition: not easily programmable; data: the history of bank operations.
- (4) Fails essence 3: there is arguably no (or not enough) data yet.

■

§1.3. Applications of Machine Learning

— slides p.12–16 —

Every application below is the same picture,



with the two ends filled in differently. So the exercise in reading an application is always: *what is the data, and what is the skill?*

§1.3.i. Daily Needs: Food, Clothing, Housing, Transportation

1. **Food** (Sadilek et al., 2013). *Data*: Twitter data (words + location). *Skill*: tell the food-poisoning likeliness of a restaurant properly.
2. **Clothing** (Abu-Mostafa, 2012). *Data*: sales figures + client surveys. *Skill*: give good fashion recommendations to clients.
3. **Housing** (Tsanas and Xifara, 2012). *Data*: characteristics of buildings and their energy load. *Skill*: predict the energy load of other buildings closely.
4. **Transportation** (Stallkamp et al., 2012). *Data*: some traffic sign images and their meanings. *Skill*: recognize traffic signs accurately.

ML is everywhere!

§1.3.ii. Education

Data: students' records on quizzes in a mathematics tutoring system. *Skill*: predict whether a student can give a correct answer to *another* quiz question.

What makes this example instructive is that the pattern one guesses is not a formula in the recorded variables at all; it is a formula in quantities nobody ever measured:

$$\text{answer correctly} \approx \llbracket \text{recent strength of student} > \text{difficulty of question} \rrbracket.$$

Notation . $\llbracket \cdot \rrbracket$ is the *Iverson bracket*: it equals 1 when the statement inside is true and 0 when it is false.

Neither “strength” nor “difficulty” appears in the data — only who answered what, and whether it was right. Feeding roughly 9 million records from 3000 students to a learning algorithm lets it *reverse-engineer* both numbers automatically, purely from the requirement that the formula above should reproduce the observed answers.

This was the key part of the world-champion system from National Taiwan University in KDDCup 2010.

§1.3.iii. Entertainment: Recommender Systems

Data: how many users have rated some movies. *Skill*: predict how a user would rate an unrated movie.


Example 1.3.1 (A hot problem). • A competition held by Netflix in 2006: 100,480,507 ratings that 480,189 users gave to 17,770 movies, with a 10% improvement worth a one-million-dollar prize.

- A similar competition (movies $\rightarrow$ songs) held by Yahoo! in KDDCup 2011: 252,800,275 ratings that 1,000,990 users gave to 624,961 songs.

Question 1.3.2. How can machines learn our preferences?

The same trick as in the education example applies, and it is worth seeing that it is literally the same trick. Postulate that a viewer and a movie are each described by the *same* list of hidden **factors** — how much comedy, how much action, how much of a blockbuster — read on

the viewer's side as *how much this viewer wants it* and on the movie's side as *how much this movie has it*.



Both rows carry the same factors in the same order, so a viewer and a movie live in one common space; the size of a dot is the size of a number. The vertical arrow pairs slot k with slot k , the horizontal one sums the pairs, and together they are a single formula: writing v_k for how much the viewer wants factor k and m_k for how much the movie has of it,

$$\text{predicted rating} \approx \sum_{k=1}^K v_k m_k,$$

the inner product of the two factor vectors. The product is the point — a factor contributes only where the two sides *agree*, so our action-loving viewer would rate this comedy rather low.

The dashes say the list is long, and the last slot says how arbitrary it may be. Nobody drew that list up: as with the students' strength and difficulty, the factors are never observed, and the learning is what invents them.

pattern: rating $\leftarrow$ viewer/movie factors;

learning: known ratings $\rightarrow$ learned factors $\rightarrow$ unknown rating prediction.

Again the key part of the world-champion system from National Taiwan University, this time in KDDCup 2011.

Exercise 1.3.3 (Fun Time). Which of the following fields *cannot* use machine learning?

- (1) Finance;
- (2) Medicine;
- (3) Law;

(4) none of the above.

Reference Answer. (4). Finance can predict stock prices from data, medicine can predict medicine effects from data, and law can summarize legal documents from data. Welcome to study this hot topic. ■

§1.4. Components of Machine Learning

— slides p.17–22 —

So far “machine learning” has been a box in a picture. This section opens the box and names its parts. The running metaphor is *credit approval*: a bank receives an application and must decide whether issuing a card to this person is good for the bank.

Applicant Information	
age	23 years
gender	female
annual salary	NTD 1,000,000
year in residence	1 year
year in job	0.5 year
current debt	200,000

The unknown pattern to be learned is: *is approving this credit card good for the bank?*

§1.4.i. Formalizing the Learning Problem

Notation .

- **Input:** $x \in \mathcal{X}$ (a customer application).
- **Output:** $y \in \mathcal{Y}$ (good/bad after approving the credit card).
- The unknown pattern to be learned $\Leftrightarrow$ the **target function**

$$f: \mathcal{X} \rightarrow \mathcal{Y}$$

(the ideal credit approval formula).

- The data $\Leftrightarrow$ the **training examples**


$$\mathcal{D} = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

(historical records in the bank).

- The skill, hopefully with good performance, $\Leftrightarrow$ the **hypothesis**

$$g: \mathcal{X} \rightarrow \mathcal{Y}$$

(the ‘learned’ formula to be used).




Remark 1.4.1. The asymmetry between f and g is the whole difficulty of the subject, so it is worth stating plainly.

- f is *unknown* — indeed, by essence 2 of the previous section, it has no programmable definition. We never see f ; we only ever see N of its values.
- g is what we produce, and we hope $g \approx f$. But g may differ from f , and perfection is ‘impossible’ precisely because f is unknown — there is no way to check the answer.

Everything in *Why can machines learn?* is an attempt to say something rigorous about $g \approx f$ despite never seeing f .

Drawing the arrows in the order in which things actually happen gives the learning flow.



§1.4.ii. The Learning Model

The diagram still hides one thing.

Question 1.4.2. What does g look like?

An algorithm cannot search over *all* functions $\mathcal{X} \rightarrow \mathcal{Y}$, so we must first decide which formulas are even candidates.


Definition 1.4.3 (Hypothesis set). The **hypothesis set** $\mathcal{H} = \{h_k\}$ is the set of candidate formulas that the algorithm is allowed to return, so that $g \in \mathcal{H}$.

Example 1.4.4. For credit approval, one might approve if

- h_1 : annual salary $>$ NTD 800,000;
- h_2 : debt $>$ NTD 100,000 (really?);
- h_3 : year in job ≤ 2 (really?).

Note that $\mathcal{H}$ happily contains bad hypotheses next to good ones. That is not a defect: it is up to $\mathcal{A}$ to pick the ‘best’ one out of $\mathcal{H}$ and call it g .

Definition 1.4.5 (Learning model). A **learning model** is the pair consisting of the hypothesis set $\mathcal{H}$ and the learning algorithm $\mathcal{A}$.



This is the picture the rest of the course elaborates. The dashed halves of the story are already visible in it: f and $\mathcal{D}$ are given by the world and we cannot choose them, while $\mathcal{H}$ and $\mathcal{A}$ are ours to design.

Machine learning: use data to compute a hypothesis g that approximates the target f .

Exercise 1.4.6 (Fun Time). How should the four sets below be used to form a learning problem for song recommendation?

- $S_1 = [0, 100]$;
- $S_2 =$ all possible (userid, songid) pairs;
- $S_3 =$ all formulas that ‘multiply’ user factors and song factors,
indexed by all possible combinations of such factors;
- $S_4 =$ 1,000,000 pairs of ((userid, songid), rating).

- (1) $S_1 = \mathcal{X}$, $S_2 = \mathcal{Y}$, $S_3 = \mathcal{H}$, $S_4 = \mathcal{D}$;
- (2) $S_1 = \mathcal{Y}$, $S_2 = \mathcal{X}$, $S_3 = \mathcal{H}$, $S_4 = \mathcal{D}$;
- (3) $S_1 = \mathcal{D}$, $S_2 = \mathcal{H}$, $S_3 = \mathcal{Y}$, $S_4 = \mathcal{X}$;
- (4) $S_1 = \mathcal{X}$, $S_2 = \mathcal{D}$, $S_3 = \mathcal{Y}$, $S_4 = \mathcal{H}$.

Reference Answer. (2). Reading the whole problem in one line,

$$S_4 \xrightarrow{\mathcal{A} \text{ on } S_3} (g: S_2 \rightarrow S_1),$$

i.e. the input space is the set of (user, song) pairs, the output space is the rating scale $[0, 100]$, the hypothesis set is the family of factor-multiplying formulas, and the data is the list of observed ratings. ■

§1.5. Machine Learning and Other Fields

— slides p.23–27 —

Machine learning did not grow up alone, and much of the confusion around the name comes from its neighbours. Comparing the one-line definitions is enough to see both the overlap and the difference in *focus* — which is usually where the real difference lies.

§1.5.i. Machine Learning and Data Mining

Machine Learning	Data Mining
use data to compute a hypothesis g that approximates a target f	use (huge) data to find a property that is <i>interesting</i>

- If the ‘interesting property’ *is* the ‘hypothesis that approximates the target’, then ML = DM. This is usually what KDDCup does.
- If the ‘interesting property’ is merely *related* to such a hypothesis, then DM can help ML and vice versa — often, but not always.
- Traditional DM also puts weight on efficient computation in large databases, which ML traditionally does not.

It is difficult to distinguish ML and DM in reality.

§1.5.ii. Machine Learning and Artificial Intelligence

Machine Learning	Artificial Intelligence
use data to compute a hypothesis g that approximates a target f	compute something that shows <i>intelligent behaviour</i>

- $g \approx f$ is something that shows intelligent behaviour, so ML can realize AI — among other routes.
- For example, in chess playing: traditional AI builds a game tree, whereas ML for AI learns from board data.

ML is one possible route to realize AI.

§1.5.iii. Machine Learning and Statistics

Machine Learning	Statistics
use data to compute a hypothesis g that approximates a target f	use data to make <i>inference</i> about an unknown process

- g is an inference outcome and f is something unknown, so statistics can be used to achieve ML.
- Traditional statistics also focuses on provable results under mathematical assumptions, and cares less about computation.

Statistics: many useful tools for ML.

Exercise 1.5.1 (Fun Time). Which of the following claims is *not* totally true?

- (1) machine learning is a route to realize artificial intelligence;
- (2) machine learning, data mining and statistics all need data;
- (3) data mining is just another name for machine learning;

(4) statistics can be used for data mining.

Reference Answer. (3). While data mining and machine learning do share a huge overlap, they are arguably not equivalent, because of the difference of focus. ■

Summary

Course Introduction foundation oriented and story-like.

What is Machine Learning use data to approximate a target.

Applications of Machine Learning almost everywhere.

Components of Machine Learning $\mathcal{A}$ takes $\mathcal{D}$ and $\mathcal{H}$ to get g .

Machine Learning and Other Fields related to DM, AI and statistics.

We have a frame, $(\mathcal{X}, \mathcal{Y}, f, \mathcal{D}, \mathcal{H}, \mathcal{A}, g)$, but not a single concrete learning model to put in it: we never said what any $h \in \mathcal{H}$ actually computes, nor how $\mathcal{A}$ searches. The next chapter fills in the simplest useful choice of $\mathcal{H}$ and $\mathcal{A}$, for the simplest useful output space $\mathcal{Y} = \{-1, +1\}$ — that is, learning to answer yes/no.

2. Learning to Answer Yes/No

Lecture 2

As previously seen, [Section 1.4](#) ended with a complete frame for the learning problem: an algorithm $\mathcal{A}$ takes the data $\mathcal{D}$ and a hypothesis set $\mathcal{H}$, and returns a final hypothesis g which we hope satisfies $g \approx f$. Nothing in that frame said what a single $h \in \mathcal{H}$ computes, nor how $\mathcal{A}$ searches.

16 Sep.


This chapter fills that gap in the simplest interesting case, where the output is a yes/no answer, and gets both an $\mathcal{H}$ and an $\mathcal{A}$ out of it — our first concrete learning model.

§2.1. Perceptron Hypothesis Set

— slides p.2–6 —

§2.1.i. The Credit Approval Problem Revisited

Put the credit approval problem of [Section 1.4](#) back into the flow, this time with the applicant record written out.



Question 2.1.1. What hypothesis set can we use?

§2.1.ii. Features of a Customer

The bank holds a customer record of d numbers,

$$\mathbf{x} = (x_1, x_2, \dots, x_d),$$

the **features** of the customer — age, annual salary, year in job, current debt, and so on — and it must answer yes or no.

Notation . From this chapter on, an input $\mathbf{x}$ is a vector of d real features, written in bold. In [Section 1.4](#) the input was an abstract element of $\mathcal{X}$; giving it coordinates is exactly the step that lets us write down formulas for h .

§2.1.iii. The Perceptron

age	23 years
annual salary	NTD 1,000,000
year in job	0.5 year
current debt	200,000

The simplest thing one can do with d numbers is to weigh them against each other. Give each feature x_i an importance w_i , form a weighted *score* $\sum_{i=1}^d w_i x_i$, and compare it against a threshold:

$$\begin{aligned} \text{approve credit if } & \sum_{i=1}^d w_i x_i > \text{threshold;} \\ \text{deny credit if } & \sum_{i=1}^d w_i x_i < \text{threshold.} \end{aligned}$$

A positive w_i means the feature speaks in the customer's favour (annual salary), a negative one means it speaks against (current debt), and a large $|w_i|$ means the feature matters a lot.

Since the answer is yes or no, take

$$\mathcal{Y} = \{+1 \text{ (good), } -1 \text{ (bad)}\},$$

and ignore the boundary case of a score exactly equal to the threshold, i.e. treat 0 as ignored. Then the rule above is a single formula.

Definition 2.1.2 (Perceptron). For a weight vector $w_1, \dots, w_d$ and a threshold, the hypothesis

$$h(\mathbf{x}) = \text{sign} \left(\left(\sum_{i=1}^d w_i x_i \right) - \text{threshold} \right)$$

is called a **perceptron**, for historical reasons. The hypothesis set $\mathcal{H}$ of this chapter is the set of all such h , i.e. one h for every choice of weights and threshold.

§2.1.iv. Vector Form

The threshold is an eyesore: it is a parameter of the same kind as the weights, but it sits outside the sum. The standard repair is to absorb it as the weight of a constant feature:

$$\begin{aligned} h(\mathbf{x}) &= \text{sign} \left(\left(\sum_{i=1}^d w_i x_i \right) - \text{threshold} \right) \\ &= \text{sign} \left(\left(\sum_{i=1}^d w_i x_i \right) + \underbrace{(-\text{threshold})}_{w_0} \cdot \underbrace{(+1)}_{x_0} \right) \\ &= \text{sign} \left(\sum_{i=0}^d w_i x_i \right) \\ &= \text{sign} \left(\mathbf{w}^\top \mathbf{x} \right). \end{aligned}$$

Notation . From now on $\mathbf{w} = (w_0, w_1, \dots, w_d)^\top$ and $\mathbf{x} = (x_0, x_1, \dots, x_d)^\top$ are the ‘tall’ versions, with $x_0 \equiv 1$ always. Each tall $\mathbf{w}$ is a hypothesis h , so from here on we say ‘the hypothesis $\mathbf{w}$ ’ and mean $h(\mathbf{x}) = \text{sign}(\mathbf{w}^\top \mathbf{x})$.

The gain is not only cosmetic: everything below manipulates $\mathbf{w}$ as a vector, and a stray additive constant would have to be carried through every one of those manipulations.

§2.1.v. Perceptrons in $\mathbb{R}^2$


Question 2.1.3. What do perceptrons h ‘look like’?

Take $d = 2$, so that

$$h(\mathbf{x}) = \text{sign}(w_0 + w_1x_1 + w_2x_2).$$

The set where the score vanishes, $w_0 + w_1x_1 + w_2x_2 = 0$, is a line; h says $+1$ on one side of it and -1 on the other. So:

- customer features $\mathbf{x}$ are points on the plane (or points in $\mathbb{R}^d$);
- labels y are drawn as $\circ$ for $+1$ and $\times$ for -1 ;
- a hypothesis h is a line (or a hyperplane in $\mathbb{R}^d$), positive on one side and negative on the other;
- a different line classifies the customers differently.



perceptrons $\Leftrightarrow$ linear (binary) classifiers.

Exercise 2.1.4 (Fun Time). Consider using a perceptron to detect spam messages. Assume that each email is represented by the frequency of keyword occurrence, and that an output of $+1$ indicates spam. Which keywords below should have large *positive* weights in a good perceptron for the task?

- (1) coffee, tea, hamburger, steak;
- (2) free, drug, fantastic, deal;
- (3) machine, learning, statistics, textbook;
- (4) national, Taiwan, university, coursera.

Reference Answer. (2). The occurrence of a keyword with a positive weight increases the ‘spam score’, and hence those keywords should be the ones that often appear in spam. ■

§2.2. Perceptron Learning Algorithm (PLA)


— slides p.7–12 —

Half of the learning model is still missing.

Question 2.2.1. $\mathcal{H}$ = all possible perceptrons, $g = ?$

§2.2.i. Selecting g from $\mathcal{H}$

- What we *want* is $g \approx f$. This is hard to even check, because f is unknown.
- What is *almost necessary* is $g \approx f$ on $\mathcal{D}$, ideally $g(\mathbf{x}_n) = f(\mathbf{x}_n) = y_n$ for every n . This we can check, because $\mathcal{D}$ is in front of us.
- What makes it *difficult* is that $\mathcal{H}$ is of infinite size, so we cannot look at the candidates one by one.
- The *idea* that resolves this: start from some g_0 , and ‘correct’ its mistakes on $\mathcal{D}$.



Infinitely many candidates are no obstacle if we never enumerate them, but instead walk from one to the next. By the previous subsection we may represent g_0 by its weight vector $\mathbf{w}_0$, so the walk happens in $\mathbb{R}^{d+1}$ rather than in the abstract set $\mathcal{H}$.

§2.2.ii. The Algorithm

Algorithm 2.1: Perceptron Learning Algorithm (PLA)

Input: training examples $\mathcal{D} = \{(\mathbf{x}_n, y_n)\}_{n=1}^N$


Output: $\mathbf{w}_{\text{PLA}}$, returned as g

```

1  $\mathbf{w}_0 \leftarrow$  some initial vector, say  $\mathbf{0}$ ;
2 for  $t = 0, 1, 2, \dots$  do
3   if  $\mathbf{w}_t$  makes no mistake on  $\mathcal{D}$  then return  $\mathbf{w}_t$  as  $\mathbf{w}_{\text{PLA}}$ ;
4   find a mistake of  $\mathbf{w}_t$ , called  $(\mathbf{x}_{n(t)}, y_{n(t)})$ , i.e. one with  $\text{sign}(\mathbf{w}_t^\top \mathbf{x}_{n(t)}) \neq y_{n(t)}$ ;
5   (try to) correct the mistake by  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + y_{n(t)} \mathbf{x}_{n(t)}$ ;
```

That is it — a fault confessed is half redressed. :-)

The update rule looks arbitrary until one draws it. A mistake at $(\mathbf{x}, y)$ with $y = +1$ means $\mathbf{w}^\top \mathbf{x} < 0$, i.e. $\mathbf{w}$ and $\mathbf{x}$ point more than 90° apart; adding $\mathbf{x}$ swings $\mathbf{w}$ *towards* $\mathbf{x}$. A mistake with $y = -1$ means $\mathbf{w}^\top \mathbf{x} > 0$; subtracting $\mathbf{x}$ swings $\mathbf{w}$ *away* from $\mathbf{x}$. Either way the correction is $\mathbf{w} + y\mathbf{x}$.



§2.2.iii. Practical Implementation of PLA

algorithm 2.1 says “find a mistake” without saying *which* one, and a literal implementation would rescan $\mathcal{D}$ from the beginning after every correction. Cyclic PLA removes both vaguenesses at once: fix an order on the examples, keep a cursor into it, and let “the next mistake” mean *the next mistake the cursor runs into*, continuing from where the last correction happened and wrapping around at the end of the list.

Written out, with $\pi(1), \dots, \pi(N)$ the chosen order:

Algorithm 2.2: Cyclic PLA

Input: training examples $\mathcal{D} = \{(x_n, y_n)\}_{n=1}^N$ and a fixed order $\pi(1), \dots, \pi(N)$ of $1, \dots, N$

Output: w_{PLA} , returned as g

```

1  $w \leftarrow 0$ ;
2  $i \leftarrow 1$ ;                                // cursor: where the next scan resumes
3  $c \leftarrow 0$ ;                                // examples seen in a row with no mistake
4 while  $c < N$  do
5    $n \leftarrow \pi(i)$ ;
6   if  $\text{sign}(w^\top x_n) \neq y_n$  then
7      $w \leftarrow w + y_n x_n$ ;                  // a mistake: correct it
8      $c \leftarrow 0$ ;                            // the clean run is broken
9   else
10     $c \leftarrow c + 1$ ;                          // this one is already right
11     $i \leftarrow (i \bmod N) + 1$ ;                  // advance, wrapping around
12 return  $w$  as  $w_{\text{PLA}}$ ;
```

Two points deserve spelling out.


- **When does the loop end?** When c reaches N , i.e. when N examples in a row have been checked and none of them was a mistake. Since the cursor moves cyclically, any N consecutive visits run through every example exactly once — so $c = N$ says precisely that a full cycle passed without encountering a mistake, which is to say that the current w is correct on all of $\mathcal{D}$. There is no need for the scan to be aligned with the start of the list.
- **Why reset c to 0 after every correction?** Because a correction changes w , and the new w may well be wrong on examples the cursor has already walked past — exactly what happens to x_{15} in the run below. Everything checked before the correction was checked against a *different* hypothesis and no longer counts, so the tally has to start over.

The order π may be the naïve cycle $1, \dots, N$, or a precomputed random cycle. Note also that the loop counter here is not the t of algorithm 2.1: this loop turns once per example *examined*, whereas t — and the T bounded in Theorem 2.3.6 — counts only the examples *actually corrected*.

§2.2.iv. Seeing is Believing

Run it on a small separable data set of $N = 20$ examples. Each frame shows the hypothesis $\mathbf{w}_t$ that PLA currently holds — its line, and the two shaded sides it calls +1 (blue) and -1 (red) — together with the example $\mathbf{x}_{n(t)}$ it still gets wrong, circled. A mistake is visible as a $\circ$ sitting in the red half or a $\times$ sitting in the blue half.

Every frame is drawn with the origin O at its centre, so that all three vectors can be read off from there: the circled $\mathbf{x}_{n(t)}$ (black), the current $\mathbf{w}_t$ (solid, the one whose line is shaded) and the $\mathbf{w}_{t+1}$ the correction is about to produce (dashed) — which is the line drawn in the next frame. All three are at 0.6 of their true length so that they stay beside the data; the dotted continuation of $\mathbf{x}_{n(t)}$ runs on to the example itself.



Because the three share one scale, the dashed segment joining the two weight tips is the step $y_{n(t)}\mathbf{x}_{n(t)}$ — parallel to $\mathbf{x}_{n(t)}$ and of the same drawn length — so the update rule is literally the triangle that closes:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + y_{n(t)}\mathbf{x}_{n(t)}.$$

The run is started from an initial $\mathbf{w}_0$ that is merely a guess — taking $\mathbf{w}_0 = \mathbf{0}$, as [algorithm 2.1](#) suggests, would leave the first frame with no line to look at — so already the first frame shows a genuine classifier, and a genuinely bad one. Each $\mathbf{w}_t$ is perpendicular to its own line and points into the +1 side, so watch it swing towards a circled $\circ$ and away from a circled $\times$. A correction made for one example happily breaks another: here $\mathbf{x}_{20}$ has to be corrected twice, in updates 1 and 3, and so does $\mathbf{x}_4$, in updates 2 and 4. After five corrections no mistake is left and the algorithm stops.

It works like a charm in fewer than 20 lines of code. But watching it work once is not a guarantee, and two kinds of issue are left open.

Question 2.2.2 (Some remaining issues of PLA). **Algorithmic:** does it *halt*, with no mistake?

- naïve cyclic: ??
- random cyclic: ??
- other variant: ??

Learning: is $g \approx f$?

- on $\mathcal{D}$, if it halts: yes, since then there is no mistake;
- outside $\mathcal{D}$: ??
- if not halting: ??

To be shown: if $(\dots)$, then after ‘enough’ corrections any PLA variant halts.

The rest of the chapter settles the first question and postpones the “outside $\mathcal{D}$ ” half of the second one to *Why can machines learn?*

Exercise 2.2.3 (Fun Time). Let us try to think about why PLA may work. Let $n = n(t)$. According to the rule

$$\text{sign}(\mathbf{w}_t^\top \mathbf{x}_n) \neq y_n, \quad \mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + y_n \mathbf{x}_n,$$

which formula is true?

- (1) $\mathbf{w}_{t+1}^\top \mathbf{x}_n = y_n$;
- (2) $\text{sign}(\mathbf{w}_{t+1}^\top \mathbf{x}_n) = y_n$;
- (3) $y_n \mathbf{w}_{t+1}^\top \mathbf{x}_n \geq y_n \mathbf{w}_t^\top \mathbf{x}_n$;
- (4) $y_n \mathbf{w}_{t+1}^\top \mathbf{x}_n < y_n \mathbf{w}_t^\top \mathbf{x}_n$.

Reference Answer. (3). Simply multiply the update rule by $y_n \mathbf{x}_n$:

$$y_n \mathbf{w}_{t+1}^\top \mathbf{x}_n = y_n \mathbf{w}_t^\top \mathbf{x}_n + y_n^2 \|\mathbf{x}_n\|^2 \geq y_n \mathbf{w}_t^\top \mathbf{x}_n,$$

since $y_n^2 = 1$. Note that (1) and (2) are *too strong*: one correction need not fix the very mistake it was made for. All we get is that the quantity $y_n \mathbf{w}^\top \mathbf{x}_n$, which is positive exactly when $\mathbf{x}_n$ is classified correctly, moves in the right direction. The rule somewhat ‘tries to correct the mistake’. ■


§2.3. Guarantee of PLA

— slides p.13–16 —

§2.3.i. Linear Separability

If PLA halts, then by construction there are no mistakes left, so $\mathcal{D}$ must have admitted a perfect $\mathbf{w}$ in the first place. That is a *necessary* condition for halting, and it deserves a name.

Definition 2.3.1 (Linear separability). A data set $\mathcal{D}$ is **linearly separable** if there exists a $\mathbf{w}$ that makes no mistake on it, i.e. $y_n = \text{sign}(\mathbf{w}^\top \mathbf{x}_n)$ for every n .



The converse is the real question.

Question 2.3.2. Assume a linearly separable $\mathcal{D}$. Does PLA always halt?

It does, and the proof is two facts about what a single correction does.

§2.3.ii. Two Facts

Fact 1: w_t gets more aligned with w_f

Restating Definition 2.3.1 gives the premise the whole argument runs on:

$$\mathcal{D} \text{ linearly separable} \Leftrightarrow \text{exists perfect } w_f \text{ such that } y_n = \text{sign}(w_f^\top x_n).$$

Fix such a w_f . Being **perfect** buys more than “no mistakes”: every example is not merely on the correct side of w_f ’s line but strictly *away* from it, so the worst example over the whole (finite) data set still has room to spare:

$$y_{n(t)} w_f^\top x_{n(t)} \geq \min_n y_n w_f^\top x_n > 0.$$

That strictly positive number is the entire engine of the proof, because it is paid out afresh at every single correction.

Lemma 2.3.3 ($w_f^\top w_t$ goes up)

Updating with any mistake $(x_{n(t)}, y_{n(t)})$ gives

$$w_f^\top w_{t+1} \geq w_f^\top w_t + \min_n y_n w_f^\top x_n > w_f^\top w_t.$$

Proof. Expand the update and apply the displayed bound above:

$$\begin{aligned} w_f^\top w_{t+1} &= w_f^\top (w_t + y_{n(t)} x_{n(t)}) \\ &\geq w_f^\top w_t + \min_n y_n w_f^\top x_n \\ &> w_f^\top w_t + 0. \end{aligned}$$

So w_t *appears* more aligned with w_f after an update — really? Not yet. An inner product can grow for two quite different reasons: because the angle between the two vectors shrank, which is what we want, or merely because w_t got longer, which would mean nothing at all. Lemma 2.3.3 on its own cannot tell these apart. The second fact closes exactly that hole. ■

Fact 2: w_t does not grow too fast

Here the premise is the other half of how PLA is built — it touches w_t *only* when there is a mistake to fix, and a mistake is precisely a negative agreement:

$$w_t \text{ changed only when mistake} \Leftrightarrow \text{sign}(w_t^\top x_{n(t)}) \neq y_{n(t)} \Leftrightarrow y_{n(t)} w_t^\top x_{n(t)} \leq 0.$$

Lemma 2.3.4 ($\|w_t\|^2$ is limited)

A mistake ‘limits’ the growth of $\|w_t\|^2$, even when the update uses the ‘longest’ x_n :

$$\|w_{t+1}\|^2 \leq \|w_t\|^2 + \max_n \|y_n x_n\|^2.$$

Proof. The cross term is the one to watch, and the characterisation of a mistake above says it is never positive:

$$\begin{aligned} \|w_{t+1}\|^2 &= \|w_t + y_{n(t)} x_{n(t)}\|^2 \\ &= \|w_t\|^2 + 2y_{n(t)} w_t^\top x_{n(t)} + \|y_{n(t)} x_{n(t)}\|^2 \\ &\leq \|w_t\|^2 + 0 + \|y_{n(t)} x_{n(t)}\|^2 \\ &\leq \|w_t\|^2 + \max_n \|y_n x_n\|^2. \end{aligned}$$

■

Intuition . Updating only on a mistake is what makes the algorithm work. A mistake means w_t and $y_{n(t)} x_{n(t)}$ disagree, so the correction is never a step in the direction w_t already points: it turns w_t rather than stretching it.

Now put the two together. Each fact describes a *single* correction, so apply each of them once per correction and let the bounds telescope. Start PLA from $w_0 = \mathbf{0}$ and suppose it has made T corrections.

Chaining Lemma 2.3.3 through $t = 0, 1, \dots, T-1$ adds the same positive amount T times over:

$$w_f^\top w_T \geq \underbrace{w_f^\top w_0}_{=0} + T \min_n y_n w_f^\top x_n = T \min_n y_n w_f^\top x_n,$$

so the inner product grows at least *linearly* in T . Chaining Lemma 2.3.4 the same way,

$$\|w_T\|^2 \leq \underbrace{\|w_0\|^2}_{=0} + T \max_n \|y_n x_n\|^2 = T \max_n \|x_n\|^2,$$

where $\|y_n x_n\| = \|x_n\|$ because $|y_n| = 1$. So the *squared* length grows at most linearly, which means $\|w_T\|$ itself grows at most like $\sqrt{T}$.

Dividing the first by the second is what finally answers the ‘(really?)’. The quantity

$$\frac{w_f^\top w_T}{\|w_f\| \|w_T\|}$$

is an inner product of two *unit* vectors, i.e. the cosine of the angle between w_f and w_T ; dividing by $\|w_T\|$ has removed precisely the loophole that made Lemma 2.3.3 inconclusive, so if this grows, the angle really does shrink. Feeding in the two displays — the numerator is bounded below, the denominator above — gives

$$\frac{w_f^\top w_T}{\|w_f\| \|w_T\|} \geq \frac{T \min_n y_n w_f^\top x_n}{\|w_f\| \sqrt{T \max_n \|x_n\|^2}} = \sqrt{T} \cdot \text{constant},$$

the constant being built from $\mathbf{w}_f$ and the data alone: every appearance of T has collected into the $T/\sqrt{T} = \sqrt{T}$ out front.

That is already the whole theorem. The left-hand side is a cosine and so never exceeds 1, while the right-hand side grows without bound as T does; the two can hold together only for finitely many T . PLA therefore corrects only finitely many mistakes — it halts. All that is left is to name the constant, which is what the next exercise does.

Exercise 2.3.5 (Fun Time). Let us upper-bound T , the number of mistakes that PLA ‘corrects’. Define

$$R^2 = \max_n \|\mathbf{x}_n\|^2, \quad \rho = \min_n y_n \frac{\mathbf{w}_f^\top}{\|\mathbf{w}_f\|} \mathbf{x}_n.$$

We want to show $T \leq \square$. Express the upper bound $\square$ by the two terms above.

- (1) R/ρ ;
- (2) R^2/ρ^2 ;
- (3) R/ρ^2 ;
- (4) ρ^2/R^2 .

Reference Answer. (2), and the calculation is exactly [Theorem 2.3.6](#) below: naming the two quantities turns the anonymous constant of the previous page into ρ/R , and a cosine’s ceiling of 1 then caps T at $1/(\rho/R)^2 = R^2/\rho^2$. ■

§2.3.iii. The Guarantee

Theorem 2.3.6 (PLA halts)

Let $\mathcal{D}$ be linearly separable, let $\mathbf{w}_f$ be a perfect weight vector for it, and put

$$R^2 = \max_n \|\mathbf{x}_n\|^2, \quad \rho = \min_n y_n \frac{\mathbf{w}_f^\top}{\|\mathbf{w}_f\|} \mathbf{x}_n > 0.$$

Then PLA started from $\mathbf{w}_0 = \mathbf{0}$ makes at most

$$T \leq \frac{R^2}{\rho^2}$$

corrections, and in particular it halts.

Proof. Everything has been derived already; naming R and ρ is all that is left. Note first that $\rho > 0$: it is $\min_n y_n \mathbf{w}_f^\top \mathbf{x}_n / \|\mathbf{w}_f\|$, and that minimum was shown strictly positive in Fact 1 because $\mathbf{w}_f$ is perfect. With this notation the two chained bounds read

$$\mathbf{w}_f^\top \mathbf{w}_T \geq T \min_n y_n \mathbf{w}_f^\top \mathbf{x}_n = T \rho \|\mathbf{w}_f\|, \quad \|\mathbf{w}_T\|^2 \leq T \max_n \|\mathbf{x}_n\|^2 = T R^2,$$

so that

$$\frac{\mathbf{w}_f^\top \mathbf{w}_T}{\|\mathbf{w}_f\| \|\mathbf{w}_T\|} \geq \frac{T \rho \|\mathbf{w}_f\|}{\|\mathbf{w}_f\| \sqrt{T R^2}} = \sqrt{T} \cdot \frac{\rho}{R}.$$

The anonymous constant was therefore ρ/R all along. By [Remark 2.3.7](#) the left-hand side is at most 1, so $\sqrt{T} \rho/R \leq 1$, i.e. $T \leq R^2/\rho^2$. Since the right-hand side is a finite number fixed by $\mathcal{D}$ and $\mathbf{w}_f$, PLA cannot correct mistakes forever. ■

Remark 2.3.7 (Is it really a cosine in $\mathbb{R}^{d+1}$?). We have been calling $\mathbf{w}_f^\top \mathbf{w}_T / (\|\mathbf{w}_f\| \|\mathbf{w}_T\|)$ a cosine, which is obvious for $d = 2$ but deserves a word once $\mathbf{w}$ lives in $\mathbb{R}^{d+1}$ for large d . It is still a cosine, for two reasons, and the proof needs only the weaker of them.

- **What the proof actually uses** is only that the quantity never exceeds 1, and that is the Cauchy–Schwarz inequality $|\mathbf{u}^\top \mathbf{v}| \leq \|\mathbf{u}\| \|\mathbf{v}\|$, valid in $\mathbb{R}^m$ for every m . No geometry is required.
- **Geometrically** it is a cosine in the ordinary sense too. Two vectors always span a subspace of dimension at most 2; inside that plane they make an ordinary angle, and $\mathbf{u}^\top \mathbf{v} / (\|\mathbf{u}\| \|\mathbf{v}\|)$ is its cosine. Indeed this is how the angle between two vectors is *defined* in any dimension — Cauchy–Schwarz is exactly the statement that makes the definition legitimate, by putting the ratio in $[-1, 1]$ so that an arccos exists.

So the picture of $\mathbf{w}_t$ slowly rotating towards $\mathbf{w}_f$ survives intact in any number of dimensions.

Remark 2.3.8. Both constants have a geometric reading, and the bound is more memorable through it. R is the radius of the data: the distance from the origin to the farthest example. And since $\mathbf{w}_f / \|\mathbf{w}_f\|$ is a unit normal of the perfect separating hyperplane, ρ is the smallest signed distance from an example to that hyperplane — the *margin*. So

$$T \leq \left(\frac{\text{radius of the data}}{\text{margin}} \right)^2 :$$

the more comfortably the data separates, the sooner PLA is done. This quantity returns as the central object of the chapter on support vector machines, which does not merely tolerate a large margin but goes looking for one.

§2.4. Non-Separable Data


— slides p.17–22 —

§2.4.i. What We Have and What We Lack

- **Guarantee:** as long as $\mathcal{D}$ is linearly separable and we correct by mistake, the inner product of $\mathbf{w}_f$ and $\mathbf{w}_t$ grows fast while the length of $\mathbf{w}_t$ grows slowly; the PLA lines get more and more aligned with $\mathbf{w}_f$, so PLA halts.
- **Pros:** simple to implement, fast, and works in any dimension d .
- **Cons:** it *assumes* a linearly separable $\mathcal{D}$ in order to halt — and that property is unknown in advance, with the circularity that if we already knew $\mathbf{w}_f$ we would not need PLA at all. Moreover we are not fully sure how long halting takes, since ρ depends on the unknown $\mathbf{w}_f$, though in practice it is fast.


Question 2.4.1. What if $\mathcal{D}$ is not linearly separable?

In practice the usual reason is not that the truth is curved, but that the data is **noisy**: the recorded y_n is sometimes simply wrong, or two customers with identical records behaved differently. The learning flow gains one term.



Question 2.4.2. How can we at least get $g \approx f$ on a noisy $\mathcal{D}$?

§2.4.ii. A Line with Noise Tolerance



If the noise is ‘little’, then $y_n = f(\mathbf{x}_n)$ *usually*, and therefore $g \approx f$ on $\mathcal{D}$ means $y_n = g(\mathbf{x}_n)$ *usually*. That suggests giving up on perfection and asking for the fewest mistakes instead:

$$\mathbf{w}_g \leftarrow \arg \min_{\mathbf{w}} \sum_{n=1}^N \mathbb{I}[y_n \neq \text{sign}(\mathbf{w}^\top \mathbf{x}_n)].$$

Remark 2.4.3. This is exactly the right thing to ask for and, unfortunately, NP-hard to solve. The obstruction is the Iverson bracket: the objective counts, so it is piecewise constant, and there is no slope anywhere to follow downhill. Much of *How can machines learn?* is about replacing this count with a smoother surrogate; here we take the cheaper route instead and modify the algorithm rather than the objective.

§2.4.iii. The Pocket Algorithm

Question 2.4.4. Can we modify PLA to get an ‘approximately good’ g ?

The fix is almost embarrassingly simple. PLA already wanders among candidate lines; it just has no memory, and on non-separable data it wanders forever. So give it a pocket: keep the best line seen so far, and keep wandering.

Algorithm 2.3: Pocket Algorithm**Input:** training examples $\mathcal{D}$, an iteration budget**Output:** $\hat{\mathbf{w}}$, called $\mathbf{w}_{\text{POCKET}}$ and returned as g

```

1 initialize the pocket weights  $\hat{\mathbf{w}}$ ;
2 for  $t = 0, 1, \dots$  until enough iterations do
3   find a (random) mistake of  $\mathbf{w}_t$ , called  $(\mathbf{x}_{n(t)}, y_{n(t)})$ ;
4   (try to) correct the mistake by  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + y_{n(t)}\mathbf{x}_{n(t)}$ ;
5   if  $\mathbf{w}_{t+1}$  makes fewer mistakes on  $\mathcal{D}$  than  $\hat{\mathbf{w}}$  then  $\hat{\mathbf{w}} \leftarrow \mathbf{w}_{t+1}$ ;
6 return  $\hat{\mathbf{w}}$  as  $g$ ;

```

Note what changed and what did not. The update rule is untouched — the black lines of PLA are still there — and the only additions are the pocket and the iteration budget, which is now needed because there is no “no more mistakes” to stop on.

A simple modification of PLA to find (somewhat) ‘best’ weights.

Exercise 2.4.5 (Fun Time). Should we use pocket or PLA? Since we do not know whether $\mathcal{D}$ is linearly separable in advance, we may decide to just go with pocket instead of PLA. If $\mathcal{D}$ is actually linearly separable, what is the difference between the two?

- (1) pocket on $\mathcal{D}$ is slower than PLA;
- (2) pocket on $\mathcal{D}$ is faster than PLA;
- (3) pocket on $\mathcal{D}$ returns a better g in approximating f than PLA;
- (4) pocket on $\mathcal{D}$ returns a worse g in approximating f than PLA.

Reference Answer. (1). Because pocket needs to check whether $\mathbf{w}_{t+1}$ is better than $\hat{\mathbf{w}}$ in each iteration — a full pass over $\mathcal{D}$ — it is slower than PLA. On a linearly separable $\mathcal{D}$, $\mathbf{w}_{\text{POCKET}}$ ends up the same as $\mathbf{w}_{\text{PLA}}$, both making no mistakes; so the price of not knowing in advance is time, not quality. ■

Summary

Perceptron Hypothesis Set hyperplanes / linear classifiers in $\mathbb{R}^d$.

Perceptron Learning Algorithm (PLA) correct mistakes and improve iteratively.

Guarantee of PLA no mistake eventually, if linearly separable.

Non-Separable Data hold somewhat ‘best’ weights in a pocket.

We now own one complete learning model, $\mathcal{H}$ and $\mathcal{A}$ both. It answers a *yes/no* question, from a vector of *real* features, given *labelled* examples handed over *all at once*. Every one of those emphasised choices can be varied, and the next chapter takes stock of the resulting zoo of learning problems.

3. Types of Learning

Lecture 3

As previously seen . Section 2.2 produced one complete learning model: PLA takes a linearly separable $\mathcal{D}$ and the perceptrons $\mathcal{H}$, and returns a hypothesis g . That is one point in a very large space of possibilities.

16 Sep.


The frame of Section 1.4 has several slots — the output space $\mathcal{Y}$, the labels y_n , the way the data reaches us, the input space $\mathcal{X}$ — and so far every one of them has been filled in the simplest way available. This chapter varies them one at a time, which is the cheapest way to see what the frame can hold. Nothing here is proved; the point is a map of the territory, and the names on it recur for the rest of the course.

§3.1. Learning with Different Output Space $\mathcal{Y}$

— slides p.2–8 —

§3.1.i. The Credit Approval Problem Revisited

Credit approval answers one yes/no question, so its output space has exactly two elements.



$\mathcal{Y} = \{-1, +1\}$: **binary classification.**


§3.1.ii. Binary Classification

Definition 3.1.1 (Binary classification). A learning problem whose output space has exactly two elements, $\mathcal{Y} = \{-1, +1\}$, is called **binary classification**. The two elements are only names: good/bad, spam/non-spam, sick/not sick.

The same output space covers an enormous range of questions:

- credit approve/disapprove;
- email spam/non-spam;
- patient sick/not sick;
- ad profitable/not profitable;
- answer correct/incorrect (KDDCup 2010).

Geometrically it is the picture we already know — a boundary with one class on each side, easy or hard depending on how the data falls.




A core and important problem, with many tools that serve as building blocks of other tools.

§3.1.iii. Multiclass Classification

Definition 3.1.2 (Multiclass classification). A learning problem whose output space is finite and explicitly listed, $\mathcal{Y} = \{1, 2, \dots, K\}$ with $K \geq 2$, is called **multiclass classification**; Definition 3.1.1 is the special case $K = 2$.

Example 3.1.3 (Coin recognition). Classify US coins — 1c, 5c, 10c, 25c — by their (size, mass). The four denominations form four clusters, because a heavier coin is usually a larger one.



Here $\mathcal{Y} = \{1c, 5c, 10c, 25c\}$, so $K = 4$.

Other multiclass problems: written digits $\Rightarrow 0, 1, \dots, 9$; pictures $\Rightarrow$ apple, orange, strawberry; emails $\Rightarrow$ spam, primary, social, promotion, update (Google).

Many applications in practice, especially for ‘recognition’.

§3.1.iv. Regression

Definition 3.1.4 (Regression). A learning problem whose output space is the whole real line, $\mathcal{Y} = \mathbb{R}$, is called **regression**; if the output is confined to an interval, $\mathcal{Y} = [\text{lower}, \text{upper}] \subset \mathbb{R}$, it is **bounded regression**. The difference from classification is that $\mathcal{Y}$ now carries an order and a distance, so being wrong by a little is better than being wrong by a lot.

Keep the patient, change the question, and the output space changes with it:

- binary classification: patient features $\Rightarrow$ sick or not;
- multiclass classification: patient features $\Rightarrow$ which type of cancer;
- *regression*: patient features $\Rightarrow$ how many days before recovery.

Regression was studied deeply in statistics long before machine learning existed, which is why so many of its tools arrive ready-made. Other regression problems: company data $\Rightarrow$ stock price; climate data $\Rightarrow$ temperature.

Also core and important, with many ‘statistical’ tools as building blocks of other tools.

§3.1.v. Structured Learning

Definition 3.1.5 (Structured learning). **Structured learning** is a learning problem whose outputs are composite objects — a tag for every word of a sentence, a folding of a protein, a parse tree — whose parts constrain one another. Formally $\mathcal{Y}$ is still a finite set of ‘classes’, but an enormous one that is defined implicitly by those constraints rather than enumerated.

Example 3.1.6 (Sequence tagging). Multiclass classification sends a word to a word class. Structured learning sends a whole sentence to a structure — the class of every word at once:

$\underbrace{\text{I}}_{\text{pronoun}} \quad \underbrace{\text{love}}_{\text{verb}} \quad \underbrace{\text{ML}}_{\text{noun}}$

so that $\mathcal{Y} = \{\text{PVN}, \text{PVP}, \text{NVN}, \text{PV}, \dots\}$, not including VVVVV.

That exclusion is the whole difficulty. The output space is a huge multiclass problem — a structure is a kind of hyperclass — but without an *explicit* class definition: the valid structures are constrained by grammar, and nobody wrote the list down. Other structured problems: protein data $\Rightarrow$ protein folding; speech data $\Rightarrow$ speech parse tree.

A fancy but complicated learning problem.

§3.1.vi. Mini Summary

Note 3.1.7 (Learning with different output space $\mathcal{Y}$).

- binary classification: $\mathcal{Y} = \{-1, +1\}$;
- multiclass classification: $\mathcal{Y} = \{1, 2, \dots, K\}$;
- regression: $\mathcal{Y} = \mathbb{R}$;
- structured learning: $\mathcal{Y} = \text{structures}$;
- ... and a lot more!!

Core tools: binary classification and regression.

Exercise 3.1.8 (Fun Time). The entrance system of the school gym, which does automatic face recognition based on machine learning, is built to charge four different groups of users differently: Staff, Student, Professor, Other. What type of learning problem best fits the need of the system?

- (1) binary classification;
- (2) multiclass classification;
- (3) regression;
- (4) structured learning.

Reference Answer. (2). There is an ‘explicit’ $\mathcal{Y}$ that contains four classes. ■

§3.2. Learning with Different Data Label y_n

— slides p.9–15 —


§3.2.i. Supervised Learning

Definition 3.2.1 (Supervised learning). Learning from data in which every input carries its own label,

$$\mathcal{D} = \{(\mathbf{x}_n, y_n)\}_{n=1}^N,$$

is called **supervised learning**. This is the only kind we have used so far.

Everything so far has quietly assumed that the data arrives fully labelled.



Supervised learning: every x_n comes with a corresponding y_n .


§3.2.ii. Unsupervised Learning

Definition 3.2.2 (Unsupervised learning). Learning from data that carries no labels at all,

$$\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N,$$

is called **unsupervised learning**. With no y_n there is no target f to approximate in the sense of Section 1.4, so what counts as a good g has to be decided separately for each problem.

Now delete the labels. The coins are still there, and they still fall into clusters — but nothing says which cluster is the quarter.



Other clustering problems: articles $\Rightarrow$ topics; consumer profiles $\Rightarrow$ consumer groups. Clustering is only one of several unsupervised problems, and they are best remembered as unsupervised shadows of the supervised ones — each keeps the *shape* of a supervised problem but has to invent its own notion of a right answer.

Definition 3.2.3 (Clustering). **Clustering** learns a map

$$\{\mathbf{x}_n\} \Rightarrow \text{cluster}(\mathbf{x})$$

assigning each input to one of finitely many groups, where the groups themselves are discovered rather than given. It is roughly ‘unsupervised multiclass classification’, with the essential difference that the classes have no names and no prescribed number — articles $\Rightarrow$ topics.

Definition 3.2.4 (Density estimation). **Density estimation** learns a map

$$\{\mathbf{x}_n\} \Rightarrow \text{density}(\mathbf{x})$$

reporting how tightly the inputs crowd around each point of $\mathcal{X}$. Since a density is a non-negative number, this is roughly ‘unsupervised bounded regression’ — traffic reports with location $\Rightarrow$ dangerous areas.

Definition 3.2.5 (Outlier detection). **Outlier detection** learns a map

$$\{\mathbf{x}_n\} \Rightarrow \text{unusual}(\mathbf{x})$$

flagging the inputs that do not look like the rest. It is roughly an extreme ‘unsupervised binary classification’: extreme because the two classes are wildly unbalanced — almost everything is normal, and it is the rare case that matters — so accuracy is a useless measure of success. Internet logs $\Rightarrow$ intrusion alert.

... and a lot more!!

Unsupervised learning: diverse, with possibly very different performance goals.


§3.2.iii. Semi-supervised Learning

Definition 3.2.6 (Semi-supervised learning). Learning from data in which only a few inputs carry labels,

$$\mathcal{D} = \{(\mathbf{x}_n, y_n)\}_{n=1}^{\ell} \cup \{\mathbf{x}_n\}_{n=\ell+1}^N, \quad \ell \ll N,$$

is called **semi-supervised learning**. The unlabelled majority is not useless: it shows where the inputs lie, and a boundary is more believable if it passes through a thin region of them.

Between the two extremes sits the case that is commonest in practice: labels exist, but only for a few examples, because labelling costs money.



Other semi-supervised problems: face images with a few labelled $\Rightarrow$ face identifier (Facebook);
medicine data with a few labelled $\Rightarrow$ medicine effect predictor.

Semi-supervised learning: leverage unlabelled data to avoid ‘expensive’ labelling.

§3.2.iv. Reinforcement Learning

Definition 3.2.7 (Reinforcement learning). Learning from data of the form

$$\mathcal{D} = \{(\mathbf{x}_n, \tilde{y}_n, \text{goodness}(\tilde{y}_n))\}_{n=1}^N,$$


where $\tilde{y}_n$ is an action that was actually taken and the third entry says only how well it turned out, is called **reinforcement learning**. The correct label y_n is never revealed — one learns that a choice was good or bad, not what the right choice would have been.

A ‘very different’ but natural way of learning, and the easiest way in is the one the slides take.

Example 3.2.8 (Teach your dog to sit).

You say ‘sit down’.

- *The dog pees on the ground.* BAD DOG. THAT’S A VERY WRONG ACTION. You cannot easily show the dog that $y_n = \text{sit}$ when $\mathbf{x}_n = \text{‘sit down’}$, but you can *punish* it, which says that $\tilde{y}_n = \text{pee}$ is wrong.
- *The dog sits down.* GOOD DOG. LET ME GIVE YOU SOME COOKIES. You still cannot show that $y_n = \text{sit}$, but you can *reward* it, which says that $\tilde{y}_n = \text{sit}$ is good.



Note what the training signal is. It is not the correct label y_n but a triple $(\mathbf{x}, \tilde{y}, \text{goodness})$: some action was taken, and we learn only how good *that* action was — never what the best one would have been. Other problems of the same shape: (customer, ad choice, ad click earning) $\Rightarrow$ ad system; (cards, strategy, winning amount) $\Rightarrow$ black jack agent.

Reinforcement: learn with ‘partial/implicit information’ (often sequentially).

§3.2.v. Mini Summary

Note 3.2.9 (Learning with different data label y_n).

- supervised: all y_n ;
- unsupervised: no y_n ;

- semi-supervised: some y_n ;
- reinforcement: implicit y_n by goodness($\tilde{y}_n$);
- ... and more!!

Core tool: supervised learning.

Exercise 3.2.10 (Fun Time). To build a tree recognition system, a company decides to gather one million pictures on the Internet. Then it asks each of its 10 members to view 100 pictures and record whether each picture contains a tree. The pictures and records are then fed to a learning algorithm to build the system. What type of learning problem does the algorithm need to solve?

- (1) supervised;
- (2) unsupervised;
- (3) semi-supervised;
- (4) reinforcement.

Reference Answer. (3). The 1,000 records are the labelled $(\mathbf{x}_n, y_n)$; the other 999,000 pictures are unlabelled $\mathbf{x}_n$. ■

§3.3. Learning with Different Protocol $f \Rightarrow (\mathbf{x}_n, y_n)$

— slides p.16–21 —

The third slot is not about what the data *is* but about how it reaches the algorithm — the arrow from f down to $\mathcal{D}$.

§3.3.i. Batch Learning

Definition 3.3.1 (Batch learning). Under the **batch** protocol the entire $\mathcal{D}$ is handed to $\mathcal{A}$ at once, and $\mathcal{A}$ returns a single hypothesis g that is then held fixed.

Example 3.3.2. 1. The coin recogniser of [Example 3.1.3](#) is handed its whole data set at once and learns from all of it: batch supervised multiclass classification.

2. a batch of (email, spam?) $\Rightarrow$ spam filter
3. a batch of (patient, cancer) $\Rightarrow$ cancer classifier; a batch of patient data $\Rightarrow$ groups of patients

Remark 3.3.3. The last one unsupervised, so the protocol is independent of the previous two sections.

Batch learning: a very common protocol.

§3.3.ii. Online Learning

Definition 3.3.4 (Online learning). Under the **online** protocol the examples arrive one at a time: at step t the learner sees $\mathbf{x}_t$, predicts with its current g_t , is then told y_t , and updates to g_{t+1} . There is no separate training phase — the hypothesis in use is the one still being learned.

Example 3.3.5 (A spam filter that improves). A *batch* spam filter learns from known (email, spam?) pairs and then predicts with a fixed g . An *online* spam filter instead repeats, for $t = 1, 2, \dots$:

1. observe an email $\mathbf{x}_t$;
2. predict its spam status with the current $g_t(\mathbf{x}_t)$;
3. receive the ‘desired label’ y_t from the user, and then update g_t with $(\mathbf{x}_t, y_t)$.


Question 3.3.6. Two connections to what we have already learned: PLA can be easily adapted to the online protocol — how? And reinforcement learning is often done online — why?

Answer. PLA already works one example at a time: it looks at a single $(\mathbf{x}_{n(t)}, y_{n(t)})$, and either leaves $\mathbf{w}$ alone or corrects it. Feed it a stream instead of a fixed $\mathcal{D}$ and nothing in [algorithm 2.1](#) has to change. Reinforcement learning is online almost by definition: the signal $(\mathbf{x}, \tilde{y}, \text{goodness})$ only exists once an action has been taken, so the data cannot be collected in advance — it is generated by the learner’s own behaviour. ■

Online: the hypothesis ‘improves’ through receiving data instances sequentially.

§3.3.iii. Active Learning

Definition 3.3.7 (Active learning). Under the **active** protocol the learner *chooses* which $\mathbf{x}_n$ it wants labelled and queries that label, instead of accepting whichever examples happen to arrive. The aim is to reach a good g from fewer labels by asking about the inputs that are most informative.



Note 3.3.8 (Protocol $\Leftrightarrow$ learning philosophy).

- batch: ‘duck feeding’ — everything is pushed in at once;
- online: ‘passive sequential’ — take what arrives, in order;
- active: ‘question asking’ (sequentially) — query the y_n of a *chosen* $\mathbf{x}_n$.

Active: improve the hypothesis with fewer labels (hopefully) by asking questions strategically.

§3.3.iv. Mini Summary

Note 3.3.9 (Learning with different protocol $f \Rightarrow (x_n, y_n)$).

- batch: all known data;
- online: sequential (passive) data;
- active: strategically-observed data;
- ... and more!!

Core protocol: batch.

Exercise 3.3.10 (Fun Time). A photographer has 100,000 pictures, each containing one baseball player. He wants to categorize the pictures automatically by the player inside. He starts by categorizing 1,000 pictures himself, and then writes an algorithm that tries to categorize the other pictures if it is ‘confident’ on the category, while pausing for — and learning from — human input if not. What protocol best describes the nature of the algorithm?

- (1) batch;
- (2) online;
- (3) active;
- (4) random.


Reference Answer. (3). The algorithm takes an active but naïve strategy: ask when ‘confused’. You should probably do the same when taking a class. :-)

§3.4. Learning with Different Input Space $\mathcal{X}$

— slides p.22–28 —

§3.4.i. Concrete Features

Definition 3.4.1 (Concrete features). $\mathcal{X} \subseteq \mathbb{R}^d$ consists of **concrete features** when each of its d dimensions carries a sophisticated physical meaning — annual salary, mass, tumour size — normally because a human who understood the task chose to measure it.



Concrete features: each dimension of $\mathcal{X} \subseteq \mathbb{R}^d$ represents ‘sophisticated physical meaning’.


Every input we have met so far is of this kind: (size, mass) for coin classification, customer information for credit approval, patient information for cancer diagnosis. What they share is that somebody decided which quantities were worth measuring — the features carry ‘human intelligence’ about the learning task, put there before learning began.

Concrete features: the ‘easy’ ones for ML.

§3.4.ii. Raw Features

Definition 3.4.2 (Raw features). $\mathcal{X}$ consists of **raw features** when each dimension has only a simple physical meaning — the darkness of one pixel, one sample of a waveform. Any dimension on its own says almost nothing; the information lives in combinations of very many of them.

Example 3.4.3 (Digit recognition). Features $\Rightarrow$ meaning of digit: a typical supervised multi-class classification problem. But what are the features?




The honest answer is that a 16×16 grey image is just a list of 256 numbers,

$$\mathbf{x} \equiv (0, 0, 0.9, 0.6, \dots) \in \mathbb{R}^{256},$$

and each of those numbers means only ‘how dark is this one pixel’. That is a *simple* physical meaning, and it makes the problem harder for machine learning, not easier — no single pixel

tells you anything about the digit. Compare what a human would hand over instead: two numbers, symmetry and density, chosen because they happen to separate the classes.



Other problems with raw features: image pixels, speech signals, and so on.

Raw features: often need humans or machines to convert them to concrete ones.

§3.4.iii. Abstract Features

Definition 3.4.4 (Abstract features). $\mathcal{X}$ consists of **abstract features** when its dimensions have no (or very little) physical meaning at all — identifiers such as a userid or an itemid, where even the ordering of the numbers is an accident of bookkeeping.

Example 3.4.5 (Rating prediction, KDDCup 2011). Given previous (userid, itemid, rating) tuples, predict the rating that some userid would give to some itemid. This is a regression problem with $\mathcal{Y} \subseteq \mathbb{R}$ as the rating and $\mathcal{X} \subseteq \mathbb{N} \times \mathbb{N}$ as (userid, itemid).

User 5,001 is not in any sense one more than user 5,000, so this is harder still than raw features. Other problems with abstract features: student IDs in an online tutoring system (KDDCup 2010), advertisement IDs in an online ad system. That these two competitions were won by extracting factors from bare IDs is exactly the story of [Section 1.3](#).

Abstract: again needs ‘feature conversion/extraction/construction’.

§3.4.iv. Mini Summary

Note 3.4.6 (Learning with different input space $\mathcal{X}$).

- concrete: sophisticated (and related) physical meaning;
- raw: simple physical meaning;
- abstract: no (or little) physical meaning;
- ... and more!!

‘Easy’ input: concrete.

Exercise 3.4.7 (Fun Time). Consider the problem of building an online image advertisement system that shows users the most relevant images. What features can you choose to use?

- (1) concrete;
- (2) concrete, raw;
- (3) concrete, abstract;
- (4) concrete, raw, abstract.

Reference Answer. (4). Concrete user features, raw image features, and maybe abstract user/image IDs. ■

Summary

Output space $\mathcal{Y}$ *classification, regression*, structured.

Data label y_n *supervised*, un/semi-supervised, reinforcement.

Protocol $f \Rightarrow (x_n, y_n)$ *batch*, online, active.

Input space $\mathcal{X}$ *concrete*, raw, abstract.

The emphasised entry in each line is the one this course treats as its core tool, and the combination of all four — batch supervised binary classification on concrete features — is exactly the setting PLA already solves. Everything else in the list is reached by relaxing one slot at a time.

Which makes the next question awkward. We have a solved case and a map of its neighbours, but not one word yet about whether any of it works on data we have never seen. The next chapter argues, seriously, that learning is impossible.

4. Feasibility of Learning

Lecture 4

As previously seen, [Section 2.2](#) produced a working learning model, and [Section 3.1](#) onwards mapped the settings around it. Every guarantee obtained so far — PLA halts, PLA makes no mistake on $\mathcal{D}$ — is a statement about the examples we were handed.

23 Sep.


$\mathcal{D}$ is not what we are paid for. The bank wants the next applicant classified, not the ones already in its files, and nothing in the previous three chapters connects the two. This chapter takes that gap seriously enough to argue first that it cannot be bridged at all, and then that it can — once we agree to settle for a guarantee of a weaker kind. Four steps: a proof that learning is impossible, a probabilistic escape, the escape applied to one hypothesis, and finally to a whole hypothesis set.

§4.1. Learning is Impossible?

— slides p.2–6 —

§4.1.i. A Learning Puzzle

Before any mathematics, a test of your own human learning, with six examples.



Each $\mathbf{x}_n$ is a 3×3 picture of black and white cells; three of them are labelled -1 and three $+1$. What does your g say about the seventh?

§4.1.ii. Two Controversial Answers

Whatever you say about $g(\mathbf{x})$, an answer with reasons behind it is available for the other side as well.

truth $f(\mathbf{x}) = +1$ because ...	truth $f(\mathbf{x}) = -1$ because ...
• symmetry $\Leftrightarrow +1$; • (black or white count = 3) or (black count = 4 and middle-top black) $\Leftrightarrow +1$.	• left-top black $\Leftrightarrow -1$; • middle column contains at most 1 black and right-top white $\Leftrightarrow -1$.

Observation . All four rules are correct on all six examples, and all four fire on the test picture. The last one is worth checking by hand: its main diagonal is black, so it is symmetric and its black count is 3, which the left-hand rules read as +1; its left-top cell is black and its middle column holds one black with a white right-top cell, which the right-hand rules read as −1.

So the data does not decide. Any answer you give can be called wrong by a teacher who had one of the other rules in mind all along, and that teacher never has to admit to cheating — the rule was consistent with everything you were shown.

All valid reasons: your adversarial teacher can always call you ‘didn’t learn’.
:-)

§4.1.iii. A ‘Simple’ Binary Classification Problem

The puzzle is easier to argue about once it is small enough to write out completely. Take

$$\mathcal{X} = \{0, 1\}^3, \quad \mathcal{Y} = \{\circ, \times\},$$

so that $\mathcal{X}$ has eight points and there are only $2^8 = 256$ candidate targets f — few enough that $\mathcal{H}$ can be taken to be *all* of them. Suppose five of the eight points come to us labelled:

$\mathbf{x}_n$	$y_n = f(\mathbf{x}_n)$
000	○
001	×
010	×
011	○
100	×

Question 4.1.1. Pick $g \in \mathcal{H}$ with $g(\mathbf{x}_n) = y_n$ for every n , the way PLA does. Does $g \approx f$?

§4.1.iv. No Free Lunch

Fixing the five labels kills 2^5 of the 256 degrees of freedom and leaves $2^3 = 8$ targets standing. Here they are, next to any g we might have picked.

	$\mathbf{x}$	y	g	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8
$\mathcal{D}$ {	000	○	○	○	○	○	○	○	○	○	○
	001	×	×	×	×	×	×	×	×	×	×
	010	×	×	×	×	×	×	×	×	×	×
	011	○	○	○	○	○	○	○	○	○	○
	100	×	×	×	×	×	×	×	×	×	×
	101		?	○	○	○	○	×	×	×	×
	110		?	○	○	×	×	○	○	×	×
	111		?	○	×	○	×	○	×	○	×

The top block is the same in every column, which is the good news:

- $g \approx f$ inside $\mathcal{D}$: sure!

- $g \approx f$ outside $\mathcal{D}$: **No!** — but that is really what we want.

The bottom block is the bad news, and it is as bad as it can be: for each of the three unseen inputs, four of the surviving targets say $\circ$ and four say $\times$. Whatever g fills into a ‘?’, half of the candidates disagree with it. This is worth saying once in general.

Proposition 4.1.2 (No free lunch)

Let $\mathcal{H}$ contain every $f: \mathcal{X} \rightarrow \{-1, +1\}$, let $\mathcal{D}$ be any set of labelled examples, and let $\mathbf{x} \notin \mathcal{D}$. Among the targets that agree with $\mathcal{D}$, exactly half satisfy $f(\mathbf{x}) = +1$ and half satisfy $f(\mathbf{x}) = -1$. Consequently every algorithm, no matter how it computes g , is wrong at $\mathbf{x}$ for exactly half of them.

Proof. Write S for the set of targets that agree with $\mathcal{D}$ — the eight columns of the table above, in the example. Given $f \in S$, let f' be the target that copies f at every input except $\mathbf{x}$, where it answers the other way: $f'(\mathbf{x}) = -f(\mathbf{x})$.

That f' is again in S is the one place the hypothesis $\mathbf{x} \notin \mathcal{D}$ is used: the single value that was changed sits at $\mathbf{x}$, which is not one of the examples, so f' still reproduces every label of $\mathcal{D}$. Flipping twice returns the original, $(f')' = f$, and f' is never f itself, since the two differ at $\mathbf{x}$ by construction. A rule that pairs each element with another one and is undone by applying it again leaves nobody over, so $f \mapsto f'$ cuts S into pairs, and the two members of a pair give opposite answers at $\mathbf{x}$.

Half of S therefore answers $+1$ at $\mathbf{x}$ and half answers -1 . Now $g(\mathbf{x})$ is a single fixed value, whatever $\mathcal{A}$ did to arrive at it, so within each pair it agrees with one member and disagrees with the other. ■

In the table, with $\mathbf{x} = 101$, that pairing is $f_1 \leftrightarrow f_5$, $f_2 \leftrightarrow f_6$, $f_3 \leftrightarrow f_7$, $f_4 \leftrightarrow f_8$: the members of a pair agree on every other row and differ only in that one.

The proof uses nothing about the algorithm — it never appears — which is exactly why no cleverer $\mathcal{A}$ can escape. The trouble is the assumption that f is completely unrestricted.

Learning from $\mathcal{D}$ (to infer something outside $\mathcal{D}$) is doomed if any ‘unknown’ f can happen. :-)

Exercise 4.1.3 (Fun Time). This is a popular ‘brain-storming’ problem, with a claim that 2% of the world’s cleverest population can crack its ‘hidden pattern’:

$$(5, 3, 2) \rightarrow 151022, \quad (7, 2, 5) \rightarrow ?$$

It is a learning problem with $N = 1$, $\mathbf{x}_1 = (5, 3, 2)$, $y_1 = 151022$. Learn a hypothesis from the one example to predict on $\mathbf{x} = (7, 2, 5)$. What is your answer?

- (1) 151026;
- (2) 143547;
- (3) I need more examples to get the correct answer;
- (4) there is no ‘correct’ answer.

Reference Answer. (4). Following the same nature of the no-free-lunch problems discussed, we cannot hope to be correct under this ‘adversarial’ setting. The designer’s answer is (2): the first two digits are $x_1 \cdot x_2$, the next two $x_1 \cdot x_3$, and the last two $x_1 \cdot x_2 + x_1 \cdot x_3 - x_2$. Knowing

that does not make (2) any more correct than (1) — it only tells you which f this particular teacher had in mind. ■


§4.2. Probability to the Rescue

— slides p.7–12 —

§4.2.i. Inferring Something Unknown

Proposition 4.1.2 closes one door: the unknown target f outside $\mathcal{D}$ cannot be inferred. Before giving up, it is worth asking whether *any* unknown quantity is ever inferred from a sample, in any other corner of mathematics — and one is, routinely.

Consider a bin holding very many marbles, each orange or green.




Do we *know* what fraction of them is orange? No: nobody counted, and the bin may hold more marbles than we could ever count.

Question 4.2.1. Can you *infer* the orange probability?

§4.2.ii. Statistics 101: Inferring Orange Probability

The statistician's move is to grab a handful without looking.



Notation . μ is the orange probability in the bin and $1 - \mu$ the green one; μ is **unknown**. Draw N marbles independently, each with replacement, and let ν be the orange fraction in that sample, $1 - \nu$ the green one; ν is **known**, because it is sitting in front of us. ‘In-sample’ always refers to ν , ‘out-of-sample’ to μ .

Question 4.2.2. Does the in-sample ν say anything about the out-of-sample μ ?

§4.2.iii. Possible versus Probable

The answer has to be given twice, because the two obvious readings of the question have opposite answers.

No!	Yes!
Possibly not: the sample can come out mostly green while the bin is mostly orange.	Probably yes: the in-sample ν is likely to be close to the unknown μ .

Note how little the two claims have to do with each other. The left column is about what a sample *can* be, and a bin of a million marbles of which 999,999 are orange can still deliver ten green ones. The right column is about what a sample *tends* to be, and puts a number on ‘tends’. Nothing is being hidden: the bad sample really does exist, we are simply going to bound how often it turns up.

Question 4.2.3. Formally, what does ν say about μ ?

§4.2.iv. Hoeffding’s Inequality

Theorem 4.2.4 (Hoeffding’s inequality)

Draw N marbles independently from a bin whose orange probability is μ , and let ν be the orange fraction among them. For every $\epsilon > 0$,

$$\mathbb{P}[|\nu - \mu| > \epsilon] \leq 2 \exp(-2\epsilon^2 N).$$

The statement is about marbles, coins, polling and anything else that can be sampled, and the three properties that make it usable are worth separating out.

- It is valid for all N and all ϵ — no ‘ N large enough’ clause, no asymptotics.
- It does not depend on μ . We never need to ‘know’ μ , which is the whole point, since μ is what we are trying to infer.
- A larger sample size N , or a looser gap ϵ , gives a higher probability for ‘ $\nu \approx \mu$ ’.

Definition 4.2.5 (PAC). A statement is **probably approximately correct** (PAC) when it may be off by a little (*approximately*, the ϵ) and may fail outright now and then (*probably*, the bound on the right). The statement ‘ $\nu = \mu$ ’ is PAC.

Remark 4.2.6. The factor 2 is a union bound over the two ways of being wrong. Hoeffding proves $\mathbb{P}[\nu - \mu > \epsilon] \leq \exp(-2\epsilon^2 N)$ and, by the same argument applied to green, $\mathbb{P}[\mu - \nu > \epsilon] \leq \exp(-2\epsilon^2 N)$; adding the two gives the two-sided form. The same cheap trick — bound each bad event, then add — reappears in [Section 4.4](#), where it is the only tool used.

Note 4.2.7. The exponent decays in N , so the bound is far better than its shape suggests. At $\epsilon = 0.1$ and $N = 1,000$ it reads $2 \exp(-20) \approx 4 \times 10^{-9}$: a thousand marbles pin down the contents of an arbitrarily large bin to within ten percentage points, and the chance of being fooled is one in a quarter of a billion.

If N is large, we can probably infer the unknown μ from the known ν .

Exercise 4.2.8 (Fun Time). Let $\mu = 0.4$. Use Hoeffding's inequality to bound the probability that a sample of 10 marbles will have $\nu \leq 0.1$. What bound do you get?

- (1) 0.67;
- (2) 0.40;
- (3) 0.33;
- (4) 0.05.

Reference Answer. (3). The event $\nu \leq 0.1$ is contained in $|\nu - \mu| \geq 0.3$, so set $N = 10$ and $\epsilon = 0.3$:

$$2 \exp(-2(0.3)^2(10)) = 2e^{-1.8} \approx 0.33.$$

The actual probability is about 0.05, which is answer (4); Hoeffding gives only an upper bound, and at $N = 10$ it is a loose one. ■

§4.3. Connection to Learning


— slides p.13–18 —

§4.3.i. The Bin Analogy

The bin is not an analogy chosen for decoration; every one of its pieces has a counterpart in learning, and the translation is exact.

bin	learning
unknown orange probability μ	a <i>fixed</i> hypothesis $h(\mathbf{x})$, asked whether it equals the target $f(\mathbf{x})$
a marble • in the bin	a point $\mathbf{x} \in \mathcal{X}$
orange marble	h is wrong : $h(\mathbf{x}) \neq f(\mathbf{x})$
green marble	h is right : $h(\mathbf{x}) = f(\mathbf{x})$
a size- N sample of i.i.d. marbles	checking h on $\mathcal{D} = \left\{ (\mathbf{x}_n, \underbrace{y_n}_{f(\mathbf{x}_n)}) \right\}$ with i.i.d. $\mathbf{x}_n$

So the input space itself is the bin, and it is h that colours it: paint $\mathbf{x}$ orange where h disagrees with f and green where it agrees.




Two conditions make the translation legal, and they are exactly the two that [Theorem 4.2.4](#) needs. The $\mathbf{x}_n$ must be drawn independently, and from the same distribution that will later produce the $\mathbf{x}$ we are judged on. Nothing is assumed about *which* distribution that is.

If N is large and the x_n are i.i.d., we can probably infer the unknown $\llbracket h(x) \neq f(x) \rrbracket$ probability from the known $\llbracket h(x_n) \neq y_n \rrbracket$ fraction.

§4.3.ii. Added Components

The learning flow of [Section 1.4](#) therefore gains one box it did not have: a probability distribution P on $\mathcal{X}$, which generates both the training inputs and the future ones.



The two quantities the bin analogy named now have proper symbols.

Definition 4.3.1 (In-sample and out-of-sample error). For a hypothesis h , the **out-of-sample error**

$$E_{\text{out}}(h) = \mathcal{E}_{x \sim P} \llbracket h(x) \neq f(x) \rrbracket$$

is the probability that h is wrong on a fresh input, and the **in-sample error**

$$E_{\text{in}}(h) = \frac{1}{N} \sum_{n=1}^N \llbracket h(x_n) \neq y_n \rrbracket$$

is the fraction of the examples it gets wrong. In the bin picture $E_{\text{out}}(h)$ is μ and $E_{\text{in}}(h)$ is ν .

For any fixed h , we can probably infer the unknown $E_{\text{out}}(h)$ from the known $E_{\text{in}}(h)$.

§4.3.iii. The Formal Guarantee

Rewriting [Theorem 4.2.4](#) in this notation costs nothing and says everything.

Theorem 4.3.2 (Hoeffding for a fixed hypothesis)

Fix h before seeing $\mathcal{D}$, and let the x_n be drawn i.i.d. from P . Then for every $\epsilon > 0$,

$$\mathbb{P} [|E_{\text{in}}(h) - E_{\text{out}}(h)| > \epsilon] \leq 2 \exp(-2\epsilon^2 N).$$

The same three remarks as before apply, and the third one has teeth here.

- Valid for all N and ϵ .

- Does not depend on $E_{\text{out}}(h)$, so there is no need to ‘know’ $E_{\text{out}}(h)$: both f and P may stay unknown.
- ‘ $E_{\text{in}}(h) = E_{\text{out}}(h)$ ’ is probably approximately correct.

If ‘ $E_{\text{in}}(h) \approx E_{\text{out}}(h)$ ’ and ‘ $E_{\text{in}}(h)$ small’, then $E_{\text{out}}(h)$ is small, i.e. $h \approx f$ with respect to P .

That last line deserves to be read slowly, because it says what $g \approx f$ has meant all along. It is not ‘ g and f agree everywhere’ — no amount of data can buy that, by [Proposition 4.1.2](#). It is ‘ g and f agree on inputs that actually occur’, weighted by how often they occur. Regions of $\mathcal{X}$ that P never visits are free to disagree, and we never find out.

§4.3.iv. Verification of One h

So for any fixed h , when the data is large enough, $E_{\text{in}}(h) \approx E_{\text{out}}(h)$. Learning, however, never reports an h of our choosing: it reports $g = \mathcal{A}(\mathcal{D})$, the single hypothesis the algorithm settles on out of $\mathcal{H}$ after seeing $\mathcal{D}$.

Question 4.3.3. Can we then claim ‘good learning’, that is, $g \approx f$ for the $g = \mathcal{A}(\mathcal{D})$ the algorithm actually returns?

Yes!

If $E_{\text{in}}(h)$ is small for the fixed h , and $\mathcal{A}$ picks that h as g , then ‘ $g = f$ ’ is PAC.

No!


If $\mathcal{A}$ is *forced* to pick THE h as g , then $E_{\text{in}}(h)$ is almost always not small, and ‘ $g \neq f$ ’ is PAC!

Both columns describe the same arrangement — one h nominated before the data, returned as g — and differ only in whether $E_{\text{in}}(h)$ turns out small, which is not ours to arrange: an h written down before anyone looks at $\mathcal{D}$ has no reason to fit it. So $E_{\text{in}}(h)$ is almost always large, [Theorem 4.3.2](#) carries *that* over to $E_{\text{out}}(h)$ just as faithfully, and what gets PAC-certified is ‘ $g \neq f$ ’. The theorem does not fail us; it works perfectly and reports that we are reliably wrong. The obvious repair — read $\mathcal{D}$ first, then return whichever h fits it — is precisely what the premise ‘fix h before seeing $\mathcal{D}$ ’ rules out. Reading the data was never the problem, since $E_{\text{in}}(h)$ is computed from it; the problem is that an h chosen that way is a function of $\mathcal{D}$, so it names no one bin whose μ stood fixed while the sample was drawn, and [Theorem 4.3.2](#) has nothing to say about it.

Real learning: $\mathcal{A}$ shall make choices $\in \mathcal{H}$ (like PLA), rather than being forced to pick one h . :-)

§4.3.v. The ‘Verification’ Flow

What [Theorem 4.3.2](#) does buy is a complete and genuinely useful procedure, one step short of learning. Somebody hands you a formula; you check it against historical records; Hoeffding tells you how much the check is worth.



There is no $\mathcal{A}$ in that picture and no $\mathcal{H}$, only a single given h that becomes g unchanged.

We can now use ‘historical records’ (data) to verify ‘one candidate formula’ h .

Exercise 4.3.4 (Fun Time). Your friend tells you her secret rule in investing in a particular stock: ‘whenever the stock goes down in the morning, it will go up in the afternoon; vice versa.’ To verify the rule, you chose 100 days uniformly at random from the past 10 years of stock data, and found that 80 of them satisfy the rule. What is the best guarantee that you can get from the verification?

- (1) You’ll definitely be rich by exploiting the rule in the next 100 days.
- (2) You’ll likely be rich by exploiting the rule in the next 100 days, if the market behaves similarly to the last 10 years.
- (3) You’ll likely be rich by exploiting the ‘best rule’ from 20 more friends in the next 100 days.
- (4) You’d definitely have been rich if you had exploited the rule in the past 10 years.


Reference Answer. (2), and every word of the hedge is doing work. (1) is no free lunch. (3) picks the best of 21 rules after looking at the data, which verification says nothing about — that is the subject of [Section 4.4](#). (4) forgets that the verification used only 100 days out of the ten years, so the rule may well be wrong on most of the rest. The clause ‘if the market behaves similarly’ is the i.i.d. assumption: the next 100 days must be drawn from the same P as the 100 you checked. ■

§4.4. Connection to Real Learning

— slides p.19–25 —

§4.4.i. Multiple h


Real learning has M candidates, not one, so there are M bins — one per hypothesis, each with its own $E_{\text{out}}(h_m)$ — and one sample, which is scooped out of all of them at once, because the same $\mathcal{D}$ is used to compute every $E_{\text{in}}(h_m)$.



Question 4.4.1. Real learning, say like PLA: BINGO when getting all green marbles?

§4.4.ii. Coin Game

The temptation is to answer yes, and the following game shows why that is a mistake. It is the same game, with bins replaced by coins: a fair coin is a bin with $\mu = \frac{1}{2}$, and five flips are a sample of size $N = 5$.



Example 4.4.2 (Coin game). Everyone in a size-150 ML class flips a coin 5 times, and one of the students gets 5 heads for her coin ‘ g ’. Is ‘ g ’ really magical?

No. Even if all 150 coins are fair, the probability that *one of them* results in 5 heads is

$$1 - \left(\frac{31}{32}\right)^{150} > 99\%.$$

A single fair coin gives five heads with probability $\frac{1}{32}$, which is small; 150 of them give it with probability 99.1%, which is a near certainty. Nothing about any individual coin changed. What changed is that we were allowed to look at all 150 results and then name the coin — and the naming is the part Hoeffding says nothing about.

BAD sample: E_{in} and E_{out} far away — and it can get worse when ‘choice’ is involved.

§4.4.iii. BAD Sample and BAD Data

Definition 4.4.3 (BAD sample, BAD data). A **BAD sample** is one on which E_{in} and E_{out} are far apart: $E_{\text{out}} = \frac{1}{2}$, say, but all heads came up, so $E_{\text{in}} = 0$. A data set $\mathcal{D}$ is **BAD data for h** when $E_{\text{out}}(h)$ and $E_{\text{in}}(h)$ are far away — for instance $E_{\text{out}}(h)$ big, so h is far from f , while $E_{\text{in}}(h)$ is small, so h is correct on most of the examples.

Intuition . Theorem 4.3.2 holds two things still and lets one vary. Held still are the distribution P on $\mathcal{X}$ and the hypothesis h ; what varies is $\mathcal{D}$, drawn i.i.d. from P . In particular $E_{\text{out}}(h)$ does not move at all — it is an expectation over P , fixed by h , f and P alone —

so every bit of the randomness sits in $E_{\text{in}}(h)$. The theorem is therefore a statement about the draw: however $\mathcal{D}$ comes out, as long as N is large enough, $E_{\text{in}}(h)$ lands within ϵ of that one fixed $E_{\text{out}}(h)$ — for all but a $2 \exp(-2\epsilon^2 N)$ fraction of the possible $\mathcal{D}$. Those rare exceptions are exactly the BAD data for h .

That is the whole content of [Theorem 4.3.2](#), restated as a property of the data set rather than of a probability: Hoeffding says the BAD rows of the following table are rare.

	$\mathcal{D}_1$	$\mathcal{D}_2$	$\dots$	$\mathcal{D}_{1126}$	$\dots$	$\mathcal{D}_{5678}$	Hoeffding
h	BAD					BAD	$\mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D} \text{ for } h] \leq \dots$

Reading a probability as a sum over data sets makes the next step obvious:

$$\mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D}] = \sum_{\text{all possible } \mathcal{D}} \mathbb{P}(\mathcal{D}) \cdot \llbracket \text{BAD } \mathcal{D} \rrbracket,$$

and Hoeffding says this sum is small.

§4.4.iv. BAD Data for Many h

With M hypotheses there are M rows, each with its own sparse scattering of BAD entries, and the algorithm's freedom is exactly the freedom to avoid them. Say that $\mathcal{D}$ is *BAD for many h* when it is BAD for at least one of them; equivalently, when $\mathcal{A}$ has no 'freedom of choice' on $\mathcal{D}$, because there exists some h whose $E_{\text{out}}(h)$ and $E_{\text{in}}(h)$ are far away.

	$\mathcal{D}_1$	$\mathcal{D}_2$	$\dots$	$\mathcal{D}_{1126}$	$\dots$	$\mathcal{D}_{5678}$	Hoeffding
h_1	BAD					BAD	$\mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D} \text{ for } h_1] \leq \dots$
h_2		BAD					$\mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D} \text{ for } h_2] \leq \dots$
h_3	BAD	BAD				BAD	$\mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D} \text{ for } h_3] \leq \dots$
$\dots$							
h_M	BAD					BAD	$\mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D} \text{ for } h_M] \leq \dots$
all	BAD	BAD				BAD	?

Question 4.4.4. For M hypotheses, what bounds $\mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D}]$?

§4.4.v. Bound of BAD Data

Theorem 4.4.5 (Finite-bin Hoeffding)

Let $\mathcal{H} = \{h_1, \dots, h_M\}$ be finite and let g be whatever $\mathcal{A}$ returns. For every $\epsilon > 0$,

$$\mathbb{P} [|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \leq 2M \exp(-2\epsilon^2 N).$$

Proof. A data set on which g fails is a data set on which *some* h_m fails, whichever h_m the algorithm happened to return, so

$$\begin{aligned} \mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D}] &= \mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D} \text{ for } h_1 \text{ or } \dots \text{ or BAD } \mathcal{D} \text{ for } h_M] \\ &\leq \mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D} \text{ for } h_1] + \dots + \mathbb{P}_{\mathcal{D}} [\text{BAD } \mathcal{D} \text{ for } h_M] && \text{(union bound)} \\ &\leq 2 \exp(-2\epsilon^2 N) + \dots + 2 \exp(-2\epsilon^2 N) \\ &= 2M \exp(-2\epsilon^2 N), \end{aligned}$$

the third line by [Theorem 4.3.2](#) applied to each fixed h_m separately. ■

- The finite-bin version is valid for all M , N and ϵ .
- It does not depend on any $E_{\text{out}}(h_m)$, so there is no need to ‘know’ them: f and P can stay unknown.
- ‘ $E_{\text{in}}(g) = E_{\text{out}}(g)$ ’ is PAC, *regardless of $\mathcal{A}$* .


Remark 4.4.6. The union bound is crude — it adds M probabilities as if the bad events never overlapped, when in a typical $\mathcal{H}$ similar hypotheses fail on exactly the same data sets. Paying M for that crudeness is harmless while M is a few thousand and fatal when M is infinite. The next chapters buy the slack back.

The last bullet is what makes the bound useful, because it frees the algorithm completely: whatever rule $\mathcal{A}$ uses to choose, the chosen g is covered. And once every hypothesis is covered, the sensible thing to do is the obvious thing.

The ‘most reasonable’ $\mathcal{A}$, like PLA or pocket: pick the h_m with the lowest $E_{\text{in}}(h_m)$ as g .

§4.4.vi. The ‘Statistical’ Learning Flow

Putting the two halves together gives the statement this chapter was after. If $|\mathcal{H}| = M$ is finite and N is large enough, then for whatever g is picked by $\mathcal{A}$ we have $E_{\text{out}}(g) \approx E_{\text{in}}(g)$; and if $\mathcal{A}$ finds one g with $E_{\text{in}}(g) \approx 0$, the PAC guarantee gives $E_{\text{out}}(g) \approx 0$. Learning is possible.



The catch is in the hypothesis of the theorem. $M = |\mathcal{H}|$ must be finite, and the one hypothesis set we own — the perceptrons of [Section 2.1](#) — is a continuum. [Theorem 4.4.5](#) with $M = \infty$ says nothing at all.

Exercise 4.4.7 (Fun Time). Consider 4 hypotheses,

$$h_1(\mathbf{x}) = \text{sign}(x_1), \quad h_2(\mathbf{x}) = \text{sign}(x_2), \quad h_3(\mathbf{x}) = \text{sign}(-x_1), \quad h_4(\mathbf{x}) = \text{sign}(-x_2).$$

For any N and ϵ , which of the following statements is *not* true?

- (1) the BAD data of h_1 and the BAD data of h_2 are exactly the same;
- (2) the BAD data of h_1 and the BAD data of h_3 are exactly the same;
- (3) $\mathbb{P}_{\mathcal{D}}[\text{BAD for some } h_k] \leq 8 \exp(-2\epsilon^2 N)$;

$$(4) \mathbb{P}_{\mathcal{D}} [\text{BAD for some } h_k] \leq 4 \exp(-2\epsilon^2 N).$$

Reference Answer. (1). Settle (2) first, because it decides everything else. Wherever $x_1 \neq 0$ we have $h_3(\mathbf{x}) = -h_1(\mathbf{x})$, and with $\mathcal{Y} = \{+1, -1\}$ that means h_3 is wrong exactly where h_1 is right:

$$E_{\text{in}}(h_3) = 1 - E_{\text{in}}(h_1), \quad E_{\text{out}}(h_3) = 1 - E_{\text{out}}(h_1).$$

Subtracting, $E_{\text{in}}(h_3) - E_{\text{out}}(h_3) = -(E_{\text{in}}(h_1) - E_{\text{out}}(h_1))$. The two gaps have the same absolute value, so $\mathcal{D}$ is BAD for h_3 exactly when it is BAD for h_1 : (2) is true, and the same argument pairs h_4 with h_2 .

The four rows of the BAD table are therefore only two distinct rows, and the union bound has two terms to add rather than four:

$$\mathbb{P}_{\mathcal{D}} [\text{BAD for some } h_k] = \mathbb{P}_{\mathcal{D}} [\text{BAD for } h_1 \text{ or BAD for } h_2] \leq 2 \cdot 2 \exp(-2\epsilon^2 N).$$

That is (4), and (3) is a weaker bound, so it holds as well. Only (1) is left, and it fails: h_1 and h_2 read different coordinates of $\mathbf{x}$, and nothing ties the examples one gets wrong to the examples the other gets wrong.

The moral is worth keeping for the next chapter. Counting hypotheses overcounts BAD data as soon as two hypotheses go wrong on the same data sets — here $M = 4$ while the honest count is 2 — and similar ideas will conquer the $M = \infty$ case. ■

Summary

Learning is impossible? Absolutely no free lunch outside $\mathcal{D}$.

Probability to the rescue Probably approximately correct outside $\mathcal{D}$.

Connection to learning Verification possible if $E_{\text{in}}(h)$ is small for a fixed h .

Connection to real learning Learning possible if $|\mathcal{H}|$ is finite and $E_{\text{in}}(g)$ is small.

The chapter's title keeps its question mark only for the first section. Learning is impossible if the target is arbitrary and the guarantee must be certain; it is possible as soon as we sample i.i.d. from a fixed P and accept a guarantee that is probabilistic, with E_{in} the quantity we minimise and E_{out} the quantity we are actually judged on.

What survives as a debt is the factor M . It is the only place where the size of $\mathcal{H}$ enters, and it is the reason a bigger hypothesis set is not free: more candidates make E_{in} easier to drive down and make the bound worse. The next chapter puts that trade-off at the centre — and starts by replacing M with something that stays finite even when $\mathcal{H}$ does not.

5. Training versus Testing

Lecture 5

23 Sep.

As previously seen . [Theorem 4.4.5](#) closed the last chapter: if $\mathcal{H} = \{h_1, \dots, h_M\}$ is finite and N is large enough, then

$$\mathbb{P}[|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \leq 2M \exp(-2\epsilon^2 N)$$

for whatever g the algorithm picks — so if $\mathcal{A}$ also finds a g with $E_{\text{in}}(g) \approx 0$, we get the PAC guarantee $E_{\text{out}}(g) \approx 0$ and learning is possible.


But the perceptrons of [Section 2.1](#) form an infinite set: $M = \infty$ and the bound says nothing. This chapter looks for a finite quantity to put in place of M . The plan is to find where M enters the proof — a union bound over one BAD event per hypothesis — observe that similar hypotheses have almost the same BAD event, and so count hypotheses only up to what N inputs can tell apart. That count is finite, and the last section asks when it grows only polynomially in N .

§5.1. Recap and Preview

— slides p.2–6 —

§5.1.i. The ‘Statistical’ Learning Flow

We start from where [Section 4.4](#) ended: the learning flow, now with its conclusion attached.



The two $\approx$ signs are achieved in different ways. The right-hand one, $E_{\text{in}}(g) \approx 0$, is about the training data: E_{in} is computed on $\mathcal{D}$, which we have, so the algorithm can measure it and push it down directly — that is what PLA and pocket do. The left-hand one, $E_{\text{out}}(g) \approx E_{\text{in}}(g)$, is about future test data: E_{out} depends on the unknown P , so it cannot be computed, and no work on $\mathcal{D}$ can check it. We can only rely on [Theorem 4.4.5](#), which says it holds with high probability when N is large and M is finite. Learning therefore splits into two separate tasks.

§5.1.ii. Two Central Questions

Each link of the chain has been supplied by an earlier chapter:

$$\underbrace{\text{for batch \& supervised binary classification,}}_{\text{Ch. 3}} \quad \underbrace{g \approx f \Leftrightarrow E_{\text{out}}(g) \approx 0,}_{\text{Ch. 1}}$$

$$\text{achieved through } \underbrace{E_{\text{out}}(g) \approx E_{\text{in}}(g)}_{\text{Ch. 4}} \quad \text{and} \quad \underbrace{E_{\text{in}}(g) \approx 0}_{\text{Ch. 2}}.$$

Chapter 3 fixed the setting and Chapter 1 the goal; Chapter 4 handled $E_{\text{out}} \approx E_{\text{in}}$ and Chapter 2 (PLA) handled $E_{\text{in}} \approx 0$. So learning comes down to two questions:

1. Can we make sure that $E_{\text{out}}(g)$ is close enough to $E_{\text{in}}(g)$?
2. Can we make $E_{\text{in}}(g)$ small enough?

Question 5.1.1. What role does $M = |\mathcal{H}|$ play for the two questions?

§5.1.iii. Trade-off on M

Recall . $\mathcal{D}$ is *BAD* (Definition 4.4.3) when it makes some hypothesis look better than it is — E_{in} and E_{out} far apart, the training error no guide to the real one. Theorem 4.4.5 bounds how often a draw of N examples does that:

$$\mathbb{P}_{\mathcal{D}}[\text{BAD } \mathcal{D}] = \mathbb{P}[|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \leq 2 \cdot M \cdot \exp(-2\epsilon^2 N),$$

so the first of the two questions is answered by making this quantity small.

M pulls the two questions in opposite directions:

small M	large M
1. Yes!, $\mathbb{P}[\text{BAD}] \leq 2 \cdot M \cdot \exp(\dots)$;	1. No!, $\mathbb{P}[\text{BAD}] \leq 2 \cdot M \cdot \exp(\dots)$;
2. No!, too few choices.	2. Yes!, many choices.

A small $\mathcal{H}$ gives a tight bound but may not contain a good fit; a large $\mathcal{H}$ fits easily but its bound is loose.

Using the right M (or $\mathcal{H}$) is important.

Question 5.1.2. Is $M = \infty$ doomed?

§5.1.iv. Preview

What we know is

$$\mathbb{P}[|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \leq 2 \cdot M \cdot \exp(-2\epsilon^2 N),$$

and the plan for the next three chapters is to keep this form while replacing the factor that becomes infinite:

- establish a finite quantity $m_{\mathcal{H}}$ that replaces M ,

$$\mathbb{P}[|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \stackrel{?}{\leq} 2 \cdot m_{\mathcal{H}} \cdot \exp(-2\epsilon^2 N); \quad (5.1)$$

- justify the feasibility of learning for infinite M ;
- study $m_{\mathcal{H}}$ to understand its trade-off for choosing the ‘right’ $\mathcal{H}$, just as we would for M .

This chapter only constructs $m_{\mathcal{H}}$ and computes it for a few hypothesis sets; the inequality itself, hence the ‘?’, is proved two chapters from now.

Note 5.1.3. The mysterious PLA is to be fully resolved after three more lectures, once the bound covers the infinite perceptron set.

Exercise 5.1.4 (Fun Time). *Data size: how large do we need?* One way to use the inequality

$$\mathbb{P}[|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \leq 2 \cdot \underbrace{M \cdot \exp(-2\epsilon^2 N)}_{\delta}$$

is to pick a tolerable difference ϵ as well as a tolerable BAD probability δ , and then gather data with size N large enough to achieve those tolerance criteria. Let $\epsilon = 0.1$, $\delta = 0.05$ and $M = 100$. What is the data size needed? (1) 215; (2) 415; (3) 615; (4) 815.

Reference Answer. (2). We can simply express N as a function of those ‘known’ variables: $2M \exp(-2\epsilon^2 N) \leq \delta$ rearranges to

$$N \geq \frac{1}{2\epsilon^2} \ln \frac{2M}{\delta} = \frac{1}{2(0.1)^2} \ln \frac{200}{0.05} = 50 \ln 4000 \approx 414.7,$$

so $N = 415$ is the first of the four that works. ■

§5.2. Effective Number of Lines

— slides p.7–15 —

To replace M , we first look at where it came from.

§5.2.i. Where Did M Come From?

Recall the proof of $\mathbb{P}[|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \leq 2M \exp(-2\epsilon^2 N)$. For each hypothesis there is a BAD event

$$\mathcal{B}_m: |E_{\text{in}}(h_m) - E_{\text{out}}(h_m)| > \epsilon,$$

and because $\mathcal{A}$ must be free to return any h_m it likes, the quantity to bound is

$$\mathbb{P}[\mathcal{B}_1 \text{ or } \mathcal{B}_2 \text{ or } \dots \text{ or } \mathcal{B}_M].$$

We bounded it by assuming the worst case — that all $\mathcal{B}_m$ are non-overlapping:

$$\mathbb{P}[\mathcal{B}_1 \text{ or } \mathcal{B}_2 \text{ or } \dots \text{ or } \mathcal{B}_M] \underbrace{\leq}_{\text{union bound}} \mathbb{P}[\mathcal{B}_1] + \mathbb{P}[\mathcal{B}_2] + \dots + \mathbb{P}[\mathcal{B}_M].$$

Each summand is at most $2 \exp(-2\epsilon^2 N)$ by Hoeffding, so M is simply the number of terms in this sum.


Question 5.2.1. What did the union bound fail to consider when $M = \infty$?

§5.2.ii. Where Did the Union Bound Fail?

The union bound is tight only when the events are disjoint. For similar hypotheses they are not: take $h_1 \approx h_2$, two lines a hair apart. Then

1. $E_{\text{out}}(h_1) \approx E_{\text{out}}(h_2)$, because the two hypotheses disagree on a set of tiny probability; and
2. for most $\mathcal{D}$, $E_{\text{in}}(h_1) = E_{\text{in}}(h_2)$, because the finitely many sample points almost never fall in the sliver where the two disagree.

Hence $\mathcal{B}_1$ and $\mathcal{B}_2$ are almost the same event, and the union bound counts their overlap twice.




With infinitely many lines this over-counting is infinite, which is exactly why the bound became useless.

Question 5.2.2. To account for the overlap, can we group similar hypotheses by *kind*?


§5.2.iii. How Many Lines Are There?

To group lines by kind, take $\mathcal{H} = \{\text{all lines in } \mathbb{R}^2\}$ (each line calls one side $\circ$ and the other $\times$). There are infinitely many lines, but on a few fixed inputs they can only produce finitely many labellings. We count those, starting from one input.



The arrow on each line points to its positive side. h_1 puts x_1 on the positive side and h_2 does not, so on one input there are two kinds of line: h_1 -like ($x_1 = \circ$) or h_2 -like ($x_1 = \times$).

With two inputs:




h_1 leaves both points on its positive side and h_2 leaves both on its negative side; the horizontal line between them, read upwards as h_3 and downwards as h_4 , splits them both ways. That gives four kinds, all 2^2 labellings.


Question 5.2.3. One input: 2; two inputs: 4. Three inputs?

§5.2.iv. Three Inputs, and Why the Answer Is Not One Number

For three inputs forming a triangle, all 2^3 labellings are achievable, since a line can cut off any one vertex (or, read the other way, any two).



If the three inputs are collinear, two of the eight are lost.




Observation . The two struck-out panels are the two that ask the middle point to differ from both of its neighbours. A line meets the segment x_1x_3 at most once, so it can never put x_2 alone on one side of it. ‘Fewer than 8’ is what degeneracy — collinear inputs, or repeated ones — costs.

So the count depends on where the inputs are: 8 or 6. Since the bound must hold for every data set, we will use the largest count, 8.

§5.2.v. Four Inputs

With four inputs the count falls below 2^N for the first time. There are $2^4 = 16$ labellings in theory, but the figure draws only the eight with $x_1 = \circ$: reading a line from its other side swaps every $\circ$ and $\times$, so labellings come in complementary pairs and each panel stands for itself and its mirror image. That is the ‘ $2\times$ ’.



Only the struck-out panel (x_1, x_3 against x_2, x_4) and its mirror image are impossible, leaving 14. This is not special to the picture:

Proposition 5.2.4 (At most 14 through four inputs)

For any four inputs $\mathbf{x}_1, \dots, \mathbf{x}_4 \in \mathbb{R}^2$, lines realise at most 14 of the $2^4 = 16$ labellings, and 14 is attained.

Proof. Radon's theorem says that any four points of $\mathbb{R}^2$ can be partitioned into two groups whose convex hulls intersect. Fix such a partition and colour one group $\circ$ and the other $\times$. A line realising that labelling would separate the two groups, hence separate their convex hulls — impossible, since the hulls share a point. The same partition with the colours exchanged is a second labelling ruled out for the same reason, and it is a different labelling because a group and its complement are never equal. So at least two of the 16 are missing.

The picture above attains the bound: the four points are in convex position in the cyclic order $\mathbf{x}_1, \mathbf{x}_4, \mathbf{x}_3, \mathbf{x}_2$, the intersecting hulls are the two diagonals $\mathbf{x}_1\mathbf{x}_3$ and $\mathbf{x}_2\mathbf{x}_4$, and every other labelling is cut off by a line. ■

§5.2.vi. Effective Number of Lines

Definition 5.2.5 (Effective number of lines). The **effective number of lines** for N inputs is the maximum number of kinds of lines with respect to any N inputs $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$, written $\text{effective}(N)$.

This is the grouping asked for in [Question 5.2.2](#): it is finite, and at most 2^N since each kind is a labelling of N inputs. So far:

lines in 2D	
N	$\text{effective}(N)$
1	2
2	4
3	8
4	$14 < 2^4$

If the union bound could be taken over kinds instead of hypotheses, we would get


$$\mathbb{P}[|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \leq 2 \cdot \text{effective}(N) \cdot \exp(-2\epsilon^2 N).$$

If ① $\text{effective}(N)$ can replace M and ② $\text{effective}(N) \ll 2^N$, then learning is possible with infinitely many lines. :-)

Intuition . Condition ② is needed because $\exp(-2\epsilon^2 N)$ decays exponentially: a factor like 2^N could cancel it, while a polynomial factor cannot, and the bound still goes to 0.

Exercise 5.2.6 (Fun Time). What is the effective number of lines for five inputs in $\mathbb{R}^2$? (1) 14; (2) 16; (3) 22; (4) 32.

Reference Answer. (3). If you put the inputs roughly around a circle,



you can then pick any *consecutive* run of inputs to be on one side of the line, and the rest on the other. There are $5 \cdot 4 = 20$ proper runs — five starting points times four lengths — plus the two labellings that use no line at all, all $\circ$ and all $\times$, for a total of 22 kinds, much smaller than $2^5 = 32$. You will find it difficult to generate more kinds by varying the inputs, and a formal proof comes in a later chapter. ■

§5.3. Effective Number of Hypotheses

— slides p.16–22 —

Nothing in [Section 5.2](#) used the geometry of lines, so the same counting applies to any $\mathcal{H}$. This section gives it general names and computes it for a few examples.

§5.3.i. Dichotomies: Mini-hypotheses

Let $\mathcal{H}$ be any set of hypotheses $h: \mathcal{X} \rightarrow \{\times, \circ\}$.

Definition 5.3.1 (Dichotomy). For fixed inputs $\mathbf{x}_1, \dots, \mathbf{x}_N$, write

$$h(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N) = (h(\mathbf{x}_1), h(\mathbf{x}_2), \dots, h(\mathbf{x}_N)) \in \{\times, \circ\}^N,$$

and call it a **dichotomy**: the hypothesis ‘limited’ to the eyes of the N inputs. Write

$$\mathcal{H}(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N)$$

for the set of all dichotomies implemented by $\mathcal{H}$ on them.

Notation . A dichotomy is written as a string, one symbol per input in index order: $\circ \circ \times \circ$ is the dichotomy that answers $\circ$ on $\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_4$ and $\times$ on $\mathbf{x}_3$. The panels of [Section 5.2](#) are exactly such strings, drawn.

A dichotomy is what is left of h when only N inputs are observed, and ‘grouping hypotheses by kind’ now means grouping them by dichotomy. This is the right grouping for the union bound because of the following.

Lemma 5.3.2 (Same dichotomy, same E_{in})

If h and h' implement the same dichotomy on the inputs of $\mathcal{D}$, then $E_{\text{in}}(h) = E_{\text{in}}(h')$ — not approximately, exactly.

Proof. $E_{\text{in}}(h) = \frac{1}{N} \sum_{n=1}^N \mathbb{I}[h(\mathbf{x}_n) \neq y_n]$ by [Definition 4.3.1](#), and by hypothesis $h(\mathbf{x}_n) = h'(\mathbf{x}_n)$ for every n , so the two sums have the same terms. ■

Hypotheses with the same dichotomy thus have the same E_{in} and, being similar, nearly the same E_{out} , so their BAD events nearly coincide — the overlap of [Question 5.2.1](#). Taking one term per dichotomy instead of one per hypothesis reduces the number of terms in the union bound from $|\mathcal{H}|$ to $|\mathcal{H}(\mathbf{x}_1, \dots, \mathbf{x}_N)| \leq 2^N$.

	hypotheses $\mathcal{H}$	dichotomies $\mathcal{H}(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N)$
e.g.	all lines in $\mathbb{R}^2$	$\{\circ \circ \circ \circ, \circ \circ \circ \times, \circ \circ \times \times, \dots\}$
size	possibly infinite	upper bounded by 2^N

$|\mathcal{H}(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N)|$ is our candidate for replacing M .

Note 5.3.3. This is a motivation, not a proof, and two gaps remain. First, $\mathcal{H}(\mathbf{x}_1, \dots, \mathbf{x}_N)$ depends on the random inputs; the next subsection removes this dependence. Second, a dichotomy fixes E_{in} exactly but E_{out} only approximately, so hypotheses in a group do not share exactly one BAD event. This second gap is the ‘?’ in Equation 5.1, closed two chapters from now.

§5.3.ii. Growth Function

The first gap is easy to close. $|\mathcal{H}(\mathbf{x}_1, \dots, \mathbf{x}_N)|$ depends on the inputs (8 or 6 for three points), so take the worst case:

Definition 5.3.4 (Growth function). The **growth function** of $\mathcal{H}$ is

$$m_{\mathcal{H}}(N) = \max_{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N \in \mathcal{X}} |\mathcal{H}(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N)|.$$

It is at most 2^N , finite however large $\mathcal{H}$ is, and for lines in the plane it is exactly $\text{effective}(N)$:


lines in 2D	
N	$m_{\mathcal{H}}(N)$
1	2
2	4
3	$\max(\dots, 6, 8) = 8$
4	$14 < 2^N$

Question 5.3.5. How does one actually ‘calculate’ the growth function?

For perceptrons we cannot yet, so we first compute it for three simpler hypothesis sets.

§5.3.iii. Growth Function for Positive Rays

Let $\mathcal{X} = \mathbb{R}$ be one dimensional and let $\mathcal{H}$ contain every $h(x) = \text{sign}(x - a)$ for a threshold $a \in \mathbb{R}$ — the ‘positive half’ of the 1D perceptrons.



Proposition 5.3.6 (Positive rays)

$$m_{\mathcal{H}}(N) = N + 1.$$

Proof. Sort the inputs as $x_1 < x_2 < \dots < x_N$. The dichotomy produced by threshold a depends only on how many inputs lie below a , so all thresholds inside one gap (x_n, x_{n+1}) give the same dichotomy, and different gaps give different ones. There are $N + 1$ gaps — the $N - 1$ internal ones together with $(-\infty, x_1)$ and (x_N, ∞) — hence $N + 1$ dichotomies, and no choice of inputs does better, since repeated inputs only merge gaps. ■


For $N = 4$, the 5 dichotomies:

x_1	x_2	x_3	x_4
○	○	○	○
×	○	○	○
×	×	○	○
×	×	×	○
×	×	×	×

$(N + 1) \ll 2^N$ when N is large.

§5.3.iv. Growth Function for Positive Intervals

Still with $\mathcal{X} = \mathbb{R}$, let $\mathcal{H}$ contain every h with $h(x) = +1$ if and only if $x \in [\ell, r)$, and -1 otherwise.



Proposition 5.3.7 (Positive intervals)

$$m_{\mathcal{H}}(N) = \binom{N+1}{2} + 1 = \frac{1}{2}N^2 + \frac{1}{2}N + 1.$$

Proof. As before, only the gaps matter: a dichotomy is decided by which of the $N + 1$ spots holds ℓ and which holds r . Choosing two *different* spots for the two ends, with ℓ in the earlier one, gives $\binom{N+1}{2}$ distinct dichotomies, each with a non-empty run of ○. Putting both ends in the same spot makes the interval miss every input and produces the all-× dichotomy, one further case however the spot is chosen. Adding the two counts gives $\binom{N+1}{2} + 1$. ■


For $N = 4$, the $10 + 1 = 11$ dichotomies:

x_1	x_2	x_3	x_4	x_1	x_2	x_3	x_4
○	×	×	×	×	○	○	○
○	○	×	×	×	×	○	×
○	○	○	×	×	×	○	○
○	○	○	○	×	×	×	○
×	○	×	×	×	×	×	×
×	○	○	×	×	×	×	×


$(\frac{1}{2}N^2 + \frac{1}{2}N + 1) \ll 2^N$ when N is large.

§5.3.v. Growth Function for Convex Sets

The last example reaches the ceiling 2^N . Let $\mathcal{X} = \mathbb{R}^2$ and let $\mathcal{H}$ contain every h with $h(\mathbf{x}) = +1$ if and only if $\mathbf{x}$ lies in a convex region, and -1 otherwise. Blue marks the region a hypothesis calls $+1$; the red surround is everything it calls -1 .




convex region in blue



non-convex region

Question 5.3.8. What is $m_{\mathcal{H}}(N)$ for convex sets?

Since $m_{\mathcal{H}}$ takes the best inputs, put $\mathbf{x}_1, \dots, \mathbf{x}_N$ on a big circle:



Proposition 5.3.9 (Convex sets)

$$m_{\mathcal{H}}(N) = 2^N.$$

Proof. Place the N inputs on a circle and fix any dichotomy. Take the convex hull of the inputs labelled $+1$ and extend it very slightly. Every positive input is inside by construction. Every negative input is outside, because a point of a circle is a vertex of the convex hull of any subset of that circle containing it — so it cannot lie in the hull of a subset that excludes it. The region is convex, so it is a legitimate h , and it implements the chosen dichotomy. All 2^N dichotomies are therefore implemented by these inputs, and 2^N is also the ceiling. ■

Definition 5.3.10 (Shattering). When $\mathcal{H}$ implements every one of the 2^N dichotomies on $\mathbf{x}_1, \dots, \mathbf{x}_N$, those N inputs are **shattered** by $\mathcal{H}$.

$m_{\mathcal{H}}(N) = 2^N$ if and only if there exist N inputs that can be shattered.

Note the quantifiers: $m_{\mathcal{H}}(N) = 2^N$ needs only *some* N inputs to be shattered, whereas $m_{\mathcal{H}}(N) < 2^N$ says that *no* N inputs are.

Exercise 5.3.11 (Fun Time). Consider positive *and* negative rays as $\mathcal{H}$, which is equivalent to the perceptron hypothesis set in 1D. The hypothesis set is often called a ‘decision stump’, to describe

the shape of its hypotheses. What is the growth function $m_{\mathcal{H}}(N)$? (1) N ; (2) $N + 1$; (3) $2N$; (4) 2^N .

Reference Answer. (3). Two dichotomies when the threshold falls in each of the $N - 1$ ‘internal’ spots, one for each reading of the sign, giving $2(N - 1)$; then two more for the all- $\circ$ and all- $\times$ cases, which the two outside spots produce whichever way they are read. In total $2(N - 1) + 2 = 2N$. ■

§5.4. Break Point

— slides p.23–26 —

§5.4.i. The Four Growth Functions

Collecting the growth functions computed so far:

- positive rays: $m_{\mathcal{H}}(N) = N + 1$;
- positive intervals: $m_{\mathcal{H}}(N) = \frac{1}{2}N^2 + \frac{1}{2}N + 1$;
- convex sets: $m_{\mathcal{H}}(N) = 2^N$;
- 2D perceptrons: $m_{\mathcal{H}}(N) < 2^N$ in some cases.

Suppose Equation 5.1 were granted and $m_{\mathcal{H}}(N)$ took the place of M :

$$\mathbb{P}[|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \stackrel{?}{\leq} 2 \cdot m_{\mathcal{H}}(N) \cdot \exp(-2\epsilon^2 N).$$

Since $\exp(-2\epsilon^2 N)$ decays exponentially, a polynomial $m_{\mathcal{H}}(N)$ still lets the bound go to 0 (rays, intervals), while 2^N does not (convex sets).

Polynomial: good; exponential: bad.

Question 5.4.1. For 2D, or general, perceptrons — is $m_{\mathcal{H}}(N)$ polynomial?

§5.4.ii. Break Point of $\mathcal{H}$

We do not know $m_{\mathcal{H}}(N)$ for 2D perceptrons, but we know that some three inputs can be shattered while *no* four inputs can (Proposition 5.2.4).



The second fact is the useful one, and we name it.

Definition 5.4.2 (Break point). If no k inputs can be shattered by $\mathcal{H}$, then k is a **break point** for $\mathcal{H}$.

Immediately:

- $m_{\mathcal{H}}(k) < 2^k$, by Definition 5.3.10.

- $k + 1, k + 2, k + 3, \dots$ are break points too.
- So the informative quantity is the *minimum* break point, and it is the one we shall study.

Proof of the second bullet. Suppose some $k + 1$ inputs were shattered. Delete one of them and restrict every dichotomy to the remaining k : each of the 2^k dichotomies on those k inputs is the restriction of at least one of the 2^{k+1} dichotomies above it, so the k inputs are shattered as well, contradicting the assumption that k is a break point. Induction gives every larger k . ■

2D perceptrons have a break point at 4.

§5.4.iii. The Four Break Points

The four hypothesis sets with their break points:

$\mathcal{H}$	$m_{\mathcal{H}}(N)$	minimum break point
positive rays	$N + 1 = O(N)$	2
positive intervals	$\frac{1}{2}N^2 + \frac{1}{2}N + 1 = O(N^2)$	3
convex sets	2^N	none
2D perceptrons	$< 2^N$ in some cases	4

(Rays: $m_{\mathcal{H}}(2) = 3 < 4$; intervals: $m_{\mathcal{H}}(3) = 7 < 8$.) Break point 2 comes with $O(N)$, break point 3 with $O(N^2)$, and no break point with 2^N . This suggests:

Conjecture 5.4.1. No break point: $m_{\mathcal{H}}(N) = 2^N$ — which is certain, being the contrapositive of [Definition 5.4.2](#). Break point k : $m_{\mathcal{H}}(N) = O(N^{k-1})$.

If true, a break point alone would guarantee polynomial growth, and hence, via [Equation 5.1](#), feasible learning for 2D perceptrons. Proving it is the next chapter.

Exercise 5.4.3 (Fun Time). Consider positive and negative rays as $\mathcal{H}$, the 1D perceptron set, whose growth function is $m_{\mathcal{H}}(N) = 2N$. What is the minimum break point for $\mathcal{H}$? (1) 1; (2) 2; (3) 3; (4) 4.

Reference Answer. (3). At $k = 3$ we get $m_{\mathcal{H}}(k) = 6$ while $2^k = 8$, so 3 is a break point; and it is minimal because $m_{\mathcal{H}}(2) = 4 = 2^2$ and $m_{\mathcal{H}}(1) = 2 = 2^1$ leave nothing broken below it. Note that [Conjecture 5.4.1](#) predicts $O(N^2)$ here while the truth is $2N$ — the conjecture is an upper bound, not an estimate. ■

Summary

Recap and Preview Two questions: $E_{\text{out}}(g) \approx E_{\text{in}}(g)$, and $E_{\text{in}}(g) \approx 0$.

Effective Number of Lines At most 14 through the eye of 4 inputs.

Effective Number of Hypotheses At most $m_{\mathcal{H}}(N)$ through the eye of N inputs.

Break Point When $m_{\mathcal{H}}(N)$ becomes ‘non-exponential’.

Training makes E_{in} small; the bound connecting E_{in} to E_{out} carried a factor M , and we have proposed replacing it by $m_{\mathcal{H}}(N)$, which counts dichotomies instead of hypotheses and is finite even for infinite $\mathcal{H}$.

Two things remain for the next two chapters: proving [Conjecture 5.4.1](#), that a break point k gives $m_{\mathcal{H}}(N) = O(N^{k-1})$; and removing the ‘?’ of [Equation 5.1](#) from

$$\mathbb{P}[|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \stackrel{?}{\leq} 2 \cdot m_{\mathcal{H}}(N) \cdot \exp(-2\epsilon^2 N),$$

Together they give a guarantee for the perceptron.

6. Theory of Generalization

Lecture 6

30 Sep.

As previously seen . Section 5.4 ended with the growth function $m_{\mathcal{H}}(N)$, the number of dichotomies $\mathcal{H}$ can implement on the best N inputs, as the (wishful) effective price of choice in training, and with Conjecture 5.4.1: a break point k should force $m_{\mathcal{H}}(N) = O(N^{k-1})$. Two things were left open — proving that conjecture, and removing the ‘?’ from

$$\mathbb{P} [|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \stackrel{?}{\leq} 2 \cdot m_{\mathcal{H}}(N) \cdot \exp(-2\epsilon^2 N).$$

This chapter does both. The first three sections prove the conjecture without ever looking at a particular $\mathcal{H}$: a break point is a purely combinatorial restriction on which tables of $\circ$ and $\times$ may occur, and counting the largest such table gives a polynomial. The last section sketches why $m_{\mathcal{H}}$, with a few changes of constant, really can stand in for M .

§6.1. Restriction of Break Point

— slides p.2–5 —

§6.1.i. The Four Break Points

The four hypothesis sets of the last chapter, each with the smallest input configuration it fails to shatter:

$\mathcal{H}$	$m_{\mathcal{H}}(N)$	inputs not shattered	break point
positive rays	$N + 1$	$\circ \times$	$m_{\mathcal{H}}(2) = 3 < 2^2$: break point at 2
positive intervals	$\frac{1}{2}N^2 + \frac{1}{2}N + 1$	$\circ \times \circ$	$m_{\mathcal{H}}(3) = 7 < 2^3$: break point at 3
convex sets	2^N	$\circ \times$ $\circ$	$m_{\mathcal{H}}(N) = 2^N$ always: no break point
2D perceptrons	$< 2^N$ in some cases	$\times \circ$ $\times \circ$ $\times \circ$	$m_{\mathcal{H}}(4) = 14 < 2^4$: break point at 4

The third column is the dichotomy that fails: a ray cannot say $\circ$ left of $\times$, an interval cannot leave a hole in the middle, a line cannot do XOR. Convex sets never fail, whatever the input.

Section 5.4 already showed one consequence of a break point: break point $k \Rightarrow$ break point $k + 1, \dots$. That concerns only the values $m_{\mathcal{H}}(N) < 2^N$ at $N \geq k$.

Question 6.1.1. What else does a break point imply?

§6.1.ii. What Must Be True When $k = 2$

Suppose the *minimum* break point is $k = 2$. Without knowing $\mathcal{H}$, a few values of $m_{\mathcal{H}}$ are forced:

- $N = 1$: every $m_{\mathcal{H}}(N) = 2$ by definition, since the minimum break point is above 1, so some single input is shattered;

- $N = 2$: every $m_{\mathcal{H}}(N) < 4$ by definition, since 2 is a break point (so the maximum possible is 3).

Question 6.1.2. What is the maximum possible $m_{\mathcal{H}}(N)$ when $N = 3$ and $k = 2$?

The only rule is that no two of the three inputs may be shattered: in the table of dichotomies, no pair of columns may show all four patterns $\circ\circ, \circ\times, \times\circ, \times\times$. So build the table row by row and reject any row that breaks the rule.

- $\circ\circ\circ$: 1 dichotomy; shatters any two points? No.
- add $\circ\circ\times$: 2 dichotomies; no.
- add $\circ\times\circ$: 3 dichotomies; no.
- add $\circ\times\times$: 4 dichotomies, but now columns $\mathbf{x}_2, \mathbf{x}_3$ read $\circ\circ, \circ\times, \times\circ, \times\times$ — shattered. Rejected.
- add $\times\circ\circ$ instead: 4 dichotomies; no pair shattered.
- a fifth row: each of the three remaining candidates shatters some pair (table below).

$\mathbf{x}_1$	$\mathbf{x}_2$	$\mathbf{x}_3$	
$\circ$	$\circ$	$\circ$	kept
$\circ$	$\circ$	$\times$	kept
$\circ$	$\times$	$\circ$	kept
$\times$	$\circ$	$\circ$	kept
$\circ$	$\times$	$\times$	4th row tried: shatters $\mathbf{x}_2, \mathbf{x}_3$
$\times$	$\circ$	$\times$	5th row tried: shatters $\mathbf{x}_1, \mathbf{x}_3$
$\times$	$\times$	$\circ$	5th row tried: shatters $\mathbf{x}_1, \mathbf{x}_2$
$\times$	$\times$	$\times$	5th row tried: shatters $\mathbf{x}_1, \mathbf{x}_2$ (every pair, in fact)

The four kept rows use up the table: $\circ\times\times$ was already refused, and $\times\circ\times, \times\times\circ, \times\times\times$ each complete the four patterns on some pair, so no column pair has room left. The maximum possible so far is 4 dichotomies. The search here started from $\circ\circ\circ$ and so is not a proof that no other table of five exists; that proof comes for free from [Theorem 6.3.3](#) below, which gives $B(3, 2) \leq 4$.

§6.1.iii. Restriction of Break Point

Putting the three cases together, when the minimum break point is $k = 2$:

- $N = 1$: every $m_{\mathcal{H}}(N) = 2$ by definition;
- $N = 2$: every $m_{\mathcal{H}}(N) < 4$ by definition (so maximum possible = 3);
- $N = 3$: maximum possible = $4 \ll 2^3$.

At $N = 3$ the break point already bites twice as hard as the definition asks: the definition only forbids shattering two points, yet it caps three points at half of 2^3 .

A break point k restricts the maximum possible $m_{\mathcal{H}}(N)$ a lot for $N > k$.

This is the answer to [Question 6.1.1](#), and it suggests how to prove [Conjecture 5.4.1](#) without knowing anything else about $\mathcal{H}$:

$$m_{\mathcal{H}}(N) \leq \text{maximum possible } m_{\mathcal{H}}(N) \text{ given } k \leq \text{poly}(N). \quad (6.1)$$

The first inequality is trivial; the work is in the second.

Exercise 6.1.3 (Fun Time). When the minimum break point is $k = 1$, what is the maximum possible $m_{\mathcal{H}}(N)$ when $N = 3$? (1) 1; (2) 2; (3) 4; (4) 8.

Reference Answer. (1). Because $k = 1$, the hypothesis set cannot even shatter one point. Thus, every ‘column’ of the table cannot contain both $\circ$ and $\times$. Then, after including the first dichotomy, it is not possible to include any other different dichotomy:

x_1	x_2	x_3	
$\circ$	$\times$	$\circ$	kept
$\circ$	$\times$	$\times$	rejected: column x_3 now shattered

Any second row differs from the first somewhere, and that column then shows both symbols. Thus, the maximum possible $m_{\mathcal{H}}(N)$ is 1. ■

§6.2. Bounding Function: Basic Cases

— slides p.6–11 —

The middle term of [Equation 6.1](#) deserves a name.

§6.2.i. Bounding Function

Definition 6.2.1 (Bounding function). The **bounding function** $B(N, k)$ is the maximum possible $m_{\mathcal{H}}(N)$ when k is a break point.

Two remarks make this concrete.

- It is a *combinatorial quantity*: the maximum number of length- N vectors with entries in $\{\circ, \times\}$ such that no length- k subvector is shattered — that is, no choice of k coordinates shows all 2^k patterns.
- It is *irrelevant of the details of $\mathcal{H}$* . For example, $B(N, 3)$ bounds both positive intervals ($k = 3$) and 1D perceptrons ($k = 3$).

The first remark is why the second holds: the dichotomies of any $\mathcal{H}$ with break point k , on any N inputs, form such a set of vectors, so

$$m_{\mathcal{H}}(N) \leq B(N, k)$$

for every $\mathcal{H}$ with break point k . The price is looseness — B takes the worst table, whether or not any $\mathcal{H}$ realises it.

Question 6.2.2. New goal: is $B(N, k) \leq \text{poly}(N)$?

§6.2.ii. Table of Bounding Function

We fill in a table of $B(N, k)$, rows N , columns k . What is known so far is:

- $B(2, 2) = 3$ (maximum < 4);
- $B(3, 2) = 4$ (the ‘pictorial’ proof above).

Three whole regions of the table then follow from short arguments.

Proposition 6.2.3 (Boundary cases)

1. $B(N, 1) = 1$ for every N .
2. $B(N, k) = 2^N$ for $N < k$ — including all dichotomies does not violate the ‘breaking condition’.
3. $B(N, k) = 2^N - 1$ for $N = k$ — removing a single dichotomy satisfies the ‘breaking condition’.

Proof. (1) is the Fun Time above, for any N : if no single coordinate may show both symbols, all rows agree everywhere, so there is one row.

(2) With $N < k$ there are no k coordinates to choose, so the condition is empty and all 2^N vectors are allowed.

(3) With $N = k$ the only k coordinates are all of them. All 2^N vectors would shatter them, so at most $2^N - 1$; and deleting any one vector leaves the full coordinate set missing a pattern, so $2^N - 1$ is attained. ■

Note that (3) contains the first known value, $B(2, 2) = 2^2 - 1 = 3$. The four stages of the table, merged into one — **blue**: column $k = 1$; **black**: $N < k$; **orange**: $N = k$; **purple**: $B(3, 2)$, the one entry from the pictorial search:

		k						
$B(N, k)$		1	2	3	4	5	6	...
N	1	1	2	2	2	2	2	...
	2	1	3	4	4	4	4	...
	3	1	4	7	8	8	8	...
	4	1			15	16	16	...
	5	1				31	32	...
	6	1					63	...
	⋮	⋮						⋱

(The diagonal’s first entry $B(1, 1) = 2^1 - 1 = 1$ belongs to both the blue column and the orange diagonal.) Everything on and above the diagonal is now known, and so is the first column; only the region $N > k \geq 2$ below the diagonal remains, which is where N grows with k fixed — the only part that matters for [Question 6.2.2](#).

More than halfway done! :-)

Exercise 6.2.4 (Fun Time). For the 2D perceptrons, which of the following claims is true?

- (1) minimum break point $k = 2$;
- (2) $m_{\mathcal{H}}(4) = 15$;
- (3) $m_{\mathcal{H}}(N) < B(N, k)$ when $N = k =$ minimum break point;
- (4) $m_{\mathcal{H}}(N) > B(N, k)$ when $N = k =$ minimum break point.

Reference Answer. (3). As discussed previously, the minimum break point for 2D perceptrons is 4, with $m_{\mathcal{H}}(4) = 14$ (Proposition 5.2.4). Also, note that $B(4, 4) = 15$. So the bounding function $B(N, k)$ can be ‘loose’ in bounding $m_{\mathcal{H}}(N)$: lines miss two of the sixteen dichotomies of four points (a Radon pair and its mirror), where the break-point rule alone forces only one to be missing. (4) is impossible, since B is by definition an upper bound. ■

§6.3. Bounding Function: Inductive Cases

— slides p.12–20 —

§6.3.i. Estimating $B(4, 3)$

The first unknown entry below the diagonal is $B(4, 3)$, marked ‘?’:

		k						
$B(N, k)$		1	2	3	4	5	6	...
N	1	1	2	2	2	2	2	...
	2	1	3	4	4	4	4	...
	3	1	4	7	8	8	8	...
	4	1		?	15	16	16	...
	5	1				31	32	...
	6	1					63	...
	$\vdots$	$\vdots$						$\ddots$

The motivation for how to get at it: a table on four inputs is a table on three inputs with one more column, so $B(4, 3)$ shall be related to $B(3, ?)$ — ‘adding’ one point to a $B(3, ?)$ table. The next step is to reduce $B(4, 3)$ to $B(3, ?)$.

§6.3.ii. ‘Achieving’ Dichotomies of $B(4, 3)$

First the answer, found the brute-force way: after checking all 2^{2^4} sets of dichotomies on four inputs — every subset of the 16 possible rows — the winner is a set of 11, and $B(4, 3) = 11$:

	x_1	x_2	x_3	x_4
01	○	○	○	○
02	×	○	○	○
03	○	×	○	○
04	○	○	×	○
05	○	○	○	×
06	×	×	○	×

	x_1	x_2	x_3	x_4
07	×	○	×	○
08	×	○	○	×
09	○	×	×	○
10	○	×	○	×
11	○	○	×	×

		k					
$B(N,k)$		1	2	3	4	5	6
N	1	1	2	2	2	2	2
	2	1	3	4	4	4	4
	3	1	4	7	8	8	8
	4	1		11	15	16	16
	5	1				31	32
	6	1					63

(I reran the search: no set of 12 avoids shattering three coordinates, and the table above passes.) The shaded cells are the hint for what comes next: the new entry 11 is to be explained by the row $N = 3$ above it.

Question 6.3.1. How to reduce $B(4, 3)$ to $B(3, ?)$ cases?

§6.3.iii. Reorganized Dichotomies of $B(4, 3)$

Sort the 11 rows by their first three entries, and look at what x_4 does. Some prefixes (x_1, x_2, x_3) occur twice, once with $x_4 = \circ$ and once with $x_4 = \times$; the others occur once.

	x_1	x_2	x_3	x_4
01	○	○	○	○
05	○	○	○	×
02	×	○	○	○
08	×	○	○	×
03	○	×	○	○
10	○	×	○	×
04	○	○	×	○
11	○	○	×	×
06	×	×	○	×
07	×	○	×	○
09	○	×	×	○

Orange: pair; purple: single. Write α for the number of pairs and β for the number of singles. Here $\alpha = 4$ and $\beta = 3$, and

$$B(4, 3) = 11 = 2\alpha + \beta.$$

The point of the split is that the two kinds of prefix face different restrictions, and each restriction is a smaller bounding function.

§6.3.iv. Estimating Part of $B(4, 3)$

Delete x_4 and merge each pair into one row. What remains is a table on (x_1, x_2, x_3) with $\alpha + \beta$ distinct rows:

	x_1	x_2	x_3
α	○	○	○
	×	○	○
	○	×	○
	○	○	×
β	×	×	○
	×	○	×
	○	×	×

	x_1	x_2	x_3
α	○	○	○
	×	○	○
	○	×	○
	○	○	×

Left: all $\alpha + \beta$ prefixes. Right: the α paired ones alone.

The whole prefix table (1/2).

- $\alpha + \beta$: dichotomies on (x_1, x_2, x_3) .
- $B(4, 3)$ ‘no shatter’ any 3 inputs $\Rightarrow \alpha + \beta$ ‘no shatter’ any 3.

The implication holds because deleting a column cannot create a pattern: if the prefix table showed all 8 patterns on x_1, x_2, x_3 , so would the original 11 rows. Hence

$$\alpha + \beta \leq B(3, 3).$$

The paired prefixes (2/2).

- α : dichotomies on (x_1, x_2, x_3) with x_4 *paired*.
- $B(4, 3)$ ‘no shatter’ any 3 inputs $\Rightarrow \alpha$ ‘no shatter’ any 2.

The claim is about the α rows alone, not the whole table: rows 01, 02, 03, 06 already shatter x_1, x_2 , and a break point at 3 says nothing about 2 in general. What makes the α rows special is that each comes with *both* values of x_4 . So if they shattered two inputs, say x_1, x_2 , the original table would show each of those four patterns with $x_4 = \circ$ and with $x_4 = \times$ — all eight patterns on x_1, x_2, x_4 , three inputs shattered. The pairing supplies the third input for free. Hence

$$\alpha \leq B(3, 2).$$

§6.3.v. Putting It All Together

The two estimates add up, since $2\alpha + \beta = (\alpha + \beta) + \alpha$:

$$\begin{aligned} B(4, 3) &= 2\alpha + \beta \\ \alpha + \beta &\leq B(3, 3) \\ \alpha &\leq B(3, 2) \\ \Rightarrow B(4, 3) &\leq B(3, 3) + B(3, 2). \end{aligned}$$

Nothing in the argument used 4 or 3 specifically. For any N and k , take a largest table for $B(N, k)$ and split it on the last input:

$$\begin{aligned} B(N, k) &= 2\alpha + \beta \\ \alpha + \beta &\leq B(N - 1, k) \\ \alpha &\leq B(N - 1, k - 1) \\ \Rightarrow B(N, k) &\leq B(N - 1, k) + B(N - 1, k - 1). \end{aligned}$$

Lemma 6.3.2 (Inductive formula)

For $N \geq 2$ and $k \geq 2$,

$$B(N, k) \leq B(N - 1, k) + B(N - 1, k - 1).$$

Proof. Let S be a set of $B(N, k)$ vectors in $\{\circ, \times\}^N$ that shatters no k coordinates. Group its vectors by their first $N - 1$ entries: α prefixes appear with both values of the last entry and β with one, so $|S| = 2\alpha + \beta$. The $\alpha + \beta$ distinct prefixes shatter no k of the first $N - 1$ coordinates, or S would shatter them too; so $\alpha + \beta \leq B(N - 1, k)$. The α paired prefixes shatter no $k - 1$ of those coordinates, or S would shatter those $k - 1$ together with the last one; so $\alpha \leq B(N - 1, k - 1)$. Add the two. ■

Running the formula down the table from the known entries fills the region below the diagonal with upper bounds (orange):

		k					
$B(N, k)$		1	2	3	4	5	6
N	1	1	2	2	2	2	2
	2	1	3	4	4	4	4
	3	1	4	7	8	8	8
	4	1	≤ 5	11	15	16	16
	5	1	≤ 6	≤ 16	≤ 26	31	32
	6	1	≤ 7	≤ 22	≤ 42	≤ 57	63

The shaded cells show the pattern: each entry is at most the one directly above plus the one above-left, as in Pascal's triangle. For instance, $B(4, 3) = 11 = 7 + 4$, so the formula is exact there, and $B(5, 3) \leq 11 + 5 = 16$.

We now have an upper bound of the bounding function.

§6.3.vi. Bounding Function: The Theorem

The Pascal-like recursion suggests binomial coefficients, and indeed:

Theorem 6.3.3 (Bounding function)

For all $N \geq 1$ and $k \geq 1$,

$$B(N, k) \leq \underbrace{\sum_{i=0}^{k-1} \binom{N}{i}}_{\text{highest term } N^{k-1}}.$$

Proof. Simple induction on N , using the boundary cases and the inductive formula. Write $\Phi(N, k)$ for the right-hand side.

Boundary. If $k = 1$, $B(N, 1) = 1 = \binom{N}{0}$ by Proposition 6.2.3. If $N = 1$ and $k \geq 2$, $B(1, k) = 2 = \binom{1}{0} + \binom{1}{1}$ (the sum stops at $\binom{1}{1}$, all later terms being 0).

Induction. For $N \geq 2$, $k \geq 2$, Lemma 6.3.2 and the hypothesis for $N - 1$ give

$$B(N, k) \leq \sum_{i=0}^{k-1} \binom{N-1}{i} + \sum_{i=0}^{k-2} \binom{N-1}{i} = \binom{N-1}{0} + \sum_{i=1}^{k-1} \left[\binom{N-1}{i} + \binom{N-1}{i-1} \right] = \sum_{i=0}^{k-1} \binom{N}{i},$$

the last step by Pascal's rule and $\binom{N-1}{0} = \binom{N}{0}$. ■

- The proof is a simple induction using the boundary and the inductive formula.
- For fixed k , $B(N, k)$ is upper bounded by $\text{poly}(N)$, the leading term being $\binom{N}{k-1} \approx N^{k-1}/(k-1)! \Rightarrow m_{\mathcal{H}}(N)$ is $\text{poly}(N)$ if a break point exists.

That settles Question 6.2.2 and, through $m_{\mathcal{H}}(N) \leq B(N, k)$, proves Conjecture 5.4.1.

Remark 6.3.4. ‘ $\leq$ ’ can be ‘ $=$ ’ actually — go play and prove it if you are a math lover! :-)
(One way: exhibit a table of the right size, e.g. all vectors with fewer than k entries $\times$. It has $\sum_{i < k} \binom{N}{i}$ rows and cannot shatter k coordinates, since it never shows $\times$ on all of them at once.)
The inequality is also known as the Sauer–Shelah lemma.

§6.3.vii. The Three Break Points

Returning to the hypothesis sets with a break point, [Theorem 6.3.3](#) bounds each growth function from its break point alone (in orange, the bound):

$\mathcal{H}$		$m_{\mathcal{H}}(N)$	break point
positive rays	$\circ \times$	$N + 1 \leq N + 1$	$m_{\mathcal{H}}(2) = 3 < 2^2$: at 2
positive intervals	$\circ \times \circ$	$\frac{1}{2}N^2 + \frac{1}{2}N + 1 \leq \frac{1}{2}N^2 + \frac{1}{2}N + 1$	$m_{\mathcal{H}}(3) = 7 < 2^3$: at 3
2D perceptrons	$\times \circ \times$ $\circ$	$m_{\mathcal{H}}(N) = ? \leq \frac{1}{6}N^3 + \frac{5}{6}N + 1$	$m_{\mathcal{H}}(4) = 14 < 2^4$: at 4

For rays and intervals the bound is exact: $\binom{N}{0} + \binom{N}{1} = N + 1$, and $\binom{N}{0} + \binom{N}{1} + \binom{N}{2} = \frac{1}{2}N^2 + \frac{1}{2}N + 1$. For 2D perceptrons we still do not know $m_{\mathcal{H}}(N)$, but with $k = 4$

$$\sum_{i=0}^3 \binom{N}{i} = 1 + N + \frac{N(N-1)}{2} + \frac{N(N-1)(N-2)}{6} = \frac{1}{6}N^3 + \frac{5}{6}N + 1,$$

a polynomial, found without computing a single dichotomy beyond $N = 4$.

We can bound $m_{\mathcal{H}}(N)$ by only one break point.

Exercise 6.3.5 (Fun Time). For 1D perceptrons (positive and negative rays), we know that $m_{\mathcal{H}}(N) = 2N$. Let k be the minimum break point. Which of the following is *not* true?

- (1) $k = 3$;
- (2) for some integers $N > 0$, $m_{\mathcal{H}}(N) = \sum_{i=0}^{k-1} \binom{N}{i}$;
- (3) for all integers $N > 0$, $m_{\mathcal{H}}(N) = \sum_{i=0}^{k-1} \binom{N}{i}$;
- (4) for all integers $N > 2$, $m_{\mathcal{H}}(N) < \sum_{i=0}^{k-1} \binom{N}{i}$.

Reference Answer. (3). The proof is generally trivial by listing the definitions. (1) is the last Fun Time of [Section 5.4](#). With $k = 3$ the sum is $1 + N + \frac{N(N-1)}{2}$, which equals 2, 4, 7 at $N = 1, 2, 3$: so for (2), $N = 1$ or 2 gives the equality, and $N = 3$ ($6 \neq 7$) breaks (3). One thing to notice is (4): the difference $1 + N + \frac{N(N-1)}{2} - 2N = \frac{(N-1)(N-2)}{2}$ is positive for $N > 2$, so the upper bound can be ‘loose’. ■

§6.4. A Pictorial Proof

— slides p.21–26 —

The growth function is now known to be polynomial whenever there is a break point. What remains is the ‘?’ of [Equation 5.1](#): may $m_{\mathcal{H}}$ take the place of M in the bound at all?

§6.4.i. BAD Bound for General $\mathcal{H}$

Because $\mathcal{A}$ may return any hypothesis, the event to control is that *some* h is BAD. What we want is

$$\mathbb{P}[\exists h \in \mathcal{H} \text{ s.t. } |E_{\text{in}}(h) - E_{\text{out}}(h)| > \epsilon] \leq 2 \cdot m_{\mathcal{H}}(N) \cdot \exp(-2\epsilon^2 N).$$

What is actually true, when N is large enough, is

$$\mathbb{P}[\exists h \in \mathcal{H} \text{ s.t. } |E_{\text{in}}(h) - E_{\text{out}}(h)| > \epsilon] \leq 2 \cdot 2m_{\mathcal{H}}(2N) \cdot \exp\left(-2 \cdot \frac{1}{16} \epsilon^2 N\right).$$


The wish survives with worse constants: $m_{\mathcal{H}}(2N)$ in place of $m_{\mathcal{H}}(N)$ (still polynomial), an extra factor 2, and $\epsilon^2/16$ in the exponent (still exponential decay). Next: a sketch of the proof, in three steps, each of which removes one obstacle and costs one of those constants.

§6.4.ii. Step 1: Replace E_{out} by E'_{in}

$$\frac{1}{2} \mathbb{P}[\exists h \in \mathcal{H} \text{ s.t. } |E_{\text{in}}(h) - E_{\text{out}}(h)| > \epsilon] \leq \mathbb{P}[\exists h \in \mathcal{H} \text{ s.t. } |E_{\text{in}}(h) - E'_{\text{in}}(h)| > \frac{\epsilon}{2}]$$

- $E_{\text{in}}(h)$ takes finitely many values across $\mathcal{H}$ (one per dichotomy), while $E_{\text{out}}(h)$ takes infinitely many — so replace the evil E_{out} first.
- How? Sample a **verification set** $\mathcal{D}'$ of size N to calculate E'_{in} , the error on $\mathcal{D}'$.
- BAD h of $E_{\text{in}} - E_{\text{out}}$ $\xrightarrow{\text{probably}}$ BAD h of $E_{\text{in}} - E'_{\text{in}}$.

$\mathcal{D}'$ is never actually collected; it exists only inside the proof.



The picture is for one fixed h . Both E_{in} and E'_{in} are averages of N independent errors, so they share the same bell-shaped distribution centred on E_{out} . Suppose $\mathcal{D}$ is BAD, with E_{in} out in the red tail, more than ϵ above E_{out} . The ghost sample is drawn independently, so E'_{in} lands in the green half, at or below E_{out} , with probability about $\frac{1}{2}$; and then E_{in} and E'_{in} are more than ϵ apart. So a BAD $\mathcal{D}$ produces a large gap $|E_{\text{in}} - E'_{\text{in}}|$ at least half the time, which is the $\frac{1}{2}$ on the left.

Remark 6.4.1. The rigorous version uses a margin instead of symmetry. On the event that $|E_{\text{in}}(h) - E_{\text{out}}(h)| > \epsilon$ for some h , fix one such h (it depends only on $\mathcal{D}$). By Hoeffding on $\mathcal{D}'$, $|E'_{\text{in}}(h) - E_{\text{out}}(h)| \leq \frac{\epsilon}{2}$ with probability at least $1 - 2\exp(-\frac{1}{2}\epsilon^2 N)$, which is at least $\frac{1}{2}$ once N is large enough — the ‘ N large enough’ in the statement. When it happens, the triangle inequality gives $|E_{\text{in}}(h) - E'_{\text{in}}(h)| > \epsilon - \frac{\epsilon}{2} = \frac{\epsilon}{2}$. The halved ϵ is the price of this margin.

The evil E_{out} is removed by verification with ‘ghost data’.


§6.4.iii. Step 2: Decompose $\mathcal{H}$ by Kind

$$\begin{aligned} \text{BAD} &\leq 2\mathbb{P}\left[\exists h \in \mathcal{H} \text{ s.t. } |E_{\text{in}}(h) - E'_{\text{in}}(h)| > \frac{\epsilon}{2}\right] \\ &\leq 2m_{\mathcal{H}}(2N)\mathbb{P}\left[\text{fixed } h \text{ s.t. } |E_{\text{in}}(h) - E'_{\text{in}}(h)| > \frac{\epsilon}{2}\right] \end{aligned}$$

- E_{in} is computed with $\mathcal{D}$ and E'_{in} with $\mathcal{D}'$ — and now $m_{\mathcal{H}}$ comes to play.

- How? Both errors depend on h only through its values on the $2N$ points of $\mathcal{D} \cup \mathcal{D}'$, so the infinite $\mathcal{H}$ becomes $|\mathcal{H}(\mathbf{x}_1, \dots, \mathbf{x}_N, \mathbf{x}'_1, \dots, \mathbf{x}'_N)|$ kinds, at most $m_{\mathcal{H}}(2N)$.
- Union bound on the $m_{\mathcal{H}}(2N)$ kinds.

The first line is Step 1 with both sides doubled; ‘BAD’ abbreviates the left-hand probability we started from. The price of this step is $m_{\mathcal{H}}(2N)$ rather than $m_{\mathcal{H}}(N)$: the dichotomies must be counted on the ghost points too.



Each box is the space of all possible data sets, and each coloured scribble the BAD region of one hypothesis — the data sets on which it is fooled. (a) Hoeffding says one hypothesis’s region is small, so a random $\mathcal{D}$ usually misses it. (b) The union bound adds up all the regions as if they were disjoint, so with many hypotheses they seem to fill the whole space. (c) In truth, hypotheses of the same kind are BAD on nearly the same data sets, and their regions pile on top of one another; counting kinds instead of hypotheses counts each pile once.

Use $m_{\mathcal{H}}(2N)$ to calculate the BAD-overlap properly.


§6.4.iv. Step 3: Use Hoeffding without Replacement

$$\begin{aligned} \text{BAD} &\leq 2m_{\mathcal{H}}(2N)\mathbb{P}\left[\text{fixed } h \text{ s.t. } |E_{\text{in}}(h) - E'_{\text{in}}(h)| > \frac{\epsilon}{2}\right] \\ &\leq 2m_{\mathcal{H}}(2N) \cdot 2 \exp\left(-2\left(\frac{\epsilon}{4}\right)^2 N\right) \end{aligned}$$

- Consider a bin of the $2N$ examples; choose N of them for E_{in} , and leave the others for E'_{in} . Since the average of the two halves is the bin’s own error rate,

$$|E_{\text{in}} - E'_{\text{in}}| > \frac{\epsilon}{2} \Leftrightarrow \left|E_{\text{in}} - \frac{E_{\text{in}} + E'_{\text{in}}}{2}\right| > \frac{\epsilon}{4}.$$

- So? It is just a ‘smaller bin’, a ‘smaller ϵ ’, and Hoeffding without replacement.



Everything is now conditioned on the $2N$ points in the bin, and h is fixed. Which N of them are $\mathcal{D}$ and which are $\mathcal{D}'$ is a uniformly random split, so E_{in} is the orange fraction of a sample of N drawn *without* replacement from a bin whose orange fraction is $\frac{E_{\text{in}} + E'_{\text{in}}}{2}$. Hoeffding's inequality holds for sampling without replacement as well (Hoeffding proved it in that form), with the tolerance now $\frac{\epsilon}{4}$: that gives $2 \exp(-2(\epsilon/4)^2 N)$.

Use Hoeffding after zooming to a fixed h .

§6.4.v. That's All!

Collecting the constants, $2m_{\mathcal{H}}(2N) \cdot 2 \exp\left(-2\frac{\epsilon^2}{16}N\right) = 4m_{\mathcal{H}}(2N) \exp\left(-\frac{1}{8}\epsilon^2 N\right)$.

Theorem 6.4.2 (Vapnik–Chervonenkis (VC) bound)

For N large enough,

$$\mathbb{P}[\exists h \in \mathcal{H} \text{ s.t. } |E_{\text{in}}(h) - E_{\text{out}}(h)| > \epsilon] \leq 4m_{\mathcal{H}}(2N) \exp\left(-\frac{1}{8}\epsilon^2 N\right).$$

Proof sketch.

- replace E_{out} by E'_{in} (Step 1);
- decompose $\mathcal{H}$ by kind (Step 2);
- use Hoeffding without replacement (Step 3). ■

Since the event covers every h , it covers the g returned by any $\mathcal{A}$, exactly as [Theorem 4.4.5](#) did for finite $\mathcal{H}$. Applied to 2D perceptrons:

- break point? 4;
- $m_{\mathcal{H}}(N)$? $O(N^3)$, by [Theorem 6.3.3](#), so $m_{\mathcal{H}}(2N)$ is $O(N^3)$ as well, and the bound goes to 0 as N grows.

Learning with 2D perceptrons is feasible! :-)

Exercise 6.4.3 (Fun Time). For positive rays, $m_{\mathcal{H}}(N) = N + 1$. Plug it into the VC bound for $\epsilon = 0.1$ and $N = 10000$. What is the VC bound of BAD events?

$$\mathbb{P}[\exists h \in \mathcal{H} \text{ s.t. } |E_{\text{in}}(h) - E_{\text{out}}(h)| > \epsilon] \leq 4m_{\mathcal{H}}(2N) \exp\left(-\frac{1}{8}\epsilon^2 N\right)$$

(1) 2.77×10^{-87} ; (2) 5.54×10^{-83} ; (3) 2.98×10^{-1} ; (4) 2.29×10^2 .

Reference Answer. (3). Simple calculation:

$$4 \cdot (2 \cdot 10000 + 1) \cdot \exp\left(-\frac{1}{8}(0.1)^2 \cdot 10000\right) = 80004 e^{-12.5} \approx 0.298.$$

Note that the BAD probability bound is not very small even with 10000 examples — the constants lost in the three steps are expensive. ■

Summary

Restriction of Break Point A break point ‘breaks’ consequent points.

Bounding Function: Basic Cases $B(N, k)$ bounds $m_{\mathcal{H}}(N)$ with break point k .

Bounding Function: Inductive Cases $B(N, k)$ is $\text{poly}(N)$.

A Pictorial Proof $m_{\mathcal{H}}(N)$ can replace M with a few changes.

The chapter closed both gaps left by the last one. A break point k caps $m_{\mathcal{H}}(N)$ by $\sum_{i < k} \binom{N}{i} = O(N^{k-1})$, and the VC bound puts $m_{\mathcal{H}}(2N)$ where M used to be; so any $\mathcal{H}$ with a break point generalises, however infinite it is. Next: how to ‘use’ the break point — the minimum break point turns out to be the one number that summarises an $\mathcal{H}$ ’s capacity.

7. The VC Dimension

Lecture 7

30 Sep.

As previously seen . [Theorem 6.3.3](#) showed that a break point k caps the growth function by a polynomial, $m_{\mathcal{H}}(N) \leq B(N, k) = \sum_{i \leq k} \binom{N}{i}$, and [Theorem 6.4.2](#) put $m_{\mathcal{H}}(2N)$ where M used to be. Together: $E_{\text{out}} \approx E_{\text{in}}$ is possible if $m_{\mathcal{H}}(N)$ breaks somewhere and N is large enough.

That conclusion depends on $\mathcal{H}$ only through one number, its minimum break point. This chapter gives that number a name (shifted by one), computes it for perceptrons in any dimension, explains what it measures physically, and then turns the VC bound around to read off what it says about model complexity and about how much data is needed.

§7.1. Definition of VC Dimension

— slides p.2–7 —

§7.1.i. Recap: More on Growth Function

The last chapter ended with

$$m_{\mathcal{H}}(N) \text{ of break point } k \leq B(N, k) = \underbrace{\sum_{i=0}^{k-1} \binom{N}{i}}_{\text{highest term } N^{k-1}}.$$

The sum is awkward to carry around. Its highest term suggests replacing it by the single power N^{k-1} , and the two tables side by side show when that is safe:

		k							k				
$B(N, k)$		1	2	3	4	5	N^{k-1}		1	2	3	4	5
N	1	1	2	2	2	2	N	1	1	1	1	1	1
	2	1	3	4	4	4		2	1	2	4	8	16
	3	1	4	7	8	8		3	1	3	9	27	81
	4	1	5	11	15	16		4	1	4	16	64	256
	5	1	6	16	26	31		5	1	5	25	125	625
	6	1	7	22	42	57		6	1	6	36	216	1296

Cell by cell, the shaded entries on the right are at least those on the left. Outside the shaded region the comparison fails: column $k = 2$ has $N + 1 > N$, and row $N = 1$ has $2 > 1$. Hence

Provably and loosely, for $N \geq 2, k \geq 3$,

$$m_{\mathcal{H}}(N) \leq B(N, k) = \sum_{i=0}^{k-1} \binom{N}{i} \leq N^{k-1}.$$

Remark 7.1.1. The table is evidence, not a proof. One proof: for $N \geq 2$, $k \geq 3$, induct on k . At $k = 3$, $1 + N + \frac{N(N-1)}{2} \leq N^2$ is $(N-2)(N+1) \geq 0$ after doubling and rearranging, which holds for all $N \geq 2$. Going from k to $k+1$ adds $\binom{N}{k} \leq \frac{N^k}{k!} \leq N^k - N^{k-1}$ whenever $N \geq 2$ and $k \geq 2$, since then $N^{k-1}(N-1) \geq N^k/2 \geq N^k/k!$.

§7.1.ii. Recap: More on Vapnik–Chervonenkis (VC) Bound

Chaining the VC bound with this estimate, for any $g = \mathcal{A}(\mathcal{D}) \in \mathcal{H}$ and a ‘statistical’ large $\mathcal{D}$, for $N \geq 2$, $k \geq 3$:

$$\begin{aligned} \mathbb{P}_{\mathcal{D}} [|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] &\leq \mathbb{P}_{\mathcal{D}} [\exists h \in \mathcal{H} \text{ s.t. } |E_{\text{in}}(h) - E_{\text{out}}(h)| > \epsilon] \\ &\leq 4m_{\mathcal{H}}(2N) \exp\left(-\frac{1}{8}\epsilon^2 N\right) \\ &\stackrel{\text{if } k \text{ exists}}{\leq} 4(2N)^{k-1} \exp\left(-\frac{1}{8}\epsilon^2 N\right). \end{aligned}$$

The first step holds because g is one of the h ; the last applies the moral above at $2N$ points. The subscript $\mathcal{D}$ is a reminder that the probability is over the draw of the data. Reading the chain as a recipe for learning:

- | | |
|--|-----------------------|
| if (1) $m_{\mathcal{H}}(N)$ breaks at k | (good $\mathcal{H}$) |
| (2) N large enough | (good $\mathcal{D}$) |
| $\Rightarrow$ probably generalized ‘ $E_{\text{out}} \approx E_{\text{in}}$ ’, and | |
| if (3) $\mathcal{A}$ picks a g with small E_{in} | (good $\mathcal{A}$) |
| $\Rightarrow$ probably learned! | (:-) good luck) |

The first two conditions are what the theory has been building; the third is an optimisation problem, and is the subject of *How can machines learn?*

§7.1.iii. VC Dimension

The chain carries the exponent $k-1$ everywhere, so it is natural to name that number directly: the formal name of the *maximum non-break* point.

Definition 7.1.2 (VC dimension). The **VC dimension** of $\mathcal{H}$, denoted $d_{\text{VC}}(\mathcal{H})$, is the

$$\text{largest } N \text{ for which } m_{\mathcal{H}}(N) = 2^N.$$

- It is the *most* inputs $\mathcal{H}$ can shatter.
- $d_{\text{VC}} = \text{‘minimum } k \text{’} - 1$.

If $m_{\mathcal{H}}(N) = 2^N$ for every N there is no break point and $d_{\text{VC}} = \infty$. The definition asks for *some* set of N inputs to be shattered, never all of them. Unwinding it,

$$\begin{aligned} N \leq d_{\text{VC}} &\Rightarrow \mathcal{H} \text{ can shatter some } N \text{ inputs,} \\ k > d_{\text{VC}} &\Rightarrow k \text{ is a break point for } \mathcal{H}. \end{aligned}$$

The second line is the one the bounds use: with $k = d_{\text{VC}} + 1$, the moral of the recap becomes

If $N \geq 2$, $d_{\text{VC}} \geq 2$, then $m_{\mathcal{H}}(N) \leq N^{d_{\text{VC}}}$.

§7.1.iv. The Four VC Dimensions

Reading off the four hypothesis sets of [Section 5.3](#), with a largest shattered configuration drawn next to each:

$\mathcal{H}$	d_{VC}	largest shattered inputs	$m_{\mathcal{H}}(N)$
positive rays	1	•	$N + 1$
positive intervals	2	• • up	$\frac{1}{2}N^2 + \frac{1}{2}N + 1$
convex sets	∞	• • • • • • • • bottom	2^N
2D perceptrons	3	• • • •	$\leq N^3$ for $N \geq 2$

Each d_{VC} is the minimum break point of the previous chapters minus one: $2 - 1$, $3 - 1$, none, $4 - 1$. The convex-set picture is the circle configuration of [Proposition 5.3.9](#), which is shattered for every N . For 2D perceptrons we still do not know $m_{\mathcal{H}}(N)$ exactly, but $d_{VC} = 3$ is enough to get $m_{\mathcal{H}}(N) \leq N^3$.

Good: finite d_{VC} .


§7.1.v. VC Dimension and Learning

Finite $d_{VC} \Rightarrow g$ ‘will’ generalize ($E_{out}(g) \approx E_{in}(g)$).

What is remarkable is what the guarantee does *not* need:

- regardless of learning algorithm $\mathcal{A}$;
- regardless of input distribution P ;
- regardless of target function f .

In the learning flow, those three boxes can be greyed out; the guarantee depends only on the hypothesis set and the size of the training set.



‘Worst case’ is the price of that generality: the bound has to hold for the least favourable $\mathcal{A}$, P and f at once, which is one reason it will turn out loose ([Section 7.4](#)).

Exercise 7.1.3 (Fun Time). If there is a set of N inputs that cannot be shattered by $\mathcal{H}$. Based only on this information, what can we conclude about $d_{VC}(\mathcal{H})$? (1) $d_{VC}(\mathcal{H}) > N$; (2) $d_{VC}(\mathcal{H}) = N$; (3) $d_{VC}(\mathcal{H}) < N$; (4) no conclusion can be made.


Reference Answer. (4). It is possible that there is another set of N inputs that can be shattered, which means $d_{VC} \geq N$. It is also possible that no set of N inputs can be shattered, which means $d_{VC} < N$. Neither case can be ruled out by one non-shattering set. (Three collinear points, which no line shatters, are the example: they say nothing against $d_{VC} = 3$ for 2D perceptrons.) ■

§7.2. VC Dimension of Perceptrons

— slides p.8–16 —

§7.2.i. 2D PLA Revisited

With d_{VC} in hand, the story of 2D PLA can finally be closed. Two independent chains lead to the same place:



The left chain is optimisation (Section 2.3): after enough iterations T , PLA makes no mistake on $\mathcal{D}$. The right chain is generalisation: with enough examples N , E_{in} is close to E_{out} for any line, including the one PLA returns. Together they give $E_{out}(g) \approx 0$.

Question 7.2.1. General PLA for $\mathbf{x}$ with more than 2 features?

§7.2.ii. VC Dimension of Perceptrons

What is known so far:

- 1D perceptron (pos/neg rays): $d_{VC} = 2$;
- 2D perceptrons: $d_{VC} = 3$
 - $d_{VC} \geq 3$:  can be shattered;
 - $d_{VC} \leq 3$:  no four inputs can be shattered;
- d -D perceptrons: $d_{VC} \stackrel{?}{=} d + 1$.

The first two fit the guess $d + 1$: $1 + 1$ and $2 + 1$. The extra 1 is the threshold, which [Section 2.1](#) absorbed as w_0 with the constant coordinate $x_0 = 1$, so a d -D perceptron really has $d + 1$ weights. Proving $d_{VC} = d + 1$ takes two steps:

- $d_{VC} \geq d + 1$;
- $d_{VC} \leq d + 1$.

The two steps need different kinds of evidence, as the next two quizzes make precise.

Exercise 7.2.2 (Extra Fun Time). What statement below shows that $d_{VC} \geq d + 1$?

- (1) There are some $d + 1$ inputs we can shatter.
- (2) We can shatter any set of $d + 1$ inputs.
- (3) There are some $d + 2$ inputs we cannot shatter.
- (4) We cannot shatter any set of $d + 2$ inputs.

Reference Answer. (1). d_{VC} is the maximum N such that $m_{\mathcal{H}}(N) = 2^N$, and $m_{\mathcal{H}}(N)$ is the most number of dichotomies of N inputs. So if we can find 2^{d+1} dichotomies on *some* $d + 1$ inputs, $m_{\mathcal{H}}(d + 1) = 2^{d+1}$ and hence $d_{VC} \geq d + 1$. (2) is also sufficient but asks for far more than needed — and is false for perceptrons, as collinear points show. ■

§7.2.iii. $d_{VC} \geq d + 1$

By the quiz, it is enough that there are *some* $d + 1$ inputs we can shatter. Pick some ‘trivial’ inputs: the origin and the d unit vectors, each with the constant coordinate $x_0 = 1$ in front. As the rows of a matrix,

$$X = \begin{bmatrix} - & \mathbf{x}_1^\top & - \\ - & \mathbf{x}_2^\top & - \\ - & \mathbf{x}_3^\top & - \\ & \vdots & \\ - & \mathbf{x}_{d+1}^\top & - \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & \dots & 0 \\ 1 & 1 & 0 & \dots & 0 \\ 1 & 0 & 1 & & 0 \\ \vdots & \vdots & & \ddots & 0 \\ 1 & 0 & \dots & 0 & 1 \end{bmatrix}.$$

Visually in 2D these are the three points $\bullet$ at $(0, 0)$, $(1, 0)$ and $(0, 1)$. Note: X is $(d+1) \times (d+1)$ and **invertible** — subtracting the first row from every other row leaves the identity matrix.

Remark 7.2.3. The $d + 1$ inputs are points of $\mathbb{R}^d$: the origin $\mathbf{0}$ and the d standard unit vectors $\mathbf{e}_1, \dots, \mathbf{e}_d$, where $\mathbf{e}_i$ has 1 in its i -th entry and 0 elsewhere. So $\mathbf{x}_1 = (1, \mathbf{0})$ and $\mathbf{x}_{i+1} = (1, \mathbf{e}_i)$ for $i = 1, \dots, d$. The first column of X is not part of the data: it is the bias coordinate $x_0 = 1$ ([Section 2.1](#)). This is why a d -D perceptron works with $(d + 1)$ -dimensional vectors, and why X is square.

§7.2.iv. Can We Shatter X ?

To shatter X means: for any $\mathbf{y} = (y_1, \dots, y_{d+1})^\top$ with entries ± 1 , find $\mathbf{w}$ such that $\text{sign}(X\mathbf{w}) = \mathbf{y}$, the n -th entry of $X\mathbf{w}$ being $\mathbf{w}^\top \mathbf{x}_n$. Demanding the signs is hard; demanding the values is easy, because X is invertible:

$$\text{sign}(X\mathbf{w}) = \mathbf{y} \quad \Longleftarrow \quad X\mathbf{w} = \mathbf{y} \quad \overset{X \text{ invertible!}}{\Longleftrightarrow} \quad \mathbf{w} = X^{-1}\mathbf{y}.$$

The first implication holds because $\text{sign}(\pm 1) = \pm 1$. So every one of the 2^{d+1} label vectors is realised by some perceptron:

‘Special’ X can be shattered $\Rightarrow d_{VC} \geq d + 1$.

Exercise 7.2.4 (Extra Fun Time). What statement below shows that $d_{VC} \leq d + 1$?

- (1) There are some $d + 1$ inputs we can shatter.
- (2) We can shatter any set of $d + 1$ inputs.
- (3) There are some $d + 2$ inputs we cannot shatter.
- (4) We cannot shatter any set of $d + 2$ inputs.

Reference Answer. (4). d_{VC} is the maximum N such that $m_{\mathcal{H}}(N) = 2^N$, and $m_{\mathcal{H}}(N)$ is the most number of dichotomies of N inputs. So if we cannot find 2^{d+2} dichotomies on *any* $d + 2$ inputs (i.e. break point), $m_{\mathcal{H}}(d + 2) < 2^{d+2}$ and hence $d_{VC} < d + 2$. That is, $d_{VC} \leq d + 1$. Statement (3) is the trap of the Fun Time in [Section 7.1](#): one bad set proves nothing. ■

The upper bound is therefore the harder half: a single well-chosen example no longer suffices, and the argument must work for *every* set of $d + 2$ inputs.

§7.2.v. $d_{VC} \leq d + 1$: A 2D Special Case

Start with four inputs in 2D, the corners of the unit square:

$$\begin{array}{cc} \bullet & \bullet \\ \bullet & \bullet \end{array} \quad X = \begin{bmatrix} - & \mathbf{x}_1^\top & - \\ - & \mathbf{x}_2^\top & - \\ - & \mathbf{x}_3^\top & - \\ - & \mathbf{x}_4^\top & - \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 1 \end{bmatrix}.$$

So $\mathbf{x}_1$ is the lower-left corner, $\mathbf{x}_2$ lower-right, $\mathbf{x}_3$ upper-left, $\mathbf{x}_4$ upper-right. Four rows in a three-dimensional space must be linearly dependent, and here the dependence is visible: $\mathbf{x}_4 = \mathbf{x}_2 + \mathbf{x}_3 - \mathbf{x}_1$. Now label $\mathbf{x}_2, \mathbf{x}_3$ as $\circ$ and $\mathbf{x}_1$ as $\times$:

$$\begin{array}{cc} \circ & ? \\ \times & \circ \end{array} \quad \text{? cannot be } \times$$

because for every $\mathbf{w}$ producing these three labels,

$$\mathbf{w}^\top \mathbf{x}_4 = \underbrace{\mathbf{w}^\top \mathbf{x}_2}_{\circ} + \underbrace{\mathbf{w}^\top \mathbf{x}_3}_{\circ} - \underbrace{\mathbf{w}^\top \mathbf{x}_1}_{\times} > 0 :$$

the first two terms are positive, and subtracting a negative one only adds. This is the XOR dichotomy again, now explained algebraically.

Linear dependence restricts dichotomy.

§7.2.vi. $d_{VC} \leq d + 1$: The d -D General Case

By the Extra Fun Time above, proving $d_{VC} \leq d + 1$ means proving that $m_{\mathcal{H}}(d + 2) < 2^{d+2}$: *no* set of $d + 2$ inputs can be shattered. One bad configuration, such as the square, is not enough; the argument has to work for every choice of $d + 2$ inputs. So take *any* $d + 2$ inputs, as rows of

$$X = \begin{bmatrix} - & \mathbf{x}_1^\top & - \\ - & \mathbf{x}_2^\top & - \\ & \vdots & \\ - & \mathbf{x}_{d+1}^\top & - \\ - & \mathbf{x}_{d+2}^\top & - \end{bmatrix}.$$

There are more rows ($d + 2$) than columns ($d + 1$), so the rows are linearly dependent: some row is a combination of the others. Relabelling so that this row is the last one,

$$\mathbf{x}_{d+2} = a_1 \mathbf{x}_1 + a_2 \mathbf{x}_2 + \cdots + a_{d+1} \mathbf{x}_{d+1}, \quad \text{some } a_i \text{ non-zero.}$$

- Can you generate $(\text{sign}(a_1), \text{sign}(a_2), \dots, \text{sign}(a_{d+1}), \times)$? If so, what $\mathbf{w}$?

$$\mathbf{w}^\top \mathbf{x}_{d+2} = \underbrace{a_1 \mathbf{w}^\top \mathbf{x}_1}_{\circ} + \underbrace{a_2 \mathbf{w}^\top \mathbf{x}_2}_{\times} + \cdots + \underbrace{a_{d+1} \mathbf{w}^\top \mathbf{x}_{d+1}}_{\times} > 0 \quad (\text{contradiction!})$$

The braces show a typical sign pattern, $a_1 > 0$ and the rest negative. Note the direction of the argument. To show that the $d + 2$ inputs are *not* shattered, it is not enough to find one $\mathbf{w}$ with a bad output; we must exhibit one dichotomy that *no* $\mathbf{w}$ produces. The dichotomy is chosen from the coefficients a_i , and then every $\mathbf{w}$ that gets the first $d + 1$ labels right is shown to be forced to say $\circ$ on $\mathbf{x}_{d+2}$.

Lemma 7.2.5 (Linear dependence blocks a dichotomy)

Let $\mathbf{x}_1, \dots, \mathbf{x}_{d+2} \in \mathbb{R}^{d+1}$ be any $d + 2$ inputs (bias coordinate $x_0 = 1$ included). Then some dichotomy $(y_1, \dots, y_{d+2})$ is not $(\text{sign}(\mathbf{w}^\top \mathbf{x}_1), \dots, \text{sign}(\mathbf{w}^\top \mathbf{x}_{d+2}))$ for any $\mathbf{w} \in \mathbb{R}^{d+1}$.

Proof. Step 1: the coefficients. $d + 2$ vectors in $\mathbb{R}^{d+1}$ are linearly dependent, so $\sum_{i=1}^{d+2} c_i \mathbf{x}_i = \mathbf{0}$ for some c_i not all zero. Pick j with $c_j \neq 0$, divide by $-c_j$, and renumber so that $j = d + 2$:

$$\mathbf{x}_{d+2} = \sum_{i=1}^{d+1} a_i \mathbf{x}_i.$$

At least one a_i is non-zero, since otherwise $\mathbf{x}_{d+2} = \mathbf{0}$, impossible because its first coordinate is $x_0 = 1$.

Step 2: the dichotomy. Define

$$y_i = \begin{cases} \text{sign}(a_i) & \text{if } a_i \neq 0, \\ +1 \text{ (say; any value works)} & \text{if } a_i = 0, \end{cases} \quad i = 1, \dots, d + 1, \quad y_{d+2} = -1.$$

Step 3: no $\mathbf{w}$ produces it. Suppose $\mathbf{w}$ gives $\text{sign}(\mathbf{w}^\top \mathbf{x}_i) = y_i$ for $i = 1, \dots, d + 1$. Take the inner product of Step 1 with $\mathbf{w}$:

$$\mathbf{w}^\top \mathbf{x}_{d+2} = \sum_{i=1}^{d+1} a_i (\mathbf{w}^\top \mathbf{x}_i).$$

Look at each term. If $a_i = 0$ the term is 0. If $a_i \neq 0$, then $\mathbf{w}^\top \mathbf{x}_i$ has the sign $y_i = \text{sign}(a_i)$, in particular it is non-zero, so $a_i(\mathbf{w}^\top \mathbf{x}_i)$ is a product of two non-zero numbers of the same sign, hence > 0 . Every term is ≥ 0 and at least one is > 0 , so $\mathbf{w}^\top \mathbf{x}_{d+2} > 0$ and $\text{sign}(\mathbf{w}^\top \mathbf{x}_{d+2}) = +1 \neq y_{d+2}$.

So every $\mathbf{w}$ that is right on the first $d + 1$ inputs is wrong on the last, and $(y_1, \dots, y_{d+2})$ is never produced. ■

The 2D special case above is the instance $d = 2$, $\mathbf{x}_4 = \mathbf{x}_2 + \mathbf{x}_3 - \mathbf{x}_1$: $a = (-1, 1, 1)$, so the blocked dichotomy is $(\times, \circ, \circ, \times)$. Since the lemma holds for *any* $d + 2$ inputs, no set of $d + 2$ inputs is shattered: $m_{\mathcal{H}}(d + 2) < 2^{d+2}$, i.e. $d + 2$ is a break point, which is statement (4) of the last Extra Fun Time.

‘General’ X no-shatter $\Rightarrow d_{\text{VC}} \leq d + 1$.

With both steps done:

Theorem 7.2.6 (VC dimension of perceptrons)

For d -D perceptrons, $d_{VC} = d + 1$.

Exercise 7.2.7 (Fun Time). Based on the proof above, what is d_{VC} of 1126-D perceptrons? (1) 1024; (2) 1126; (3) 1127; (4) 6211.


Reference Answer. (3). $1126 + 1 = 1127$. Well, too much fun for this section! :-)

§7.3. Physical Intuition of VC Dimension

— slides p.17–20 —

Theorem 7.2.6 says a perceptron has d_{VC} equal to its number of weights. That is suggestive enough to ask what d_{VC} is measuring in general.

§7.3.i. Degrees of Freedom



(modified from the work of Hugues Vermeiren on <http://www.texample.net>)

Think of a hypothesis as a machine with knobs, each weight being one knob. The three quantities met so far are three ways of counting what the knobs can do:



- hypothesis parameters $\mathbf{w} = (w_0, w_1, \dots, w_d)$: **create degrees of freedom**;
- hypothesis quantity $M = |\mathcal{H}|$: ‘analog’ degrees of freedom;
- hypothesis ‘power’ $d_{VC} = d + 1$: **effective ‘binary’ degrees of freedom**.

M counts every knob setting, and is infinite as soon as a knob turns continuously. d_{VC} counts only how many inputs the knobs can label independently in *every* possible way — how many free binary choices the knobs really control.

$d_{VC}(\mathcal{H})$: **powerfulness of $\mathcal{H}$.**

§7.3.ii. Two Old Friends

The hypothesis sets of [Section 5.3](#) fit the same picture.

Positive Rays ($d_{VC} = 1$)**Positive Intervals ($d_{VC} = 2$)**

One threshold, one binary degree of freedom; two endpoints, two. This suggests

Practical rule of thumb: $d_{VC} \approx \# \text{free parameters (but not always)}$.

‘Not always’ in both directions: parameters can be redundant ($h(x) = \text{sign}(w_1 w_2 x)$ has two parameters but only the sign of the product matters), and a single parameter can be enormously powerful ($h(x) = \text{sign}(\sin(\alpha x))$ has one parameter and $d_{VC} = \infty$).

§7.3.iii. M and d_{VC}

Copied from [Section 5.1](#) :-), the two central questions are:

- (1) can we make sure that $E_{\text{out}}(g)$ is close enough to $E_{\text{in}}(g)$?
- (2) can we make $E_{\text{in}}(g)$ small enough?

d_{VC} now plays exactly the role M played there, with M replaced by the polynomial $(2N)^{d_{VC}}$:

	(1) $E_{\text{out}} \approx E_{\text{in}}?$	(2) E_{in} small?
small M	Yes!, $\mathbb{P}[\mathbf{BAD}] \leq 2 \cdot M \cdot \exp(\dots)$	No!, too few choices
large M	No!, $\mathbb{P}[\mathbf{BAD}] \leq 2 \cdot M \cdot \exp(\dots)$	Yes!, many choices
small d_{VC}	Yes!, $\mathbb{P}[\mathbf{BAD}] \leq 4 \cdot (2N)^{d_{VC}} \cdot \exp(\dots)$	No!, too limited power
large d_{VC}	No!, $\mathbb{P}[\mathbf{BAD}] \leq 4 \cdot (2N)^{d_{VC}} \cdot \exp(\dots)$	Yes!, lots of power

Using the right d_{VC} (or $\mathcal{H}$) is important.

Exercise 7.3.1 (Fun Time). Origin-crossing hyperplanes are essentially perceptrons with w_0 fixed at 0. Make a guess about the d_{VC} of origin-crossing hyperplanes in $\mathbb{R}^d$. (1) 1; (2) d ; (3) $d + 1$; (4) ∞ .

Reference Answer. (2). The proof is almost the same as proving the d_{VC} for usual perceptrons, but it is the intuition ($d_{VC} \approx \# \text{free parameters}$) that you shall use to answer this quiz: fixing w_0 removes one of the $d + 1$ knobs. (In the proof, drop the constant column of X : the d unit vectors are shattered, and any $d + 1$ vectors in $\mathbb{R}^d$ are linearly dependent.) ■

§7.4. Interpreting VC Dimension

— slides p.21–25 —

Everything so far has been about *whether* generalisation happens. Turning the VC bound around shows *how much* it costs, in two currencies: error and data.

§7.4.i. VC Bound Rephrase: Penalty for Model Complexity

For any $g = \mathcal{A}(\mathcal{D}) \in \mathcal{H}$ and a ‘statistical’ large $\mathcal{D}$, for $N \geq 2$, $d_{\text{VC}} \geq 2$ (the condition ‘ $k \geq 3$ ’ of [Section 7.1](#), now in terms of d_{VC}),

$$\mathbb{P}_{\mathcal{D}} \left[\underbrace{|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon}_{\text{BAD}} \right] \leq \underbrace{4(2N)^{d_{\text{VC}}} \exp \left(-\frac{1}{8} \epsilon^2 N \right)}_{\delta}.$$

Call the right-hand side δ . The bound then says: ..., with probability $\geq 1 - \delta$, **GOOD**: $|E_{\text{in}}(g) - E_{\text{out}}(g)| \leq \epsilon$. Instead of fixing ϵ and asking for δ , fix the confidence δ and solve for ϵ :

$$\begin{aligned} \text{set } \delta &= 4(2N)^{d_{\text{VC}}} \exp \left(-\frac{1}{8} \epsilon^2 N \right) \\ \frac{\delta}{4(2N)^{d_{\text{VC}}}} &= \exp \left(-\frac{1}{8} \epsilon^2 N \right) \\ \ln \left(\frac{4(2N)^{d_{\text{VC}}}}{\delta} \right) &= \frac{1}{8} \epsilon^2 N \\ \sqrt{\frac{8}{N} \ln \left(\frac{4(2N)^{d_{\text{VC}}}}{\delta} \right)} &= \epsilon. \end{aligned}$$

Rephrased: ..., with probability $\geq 1 - \delta$, **GOOD!**

$$\text{gen. error } |E_{\text{in}}(g) - E_{\text{out}}(g)| \leq \sqrt{\frac{8}{N} \ln \left(\frac{4(2N)^{d_{\text{VC}}}}{\delta} \right)},$$

that is,

$$E_{\text{in}}(g) - \sqrt{\frac{8}{N} \ln \left(\frac{4(2N)^{d_{\text{VC}}}}{\delta} \right)} \leq E_{\text{out}}(g) \leq E_{\text{in}}(g) + \sqrt{\frac{8}{N} \ln \left(\frac{4(2N)^{d_{\text{VC}}}}{\delta} \right)}.$$

The left inequality is greyed out because nobody minds E_{out} being *smaller* than expected; only the upper bound matters.

Definition 7.4.1 (Penalty for model complexity). The square root

$$\Omega(N, \mathcal{H}, \delta) = \sqrt{\frac{8}{N} \ln \left(\frac{4(2N)^{d_{\text{VC}}}}{\delta} \right)}$$


is the **penalty for model complexity**.

It depends on the data through N , on the hypothesis set through d_{VC} , and on how sure we want to be through δ .

§7.4.ii. THE VC Message

With a high probability,

$$E_{\text{out}}(g) \leq E_{\text{in}}(g) + \underbrace{\sqrt{\frac{8}{N} \ln \left(\frac{4(2N)^{d_{\text{VC}}}}{\delta} \right)}}_{\Omega(N, \mathcal{H}, \delta)}.$$



The picture reads the inequality as a function of d_{VC} , with the bound standing in for E_{out} :

- $d_{\text{VC}} \uparrow$: $E_{\text{in}} \downarrow$ but $\Omega \uparrow$;
- $d_{\text{VC}} \downarrow$: $\Omega \downarrow$ but $E_{\text{in}} \uparrow$;
- best d_{VC}^* in the middle.

The first bullet is the intuition of [Section 7.3](#): d_{VC} counts the effective binary degrees of freedom, the powerfulness of $\mathcal{H}$, and a more powerful $\mathcal{H}$ can fit the labels better — the ‘Yes!, lots of power’ entry of the M -and- d_{VC} table. But the penalty grows with it, and their sum is smallest somewhere in between.

Remark 7.4.2. ‘More powerful, hence smaller E_{in} ’ is an intuition, not a fact about d_{VC} alone. Making it precise takes two assumptions that the picture leaves implicit.

1. The hypothesis sets are *nested*: $\mathcal{H}_1 \subset \mathcal{H}_2 \subset \dots$ with $d_{\text{VC}}(\mathcal{H}_1) \leq d_{\text{VC}}(\mathcal{H}_2) \leq \dots$, as when more weights are added to a perceptron (setting the new ones to 0 recovers the old hypotheses).
2. $\mathcal{A}$ returns the g of smallest E_{in} in $\mathcal{H}$.

Under these,

$$\min_{h \in \mathcal{H}_2} E_{\text{in}}(h) \leq \min_{h \in \mathcal{H}_1} E_{\text{in}}(h),$$

since the best hypothesis of $\mathcal{H}_1$ is still available in $\mathcal{H}_2$. A larger d_{VC} means more dichotomies on the data, so a better match to $y_1, \dots, y_N$ is usually found and the inequality is usually strict. Without nesting there is no such guarantee: on three collinear inputs labelled $\times \circ \times$, positive intervals along that line ($d_{\text{VC}} = 2$) reach $E_{\text{in}} = 0$, while 2D perceptrons ($d_{\text{VC}} = 3$) cannot.

Powerful $\mathcal{H}$ not always good!

§7.4.iii. VC Bound Rephrase: Sample Complexity

The same bound can be solved for N instead. Given specs $\epsilon = 0.1$, $\delta = 0.1$, $d_{VC} = 3$, we want

$$4(2N)^{d_{VC}} \exp\left(-\frac{1}{8}\epsilon^2 N\right) \leq \delta.$$

No closed form, but tabulating the left-hand side is enough:

N	bound
100	2.82×10^7
1,000	9.17×10^9
10,000	1.19×10^8
100,000	1.65×10^{-38}
29,300	9.99×10^{-2}

For small N the polynomial factor wins and the ‘probability’ bound is in the millions; the exponential only takes over between 10^4 and 10^5 , and the spec is first met around $N = 29,300$. This gives the **sample complexity**: need $N \approx 10,000 d_{VC}$ in theory.

Practical rule of thumb: $N \approx 10 d_{VC}$ often enough!

§7.4.iv. Looseness of VC Bound

$$\mathbb{P}_{\mathcal{D}} [|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \leq 4(2N)^{d_{VC}} \exp\left(-\frac{1}{8}\epsilon^2 N\right)$$

Theory: $N \approx 10,000 d_{VC}$; practice: $N \approx 10 d_{VC}$.

Question 7.4.3. Why is the VC bound so loose?

Every step of the derivation paid for generality:

step	it must hold for
Hoeffding for unknown E_{out}	any distribution, any target
$m_{\mathcal{H}}(N)$ instead of $ \mathcal{H}(\mathbf{x}_1, \dots, \mathbf{x}_N) $	‘any’ data
$N^{d_{VC}}$ instead of $m_{\mathcal{H}}(N)$	‘any’ $\mathcal{H}$ of same d_{VC}
union bound on worst cases	any choice made by $\mathcal{A}$

— but hardly better, and ‘similarly loose for all models’. Tightening one step usually means assuming something about P , f , the data or $\mathcal{A}$, which gives up the very generality that made the bound useful. And because the looseness is roughly the same for every model, the bound still ranks models correctly even when its numbers are useless.

Philosophical message of VC bound: important for improving ML.

Exercise 7.4.4 (Fun Time). Consider the VC bound below. How can we decrease the probability of getting BAD data?

$$\mathbb{P}_{\mathcal{D}} [|E_{\text{in}}(g) - E_{\text{out}}(g)| > \epsilon] \leq 4(2N)^{d_{VC}} \exp\left(-\frac{1}{8}\epsilon^2 N\right)$$

- (1) decrease model complexity d_{VC} ;
- (2) increase data size N a lot;
- (3) increase generalization error tolerance ϵ ;
- (4) all of the above.

Reference Answer. (4). Congratulations on being Master of VC bound! :-) Each of (1)–(3) shrinks the right-hand side; ‘a lot’ in (2) because the polynomial factor grows with N too, and only wins once N is large (the table above). ■

Summary

Definition of VC Dimension maximum non-break point.

VC Dimension of Perceptrons $d_{VC}(\mathcal{H}) = d + 1$.

Physical Intuition of VC Dimension $d_{VC} \approx \# \text{free parameters}$.

Interpreting VC Dimension loosely: model complexity & sample complexity.

Learning happens if d_{VC} is finite, N is large and E_{in} is low. All of that was derived for the cleanest possible setting: binary classification with a deterministic target and no noise. Next: more than noiseless binary classification?

8. Noise and Error

Lecture 8

As previously seen . Section 7.4 ended with the VC message: learning happens if d_{VC} is finite, N is large, and E_{in} is low. The whole derivation, however, assumed a deterministic target f , labels $y_n = f(\mathbf{x}_n)$ with no noise, and binary classification measured by the fraction of mistakes.


30 Sep.

This chapter relaxes the first two assumptions and questions the third. Noise turns the target function into a target *distribution*, and the VC bound survives the change. Once the target is a distribution, ‘the best prediction’ depends on how errors are measured, so the error measure becomes a component of the learning problem in its own right — chosen by the application, and then often replaced by something the algorithm can actually optimise. The last section works one such case, weighted classification, end to end.

§8.1. Noise and Probabilistic Target

— slides p.2–7 —


§8.1.i. Recap: The Learning Flow



This is the flow of Section 4.3, with one addition in the target box.

Question 8.1.1. What if there is noise?

§8.1.ii. Noise



Noise was briefly introduced before the **pocket** algorithm (Section 2.4), as the reason a data set might not be linearly separable even when the underlying rule is a line: a few examples sit on the wrong side of the best line.

§8.1.iv. Target Distribution $P(y \mid \mathbf{x})$

Definition 8.1.3 (Target distribution). The **target distribution** $P(y \mid \mathbf{x})$ characterizes the behaviour of the ‘mini-target’ on one $\mathbf{x}$.

- It can be viewed as ‘ideal mini-target’ + noise, e.g.
 - $P(\circ \mid \mathbf{x}) = 0.7$, $P(\times \mid \mathbf{x}) = 0.3$;
 - ideal mini-target $f(\mathbf{x}) = \circ$;
 - ‘flipping’ noise level = 0.3.
- A deterministic target f is a **special case of target distribution**:
 - $P(y \mid \mathbf{x}) = 1$ for $y = f(\mathbf{x})$;
 - $P(y \mid \mathbf{x}) = 0$ for $y \neq f(\mathbf{x})$.


So nothing is lost by switching: every earlier setting is a target distribution concentrated on one label. The two distributions divide the work:

Goal of learning: predict the ideal mini-target (w.r.t. $P(y \mid \mathbf{x})$) on often-seen inputs (w.r.t. $P(\mathbf{x})$).

$P(y \mid \mathbf{x})$ is what we try to learn; $P(\mathbf{x})$ only decides where it matters to get it right.

§8.1.v. The New Learning Flow

In the flow, the target function box becomes a target distribution, which now supplies the labels $y_1, \dots, y_N$ of the training set *and* the label y against which g is judged on a future $\mathbf{x}$.



VC still works, pocket algorithm explained :-)

Pocket was introduced for data that no line fits. With noise, that is the normal case even when the ideal mini-target is linear; the VC bound says the line with the smallest E_{in} still has $E_{\text{out}} \approx E_{\text{in}}$, so pocket’s choice is justified.

Exercise 8.1.4 (Fun Time). Let's revisit PLA/pocket. Which of the following claim is true?

- (1) In practice, we should try to compute if $\mathcal{D}$ is linear separable before deciding to use PLA.
- (2) If we know that $\mathcal{D}$ is not linear separable, then the target function f must not be a linear function.
- (3) If we know that $\mathcal{D}$ is linear separable, then the target function f must be a linear function.
- (4) None of the above.

Reference Answer. (4). (1) After computing if $\mathcal{D}$ is linear separable, we shall know $\mathbf{w}^*$ and then there is no need to use PLA. (2) What about noise? A linear target plus flipping noise gives non-separable data. (3) What about 'sampling luck'? :-) A curved target can happen to produce a sample that some line separates. ■

§8.2. Error Measure

— slides p.8–13 —

The target is now a distribution, and g is supposed to be $\approx f$. Before learning anything, we need to say what ' $\approx$ ' means.

§8.2.i. Error Measure

final hypothesis
 $g \approx f$

- How well? Previously, we considered the out-of-sample measure

$$E_{\text{out}}(g) = \mathcal{E}_{\mathbf{x} \sim P} \llbracket g(\mathbf{x}) \neq f(\mathbf{x}) \rrbracket.$$

- More generally, an **error measure** $E(g, f)$.
- Naturally considered:
 - out-of-sample: averaged over unknown $\mathbf{x}$;
 - pointwise: evaluated on one $\mathbf{x}$;
 - classification: $\llbracket \text{prediction} \neq \text{target} \rrbracket$.

Here $\mathcal{E}$ denotes expectation. The three adjectives are three independent choices that E_{out} happened to make at once, and each could be made differently.

Classification error $\llbracket \cdot \rrbracket$: often also called '0/1 error'.

§8.2.ii. Pointwise Error Measure

We can often express $E(g, f)$ as an average of $\text{err}(g(\mathbf{x}), f(\mathbf{x}))$, like

$$E_{\text{out}}(g) = \mathcal{E}_{\mathbf{x} \sim P} \underbrace{\llbracket g(\mathbf{x}) \neq f(\mathbf{x}) \rrbracket}_{\text{err}(g(\mathbf{x}), f(\mathbf{x}))}.$$

Definition 8.2.1 (Pointwise error measure). The function err above is called a **pointwise error measure**. With it,

$$\underbrace{E_{\text{in}}(g) = \frac{1}{N} \sum_{n=1}^N \text{err}(g(\mathbf{x}_n), f(\mathbf{x}_n))}_{\text{in-sample}} \quad \underbrace{E_{\text{out}}(g) = \mathcal{E}_{\mathbf{x} \sim P} \text{err}(g(\mathbf{x}), f(\mathbf{x}))}_{\text{out-of-sample}}.$$

Not every sensible measure is pointwise — the area under an ROC curve, say, compares predictions on different inputs — but we will mainly consider pointwise err for simplicity.

§8.2.iii. Two Important Pointwise Error Measures

Write the prediction and the target as

$$\text{err}(\underbrace{g(\mathbf{x})}_{\tilde{y}}, \underbrace{f(\mathbf{x})}_{y}).$$

0/1 error	squared error
$\text{err}(\tilde{y}, y) = \mathbb{I}[\tilde{y} \neq y]$	$\text{err}(\tilde{y}, y) = (\tilde{y} - y)^2$
correct or incorrect?	how far is $\tilde{y}$ from y ?
often for classification	often for regression

Question 8.2.2. How does err ‘guide’ learning?

§8.2.iv. Ideal Mini-Target

The interplay between noise and error:

$P(y | \mathbf{x})$ and err define the ideal mini-target $f(\mathbf{x})$.

With noise there is no single ‘true’ label to recover at $\mathbf{x}$: what nature gives is the whole distribution $P(y | \mathbf{x})$. The f we aim at is therefore not a target function handed to us but the best *prediction* against that distribution, and ‘best’ only means something once err is fixed.

Definition 8.2.3 (Ideal mini-target). Given $P(y | \mathbf{x})$ and err , the **ideal mini-target** at $\mathbf{x}$ is the prediction of smallest expected error,

$$f(\mathbf{x}) = \arg \min_{\tilde{y}} \mathcal{E}_{y \sim P(y|\mathbf{x})} \text{err}(\tilde{y}, y).$$

So from here on f denotes the error-minimising function, not a function that generated the labels; for a deterministic target ($P(y | \mathbf{x}) = 1$ at $y = f(\mathbf{x})$) and any err with $\text{err}(y, y) = 0 \leq \text{err}(\tilde{y}, y)$, the two coincide.

Take one input whose labels are

$$P(y = 1 | \mathbf{x}) = 0.2, \quad P(y = 2 | \mathbf{x}) = 0.7, \quad P(y = 3 | \mathbf{x}) = 0.1,$$

and ask which prediction $\tilde{y}$ has the smallest average error under each measure:

$\tilde{y}$	$\text{err}(\tilde{y}, y) = \mathbb{I}[\tilde{y} \neq y]$	$\text{err}(\tilde{y}, y) = (\tilde{y} - y)^2$
1	avg. err 0.8	avg. err 1.1
2	avg. err 0.3 (*)	avg. err 0.3
3	avg. err 0.9	avg. err 1.5
1.9	avg. err 1.0 (really? :-))	avg. err 0.29 (*)
	$f(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} P(y \mid \mathbf{x})$	$f(\mathbf{x}) = \sum_{y \in \mathcal{Y}} y \cdot P(y \mid \mathbf{x})$

The two starred rows are no accident: they are the two minimisers below.

Proposition 8.2.4 (Ideal mini-targets of 0/1 and squared error)

Fix $\mathbf{x}$ and write $\mathcal{E}$ for expectation over $y \sim P(y \mid \mathbf{x})$.

1. For $\text{err}(\tilde{y}, y) = \mathbb{I}[\tilde{y} \neq y]$, $f(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} P(y \mid \mathbf{x})$ (any maximiser, if there are ties).
2. For $\text{err}(\tilde{y}, y) = (\tilde{y} - y)^2$, $f(\mathbf{x}) = \sum_{y \in \mathcal{Y}} y \cdot P(y \mid \mathbf{x}) = \mathcal{E}[y]$, and this minimiser is unique.

Proof. (1) For any real $\tilde{y}$,

$$\mathcal{E}[\mathbb{I}[\tilde{y} \neq y]] = \mathbb{P}[y \neq \tilde{y}] = 1 - P(y = \tilde{y} \mid \mathbf{x}),$$

where $P(y = \tilde{y} \mid \mathbf{x}) = 0$ if $\tilde{y} \notin \mathcal{Y}$. Minimising the left side is maximising $P(y = \tilde{y} \mid \mathbf{x})$. A $\tilde{y}$ outside $\mathcal{Y}$ scores the worst possible value 1, so the minimum is attained exactly at the labels of largest probability.

(2) Let $\mu = \mathcal{E}[y]$. Add and subtract μ :

$$\mathcal{E}(\tilde{y} - y)^2 = \mathcal{E}[(\tilde{y} - \mu) + (\mu - y)]^2 = (\tilde{y} - \mu)^2 + 2(\tilde{y} - \mu) \mathcal{E}[\mu - y] + \mathcal{E}(\mu - y)^2.$$

The middle term vanishes because $\mathcal{E}[\mu - y] = \mu - \mu = 0$, and the last term is $\text{Var}(y)$, which does not depend on $\tilde{y}$. Hence


$$\mathcal{E}(\tilde{y} - y)^2 = (\tilde{y} - \mu)^2 + \text{Var}(y) \geq \text{Var}(y),$$

with equality if and only if $\tilde{y} = \mu$. ■

On the example, (1) gives $1 - 0.7 = 0.3$ at $\tilde{y} = 2$, and 1 at $\tilde{y} = 1.9$, which is never a label and so is always wrong. (2) gives $\mu = 0.2 + 1.4 + 0.3 = 1.9$ and $\text{Var}(y) = 3.9 - 1.9^2 = 0.29$, the table's minimum; predicting 2 instead costs $(2 - 1.9)^2 + 0.29 = 0.3$. Same noise, different err, different ideal mini-target.

§8.2.v. Learning Flow with Error Measure

The error measure therefore joins the flow as a component: it is what ' $g \approx f$ ' is judged by.



Extended VC theory/‘philosophy’ works for most $\mathcal{H}$ and err .

The proof in [Section 6.4](#) was for 0/1 error; analogous bounds exist for other bounded error measures, with d_{VC} replaced by suitable generalisations. We take the philosophy on trust from here.

Exercise 8.2.5 (Fun Time). Consider the following $P(y | \mathbf{x})$ and $\text{err}(\tilde{y}, y) = |\tilde{y} - y|$. Which of the following is the ideal mini-target $f(\mathbf{x})$?

$$\begin{aligned} P(y = 1 | \mathbf{x}) &= 0.10, & P(y = 2 | \mathbf{x}) &= 0.35, \\ P(y = 3 | \mathbf{x}) &= 0.15, & P(y = 4 | \mathbf{x}) &= 0.40. \end{aligned}$$

- (1) $2.5 = \text{average within } \mathcal{Y} = \{1, 2, 3, 4\}$;
- (2) $2.85 = \text{weighted mean from } P(y | \mathbf{x})$;
- (3) $3 = \text{weighted median from } P(y | \mathbf{x})$;
- (4) $4 = \arg \max P(y | \mathbf{x})$.

Reference Answer. (3). For the ‘absolute error’, the weighted median provably results in the minimum average error. The average errors of the four candidates are 1.0, 0.965, 0.95 and 1.15. Moving $\tilde{y}$ to the right by a little changes the average by (mass to the left) – (mass to the right); that is negative until half the mass is on the left, which happens at 3 ($0.10 + 0.35 + 0.15 = 0.6 \geq 0.5$). ■


§8.3. Algorithmic Error Measure

— slides p.14–19 —

The last section treated err as given. Who gives it?

§8.3.i. Choice of Error Measure

Fingerprint verification. A fingerprint goes in, and the target says whether it belongs to you:



There are two types of error: **false accept** and **false reject**.

		g	
		+1	-1
f	+1	no error	false reject
	-1	false accept	no error

0/1 error penalizes both types equally.

Whether that is right depends entirely on who uses the verifier.

§8.3.ii. Fingerprint Verification for Supermarket

		g	
		+1	-1
f	+1	0	10
	-1	1	0

- Supermarket: fingerprint for discount.
- False reject: **very unhappy customer, lose future business.**
- False accept: give away a minor discount, intruder left fingerprint :-)

§8.3.iii. Fingerprint Verification for CIA

		g	
		+1	-1
f	+1	0	1
	-1	1000	0

- CIA: fingerprint for entrance.
- False accept: **very serious consequences!**
- False reject: unhappy employee, but so what? :-)

Same inputs, same target, same two kinds of mistake — and cost ratios of 10 : 1 one way and 1 : 1000 the other. The best g for one application is a poor g for the other.

§8.3.iv. Take-home Message for Now

err is application/user-dependent.

In practice the user rarely writes the true err down, so the algorithm works with a stand-in.


Definition 8.3.1 (Algorithmic error measure). An **algorithmic error measure** $\widehat{\text{err}}$ is the error measure $\mathcal{A}$ actually uses in place of err . It may be

- true: just err ;
- plausible:
 - 0/1: minimum ‘flipping noise’ — NP-hard to optimize, remember? :-)
 - squared: minimum Gaussian noise;
- friendly: easy to optimize for $\mathcal{A}$:
 - closed-form solution;
 - convex objective function.

‘Plausible’ means justified by a belief about the noise: if labels are flipped at random, the best one can do is make the fewest disagreements, which is 0/1 error — and minimising it over perceptrons is the NP-hard problem of [Section 2.4](#); if the noise is Gaussian, maximum likelihood is least squares. ‘Friendly’ means justified by what the optimiser can do.

$\widehat{\text{err}}$: more in next lectures.

§8.3.v. Learning Flow with Algorithmic Error Measure



err : application goal; $\widehat{\text{err}}$: a key part of many $\mathcal{A}$.

err sits outside the algorithm and judges g ; $\widehat{\text{err}}$ sits inside and steers the search. The art is choosing $\widehat{\text{err}}$ close enough to err to aim at the right goal, and friendly enough to be optimised.

Exercise 8.3.2 (Fun Time). Consider err below for CIA. What is $E_{\text{in}}(g)$ when using this err?

		g	
		+1	-1
f	+1	0	1
	-1	1000	0

- (1) $\frac{1}{N} \sum_{n=1}^N \mathbb{I}[y_n \neq g(\mathbf{x}_n)];$
- (2) $\frac{1}{N} \left(\sum_{y_n=+1} \mathbb{I}[y_n \neq g(\mathbf{x}_n)] + 1000 \sum_{y_n=-1} \mathbb{I}[y_n \neq g(\mathbf{x}_n)] \right);$
- (3) $\frac{1}{N} \left(\sum_{y_n=+1} \mathbb{I}[y_n \neq g(\mathbf{x}_n)] - 1000 \sum_{y_n=-1} \mathbb{I}[y_n \neq g(\mathbf{x}_n)] \right);$
- (4) $\frac{1}{N} \left(1000 \sum_{y_n=+1} \mathbb{I}[y_n \neq g(\mathbf{x}_n)] + \sum_{y_n=-1} \mathbb{I}[y_n \neq g(\mathbf{x}_n)] \right).$

Reference Answer. (2). When $y_n = -1$, the false positive made on such $(\mathbf{x}_n, y_n)$ is penalized 1000 times more! ■

§8.4. Weighted Classification

— slides p.20–24 —

The CIA error of the quiz is a concrete, useful err that is not 0/1. This section shows how to learn with it using nothing but the tools of [Section 2.4](#).

§8.4.i. Weighted Classification

CIA cost (error, loss, ...) matrix.

		$h(\mathbf{x})$	
		+1	-1
y	+1	0	1
	-1	1000	0

Out-of-sample,

$$E_{\text{out}}(h) = \mathcal{E}_{(\mathbf{x}, y) \sim P} \left\{ \begin{matrix} 1 & \text{if } y = +1 \\ 1000 & \text{if } y = -1 \end{matrix} \right\} \cdot \mathbb{I}[y \neq h(\mathbf{x})];$$

in-sample,

$$E_{\text{in}}(h) = \frac{1}{N} \sum_{n=1}^N \left\{ \begin{matrix} 1 & \text{if } y_n = +1 \\ 1000 & \text{if } y_n = -1 \end{matrix} \right\} \cdot \mathbb{I}[y_n \neq h(\mathbf{x}_n)].$$

The cost of a mistake depends only on the true label, so the matrix collapses into one weight per example.

Definition 8.4.1 (Weighted classification). **Weighted classification** is classification with a different ‘weight’ for different $(\mathbf{x}, y)$: the pointwise error is $w(\mathbf{x}, y) \cdot \mathbb{I}[y \neq h(\mathbf{x})]$ for a given weight $w(\mathbf{x}, y) \geq 0$.

§8.4.ii. Minimizing E_{in}^w for Weighted Classification

Write E_{in}^w for the weighted in-sample error,

$$E_{\text{in}}^w(h) = \frac{1}{N} \sum_{n=1}^N \begin{cases} 1 & \text{if } y_n = +1 \\ 1000 & \text{if } y_n = -1 \end{cases} \cdot \mathbb{I}[y_n \neq h(\mathbf{x}_n)],$$

and $E_{\text{in}}^{0/1}$ for the plain fraction of mistakes.

Naïve thoughts.

- PLA: doesn't matter if linear separable. :-) Then some $\mathbf{w}$ has $E_{\text{in}}^w = 0$, whatever the weights, and PLA finds it.
- Pocket: modify the pocket-replacement rule — if $\mathbf{w}_{t+1}$ reaches smaller E_{in}^w than $\hat{\mathbf{w}}$, replace $\hat{\mathbf{w}}$ by $\mathbf{w}_{t+1}$.

Question 8.4.2. Pocket has some guarantee on $E_{\text{in}}^{0/1}$; does the modified pocket have a similar guarantee on E_{in}^w ?

§8.4.iii. Systematic Route: Connect E_{in}^w and $E_{\text{in}}^{0/1}$

Rather than re-prove anything about pocket, turn the weighted problem into an unweighted one.

original problem

		$h(\mathbf{x})$	
		+1	-1
y	+1	0	1
	-1	1000	0

$\mathcal{D}$:

$(\mathbf{x}_1, +1)$

$(\mathbf{x}_2, -1)$

$(\mathbf{x}_3, -1)$

...

$(\mathbf{x}_{N-1}, +1)$

$(\mathbf{x}_N, +1)$

equivalent problem

		$h(\mathbf{x})$	
		+1	-1
y	+1	0	1
	-1	1	0

$(\mathbf{x}_1, +1)$

$(\mathbf{x}_2, -1), (\mathbf{x}_2, -1), \dots, (\mathbf{x}_2, -1)$

$(\mathbf{x}_3, -1), (\mathbf{x}_3, -1), \dots, (\mathbf{x}_3, -1)$

...

$(\mathbf{x}_{N-1}, +1)$


$(\mathbf{x}_N, +1)$

After copying -1 examples 1000 times, E_{in}^w for LHS $\equiv E_{\text{in}}^{0/1}$ for RHS!

A mistake on $(\mathbf{x}_2, -1)$ costs 1000 on the left and is repeated on all 1000 copies on the right, so the two error counts agree term by term for every h . (The normalisations differ, N against the size of the copied set, but that is a constant factor and does not change which h minimises.) Pocket's guarantee for $E_{\text{in}}^{0/1}$ on the copied data therefore *is* a guarantee for E_{in}^w on the original.

§8.4.iv. Weighted Pocket Algorithm

Copying 1000 times is wasteful; the copies only need to exist in how often examples are picked and how errors are counted.



Using ‘virtual copying’, the **weighted pocket algorithm** includes:

- weighted PLA: randomly check -1 example mistakes with 1000 times more probability;
- weighted pocket replacement: if $\mathbf{w}_{t+1}$ reaches smaller $E_{\text{in}}^{\mathbf{w}}$ than $\hat{\mathbf{w}}$, replace $\hat{\mathbf{w}}$ by $\mathbf{w}_{t+1}$.

Against [algorithm 2.3](#), only two lines change:

Algorithm 8.1: Weighted Pocket Algorithm

Input: training examples $\mathcal{D}$, weights 1 for $y_n = +1$ and 1000 for $y_n = -1$, an iteration budget

Output: $\hat{\mathbf{w}}$, returned as g

```

1 initialize the pocket weights  $\hat{\mathbf{w}}$ ;
2 for  $t = 0, 1, \dots$  until enough iterations do
3   find a random mistake of  $\mathbf{w}_t$ ,  $(\mathbf{x}_{n(t)}, y_{n(t)})$ , picking each mistake with probability
   proportional to its weight;
4   (try to) correct the mistake by  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + y_{n(t)}\mathbf{x}_{n(t)}$ ;
5   if  $E_{\text{in}}^{\mathbf{w}}(\mathbf{w}_{t+1}) < E_{\text{in}}^{\mathbf{w}}(\hat{\mathbf{w}})$  then  $\hat{\mathbf{w}} \leftarrow \mathbf{w}_{t+1}$ ;
6 return  $\hat{\mathbf{w}}$  as  $g$ ;
```

Picking a random mistake of the copied data set is exactly picking an original mistake with probability proportional to its weight, and the copied data’s $E_{\text{in}}^{0/1}$ is $E_{\text{in}}^{\mathbf{w}}$ up to a constant; so this *is* pocket on the copied data, run without making the copies.

Systematic route (called ‘reduction’): can be applied to many other algorithms!

Exercise 8.4.3 (Fun Time). Consider the CIA cost matrix. If there are 10 examples with $y_n = -1$ (intruder) and 999,990 examples with $y_n = +1$ (you). What would $E_{\text{in}}^{\mathbf{w}}(h)$ be for a constant $h(\mathbf{x})$ that always returns $+1$?

		$h(\mathbf{x})$	
		+1	-1
y	+1	0	1
	-1	1000	0

(1) 0.001; (2) 0.01; (3) 0.1; (4) 1.

Reference Answer. (2). The constant h is wrong on exactly the 10 intruders, each costing 1000, out of $N = 10^6$ examples: $E_{\text{in}}^{\mathbf{w}} = 10 \cdot 1000 / 10^6 = 0.01$. While the quiz is a simple evaluation, it is not uncommon that the data is very unbalanced for such an application. Properly ‘setting’ the weights can be used to avoid the lazy constant prediction: under plain 0/1 error the same h would score 10^{-5} , and look nearly perfect. ■

Summary

Noise and Probabilistic Target can replace $f(\mathbf{x})$ by $P(y \mid \mathbf{x})$.

Error Measure affects the ‘ideal’ target.

Algorithmic Error Measure user-dependent $\Rightarrow$ plausible or friendly.

Weighted Classification easily done by virtual ‘example copying’.

The *Why can machines learn?* part closes here: with noise and a general error measure, the VC philosophy still says a finite d_{VC} and enough data make E_{in} a trustworthy proxy for E_{out} . What remains is to make E_{in} small, and so far the only tools are PLA and pocket. Next: more algorithms, please? :-)